Course Reminders

- Guest Lecture Thursday
- A3 due Friday (11:59 PM)
- Project Proposal Feedback Available
  - Median: 8.8/10; Mean: 8.5/10
  - Regrades: Piazza *proposal-regrade* ; all handled by Prof Ellis
- Updates:
  - Piazza - separate folders for each assignment (i.e. A3 = *a3* folder/tag)
  - Sam's OH: Fri 11 AM-12 PM

Project Proposal: General Themes (cont'd)

- Overall:
  - Excellent proposal but research question seems overly simple — you could feasibly answer your question with a single correlation table. How will your project go beyond this?
- Data:
  - For each dataset chosen, the proposal should describe the observations (what each row means). This proposal only describes the variables for each dataset. The proposal is also missing the number of observations within each dataset needed for this analysis.
  - It's unclear how {extra_columns} are related to your research question and hypotheses.
- Background:
  - Missing introduction to topic
  - Missing why this question is of interest to your group
  - Missing how the sources cited lead to the hypotheses of this proposal.
- Ethics:
  - How might potential sources of bias affect the analysis?

Inferential Analysis

Shannon E. Ellis, Ph.D
UC San Diego

Department of Cognitive Science
sellis@ucsd.edu

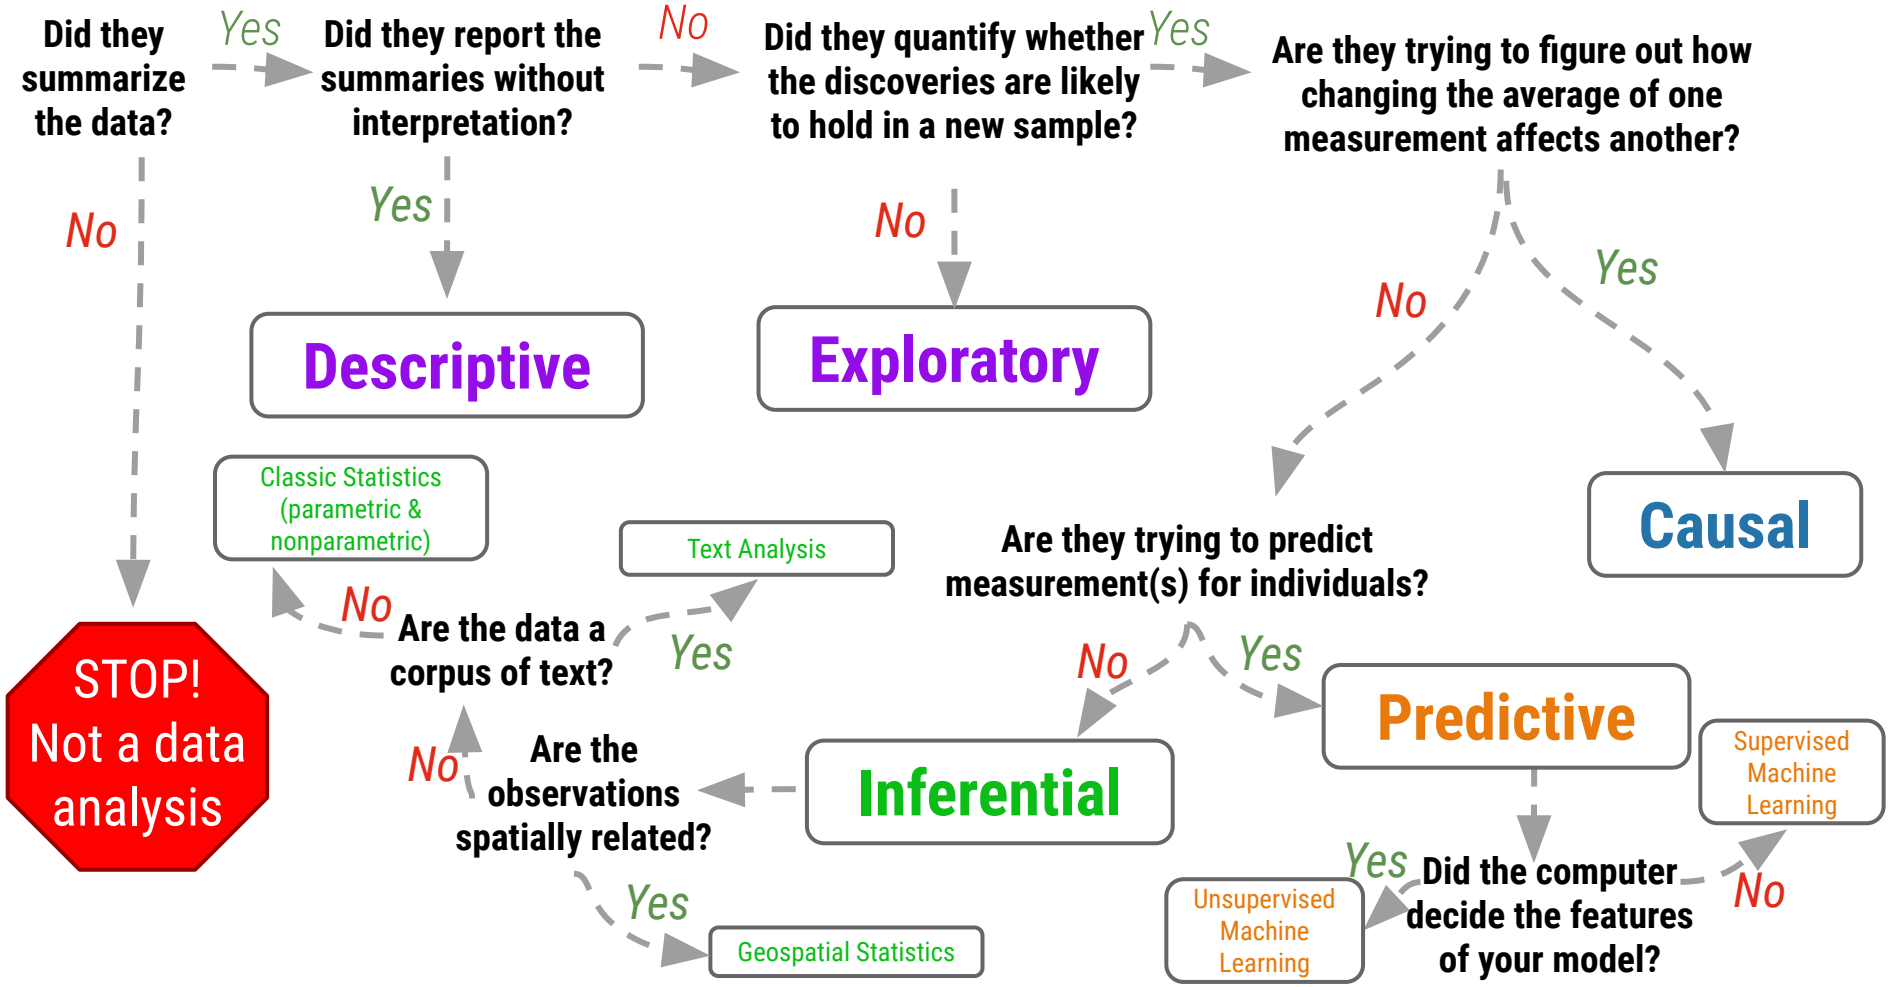
Did they summarize the data?
Yes
Did they report the summaries without interpretation?
No
Did they quantify whether the discoveries are likely to hold in a new sample?
Yes
Are they trying to figure out how changing the average of one measurement affects another?
No
Yes
Causal
No
Descriptive
Yes
No
Exploratory
Yes
Classic Statistics (parametric & nonparametric)
Text Analysis
Are they trying to predict measurement(s) for individuals?
No
Are the data a corpus of text?
Yes
No
Yes
Predictive
Supervised Machine Learning
STOP! Not a data analysis
No
Are the observations spatially related?
Inferential
Yes
Did the computer decide the features of your model?
Yes
No
Unsupervised Machine Learning
Yes
Geospatial Statistics

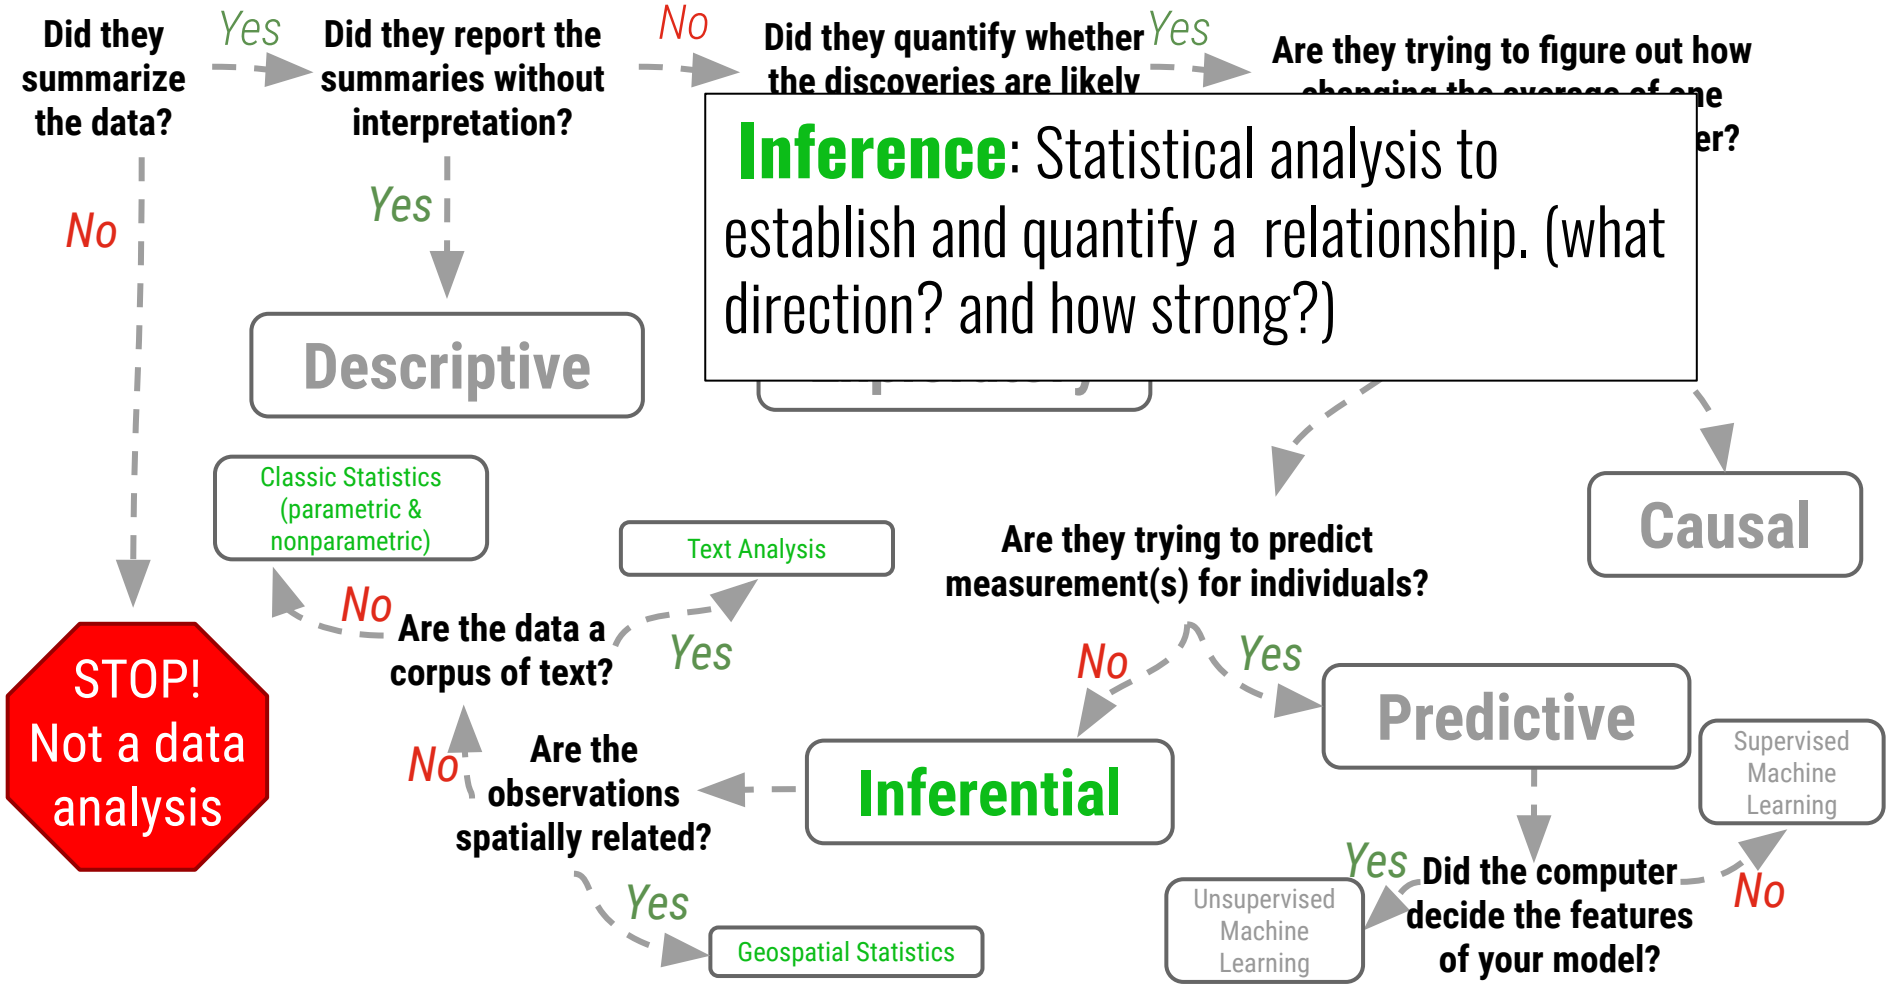
Did they summarize the data?
Yes
Did they report the summaries without interpretation?
No
Did they quantify whether the discoveries are likely
Yes
Are they trying to figure out how
No
Yes
Descriptive
Inference: Statistical analysis to establish and quantify a relationship. (what direction? and how strong?)
Classic Statistics (parametric & nonparametric)
Text Analysis
Are they trying to predict measurement(s) for individuals?
Causal
No
Are the data a corpus of text?
Yes
STOP! Not a data analysis
No
Yes
Predictive
No
Are the observations spatially related?
Inferential
Supervised Machine Learning
Yes
Did the computer decide the features of your model?
No
Unsupervised Machine Learning
Yes
Geospatial Statistics

- **Problem:** Does Sesame Street affect kids brain development?
- **Data science question:** What is the relationship between watching Sesame Street and test scores among children?
- **Type of analysis:** Inferential analysis

Sesame Street viewership

??

Test scores

Establishing & Stating Your Null and Alternative Hypotheses Helps Guide Your Analysis

Null Hypothesis:

H_0: Sesame Street has *no effect* on kids brain development

Alternative Hypothesis:

H_a: Watching Sesame Street *has an effect* on kids' brain development

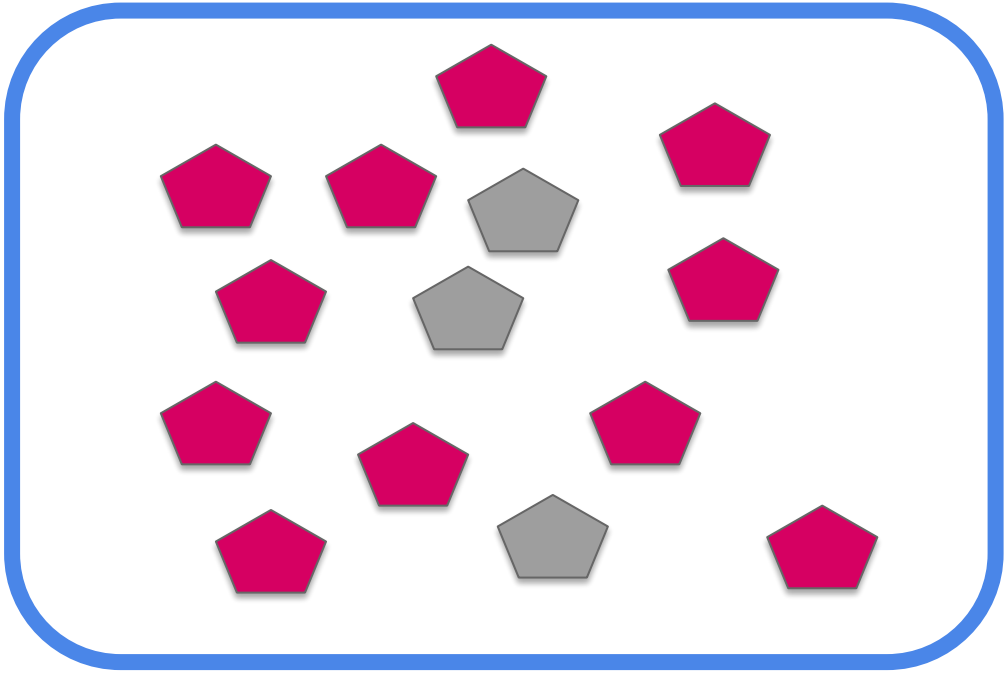

Population

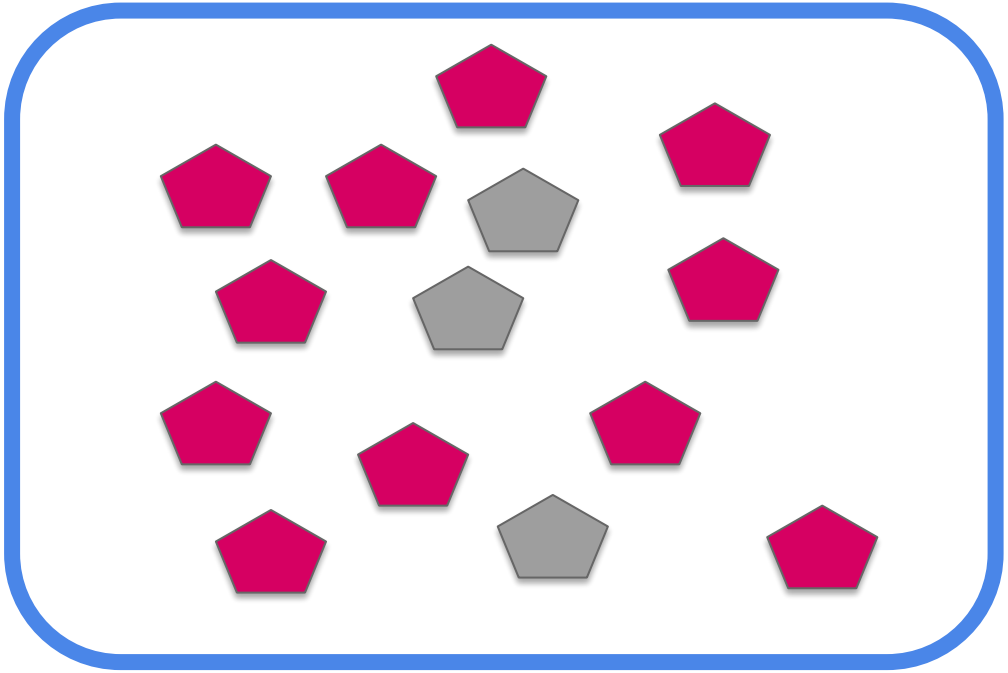

Population

In our Sesame street example, the population would be all children

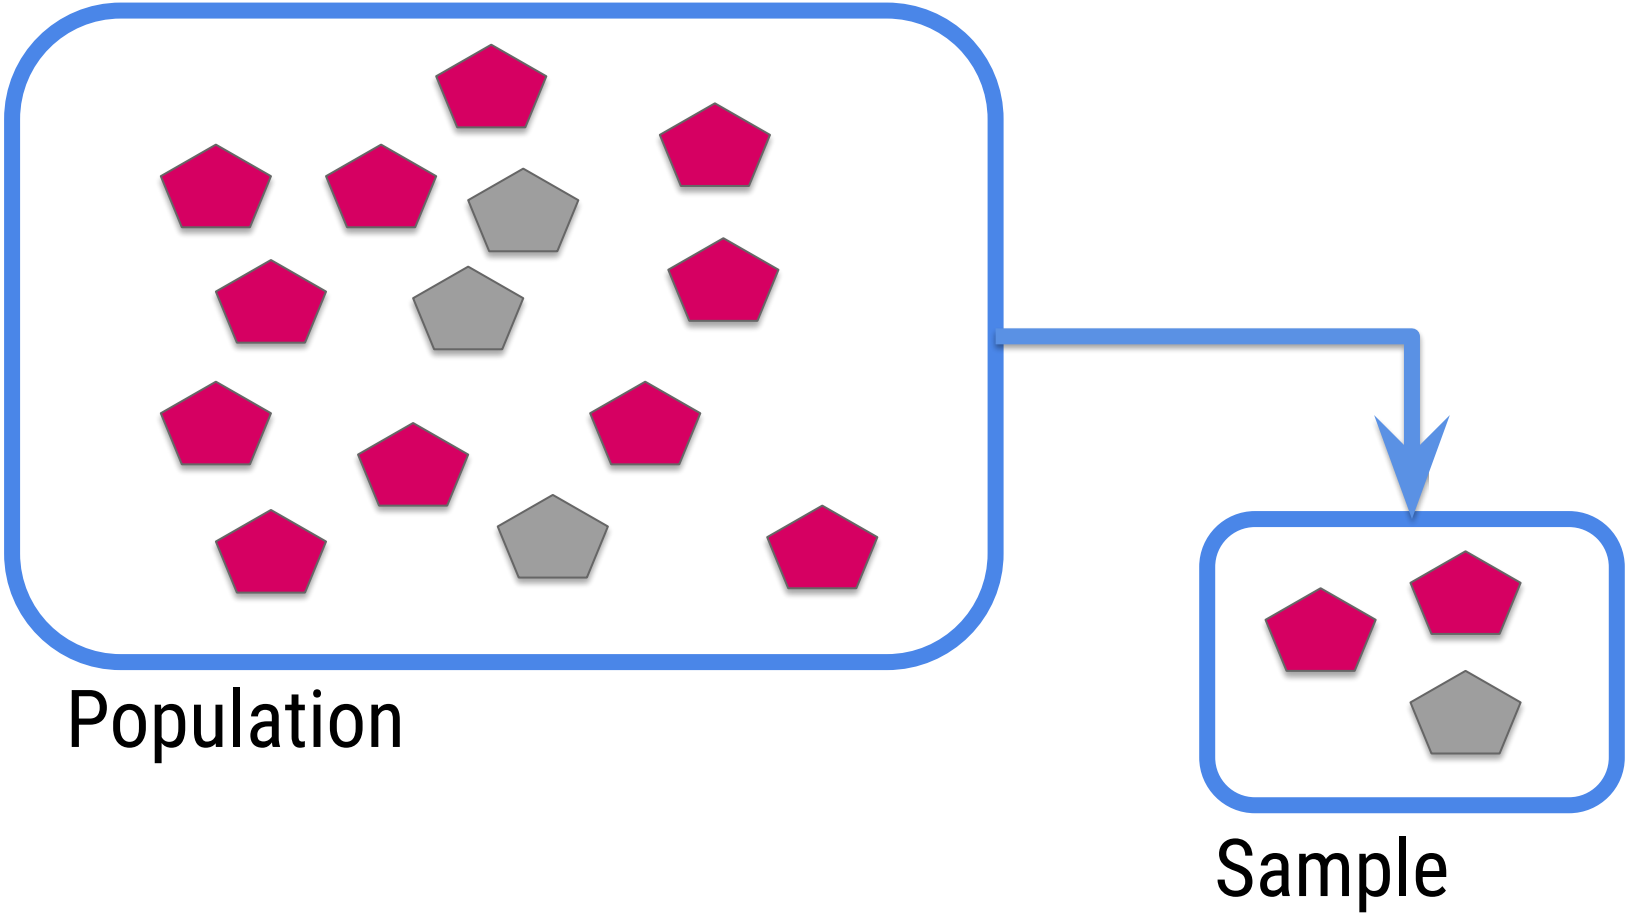
Population
Sample

Population

In our Sesame street example, the sample would be the children included in the study

Sample

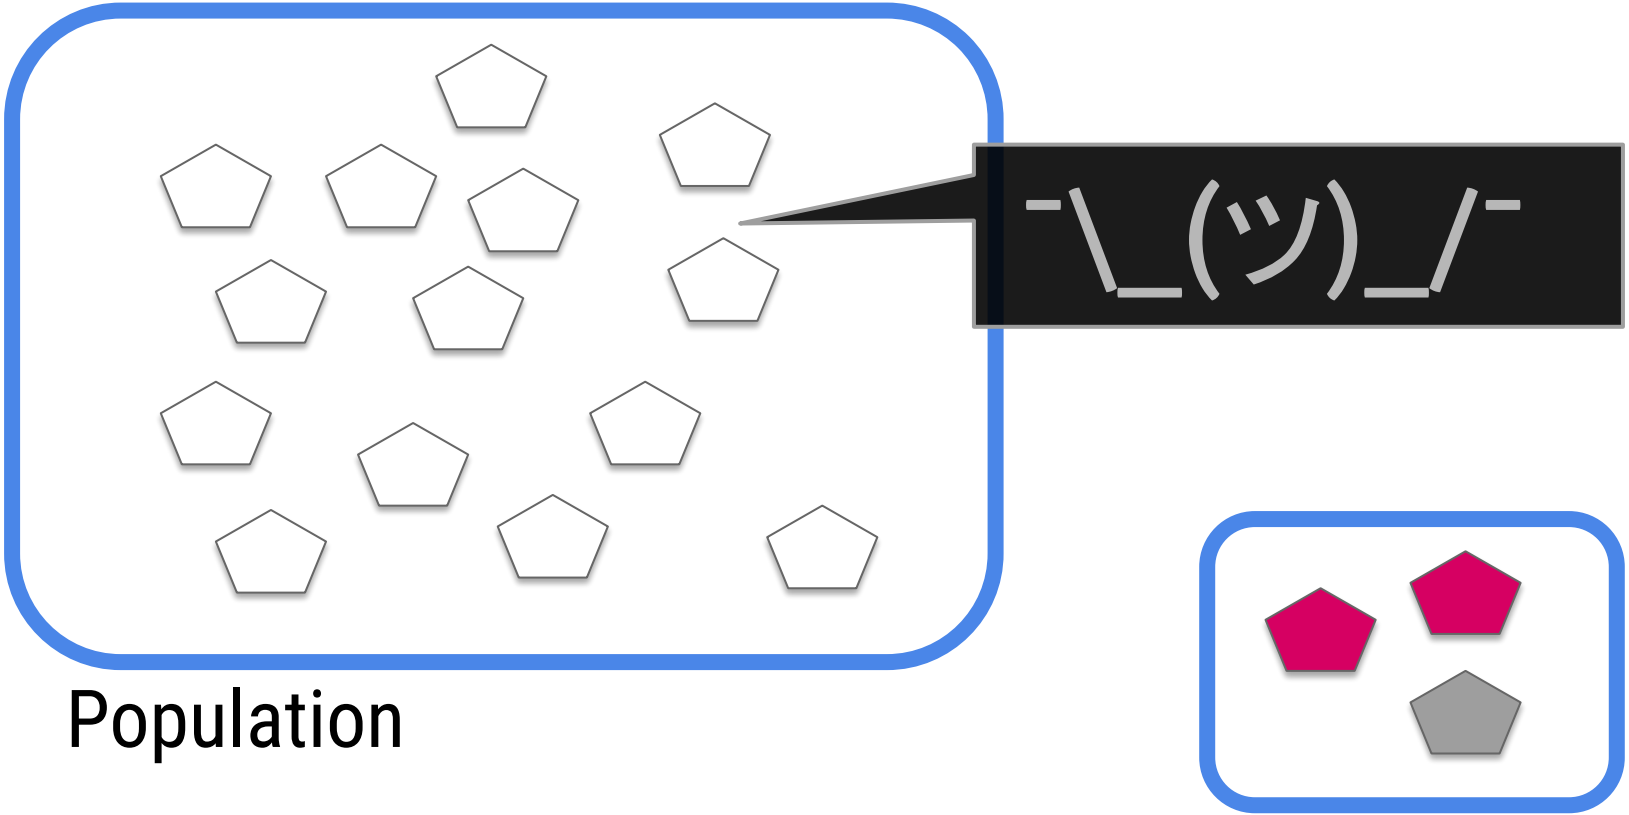

Population

Sample

We don't know how much Sesame street was watched by or the tests scores of <u>all kids</u>

¯_(ツ)_/¯

Population

Sample

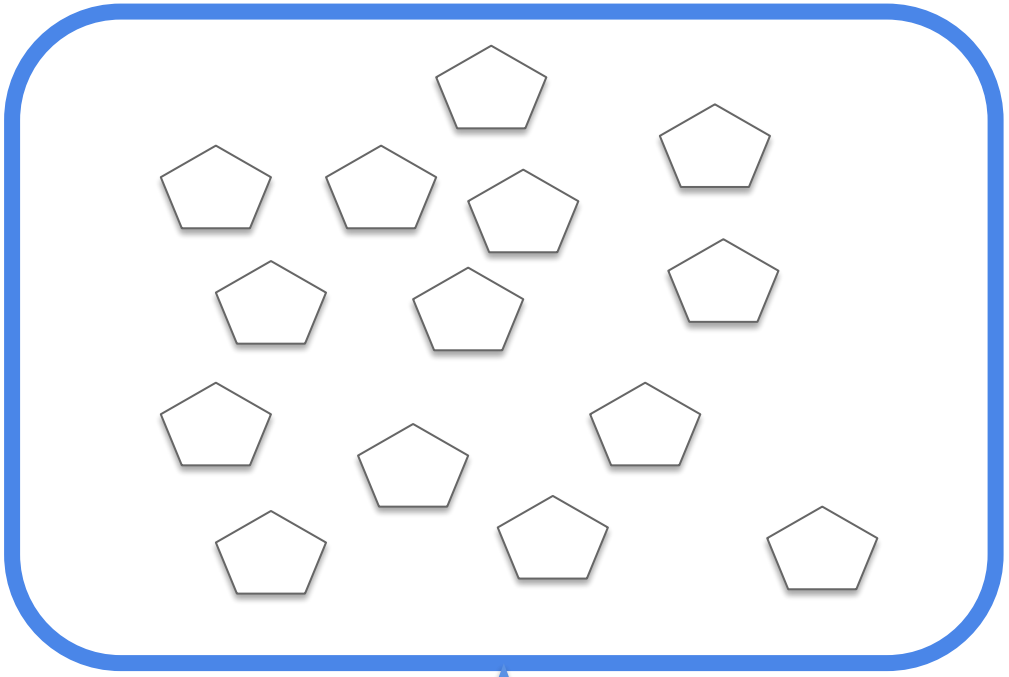

Population

Based on the relationship we see in our sample, we can infer the answer to our question in our population

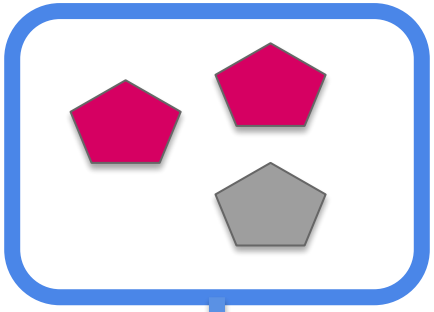

Sample

Inference!

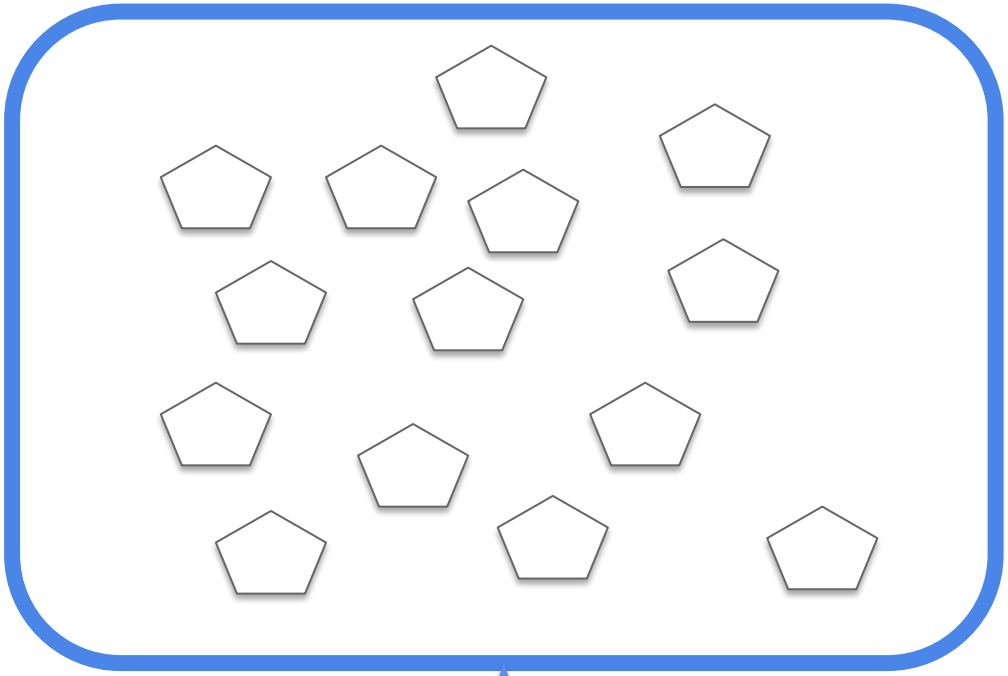

Population

So we look at Sesame street viewing and test scores in a representative sample of kids

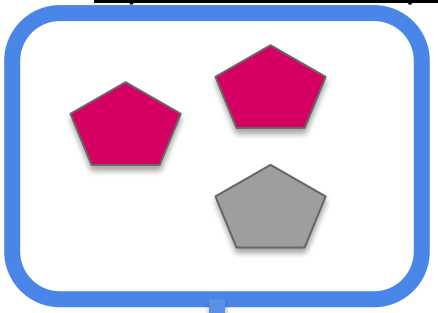

Sample

Inference!

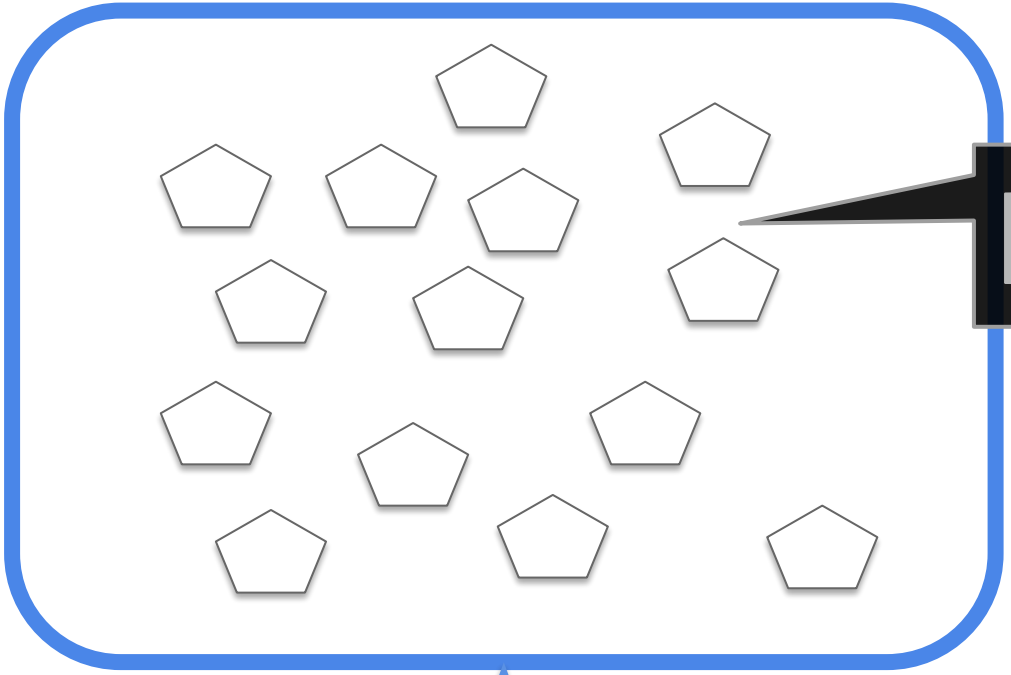

Best guess

So we look at Sesame street viewing and test scores in a representative sample of kids

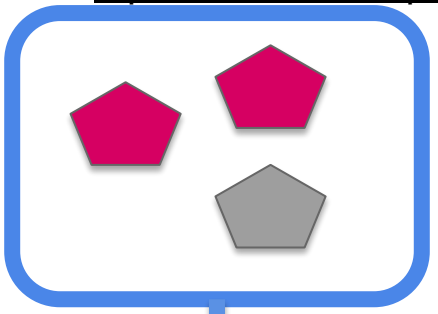

Population

Sample

Inference!

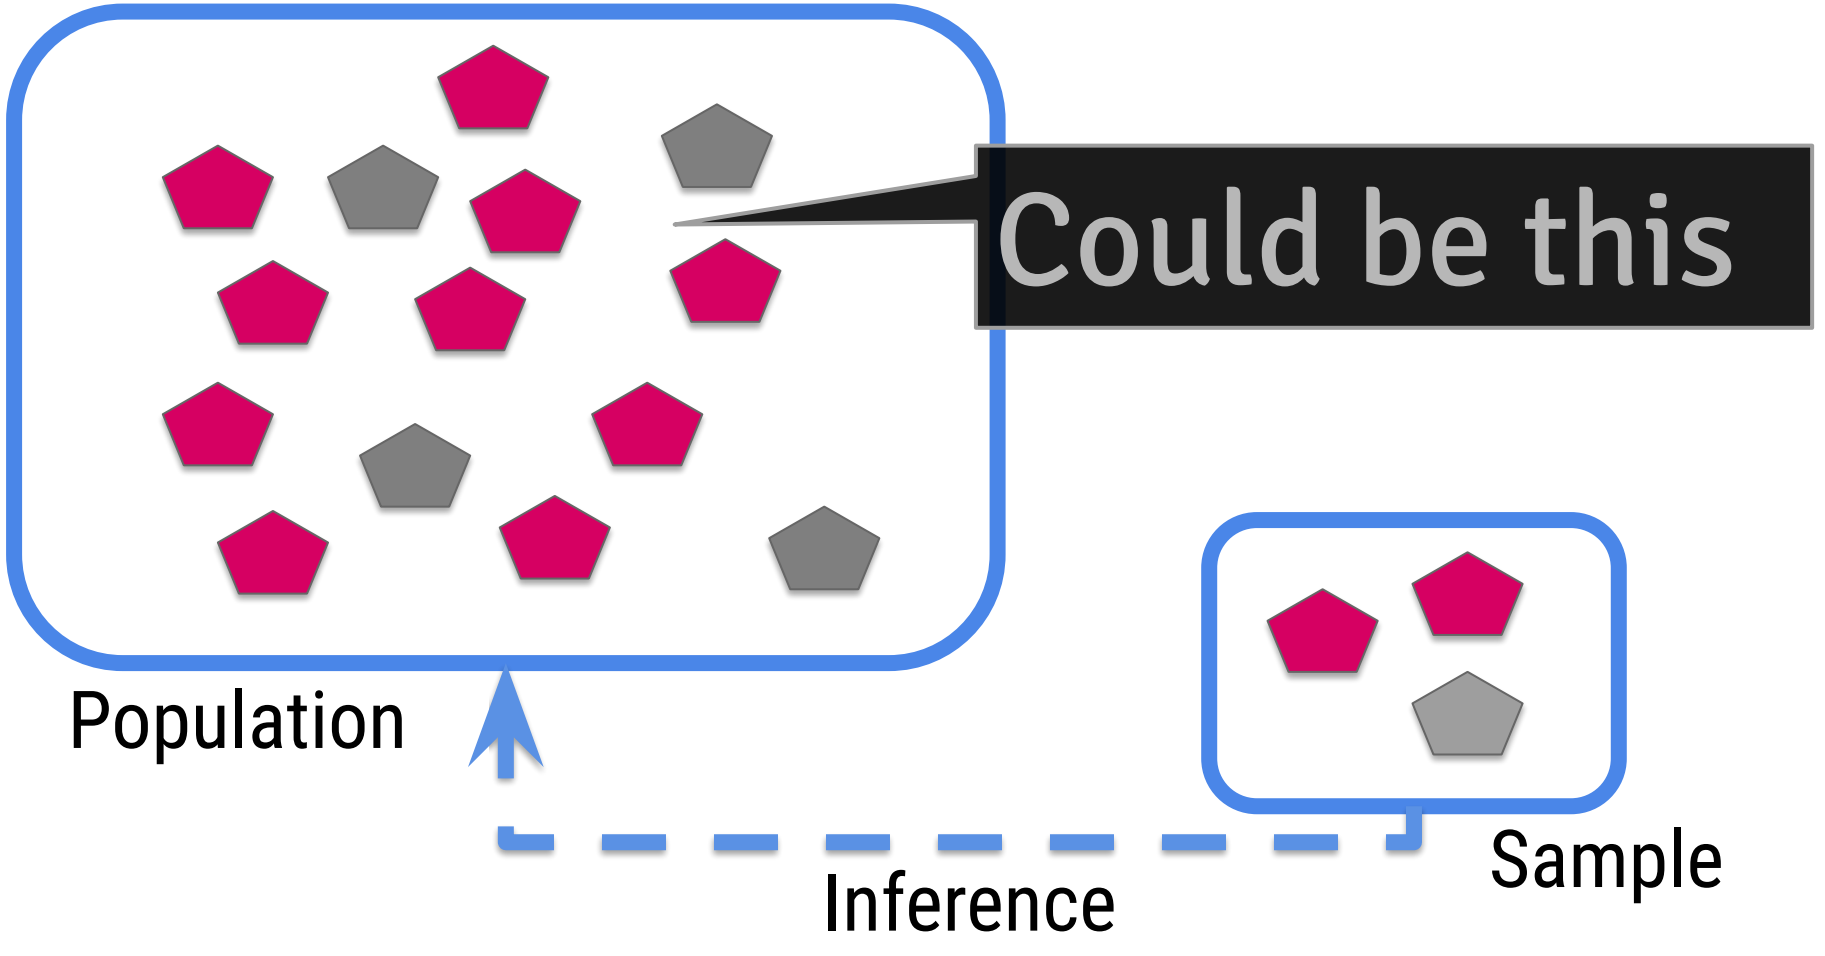
Could be this
Population
Inference
Sample

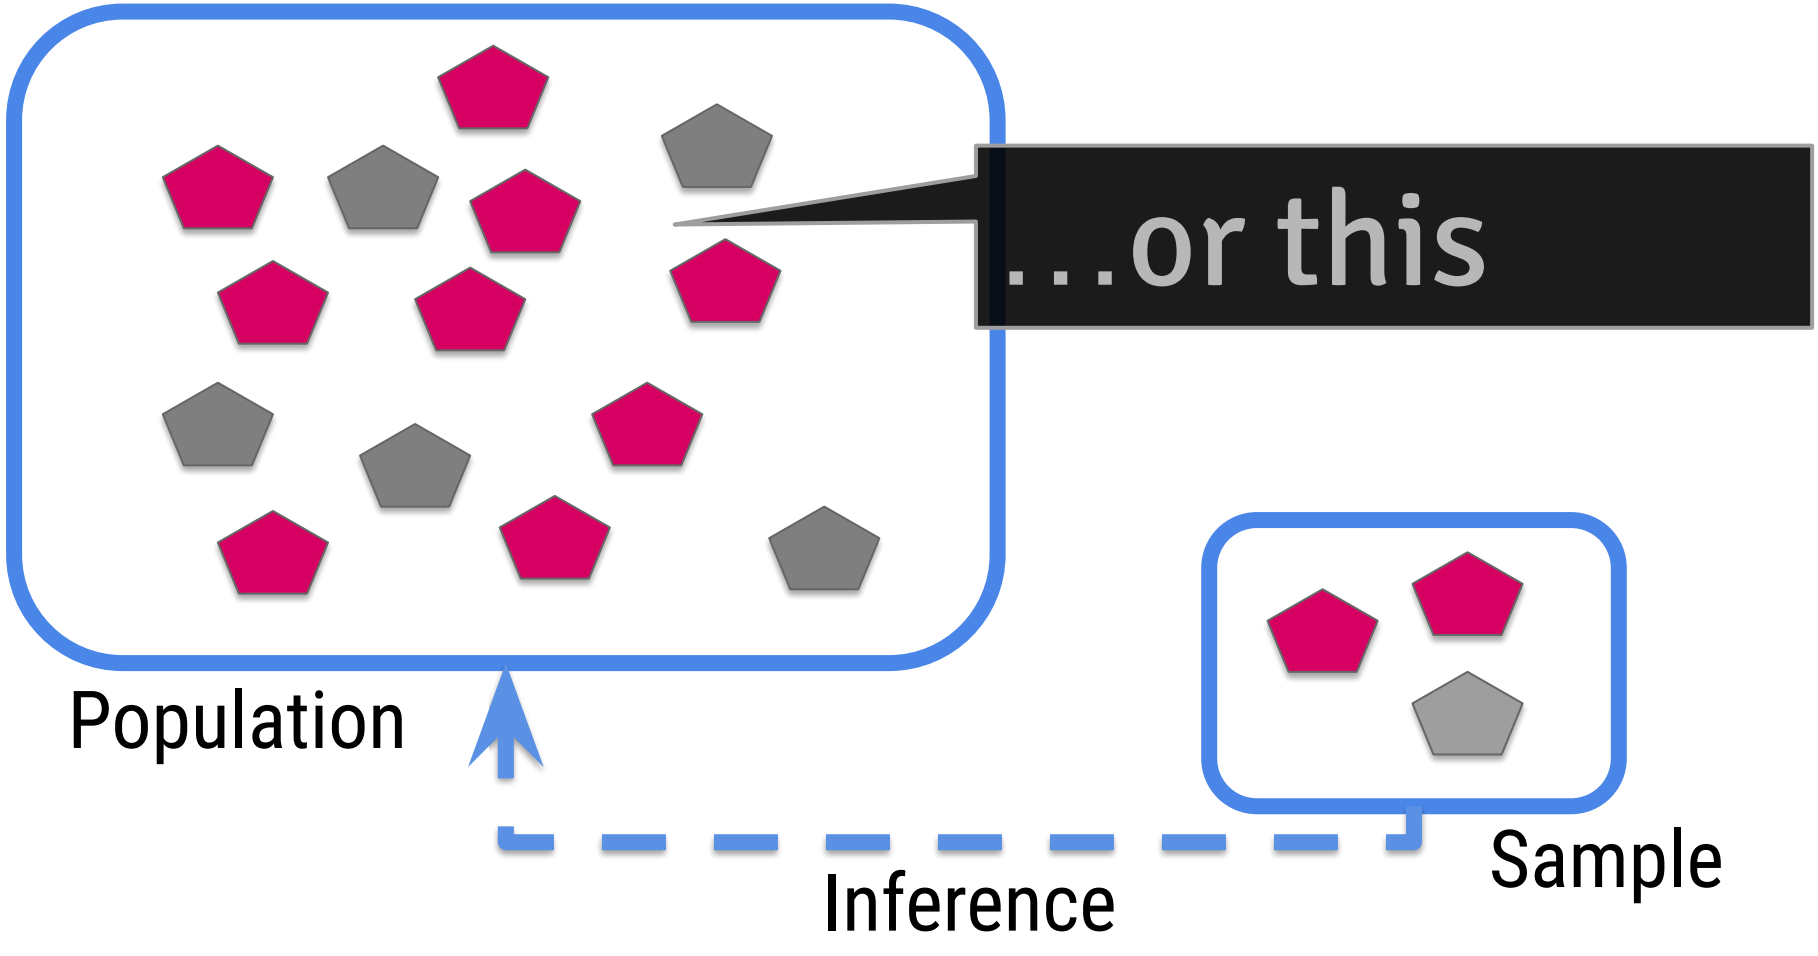
…or this
Population
Sample
Inference

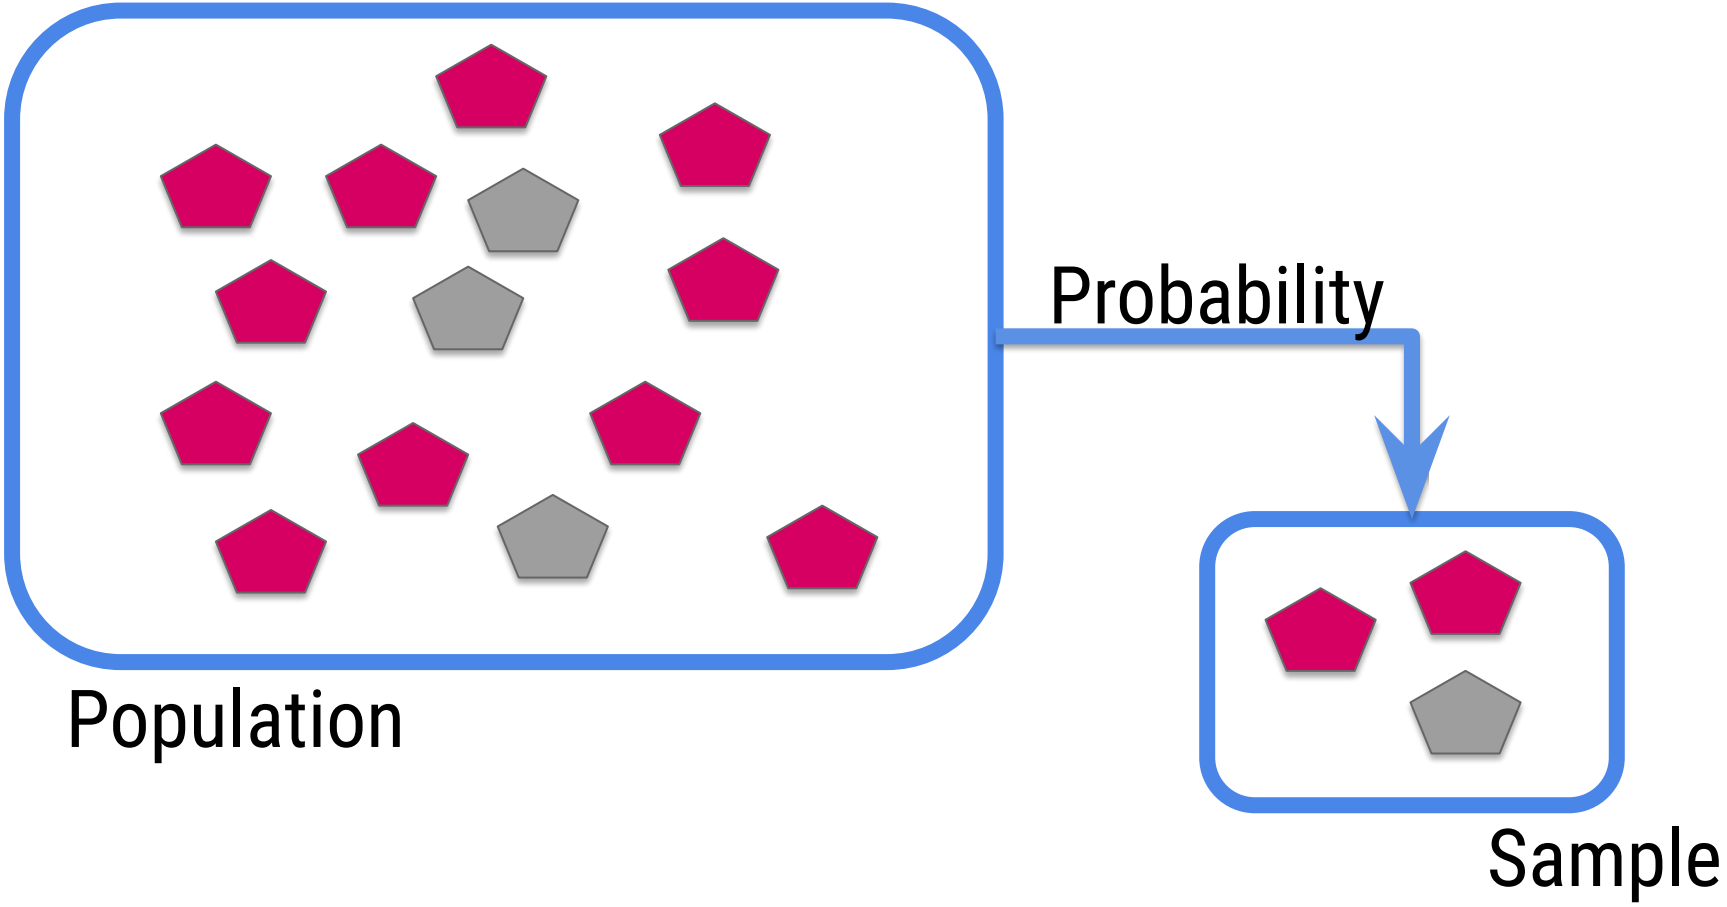
Probability
Population
Sample

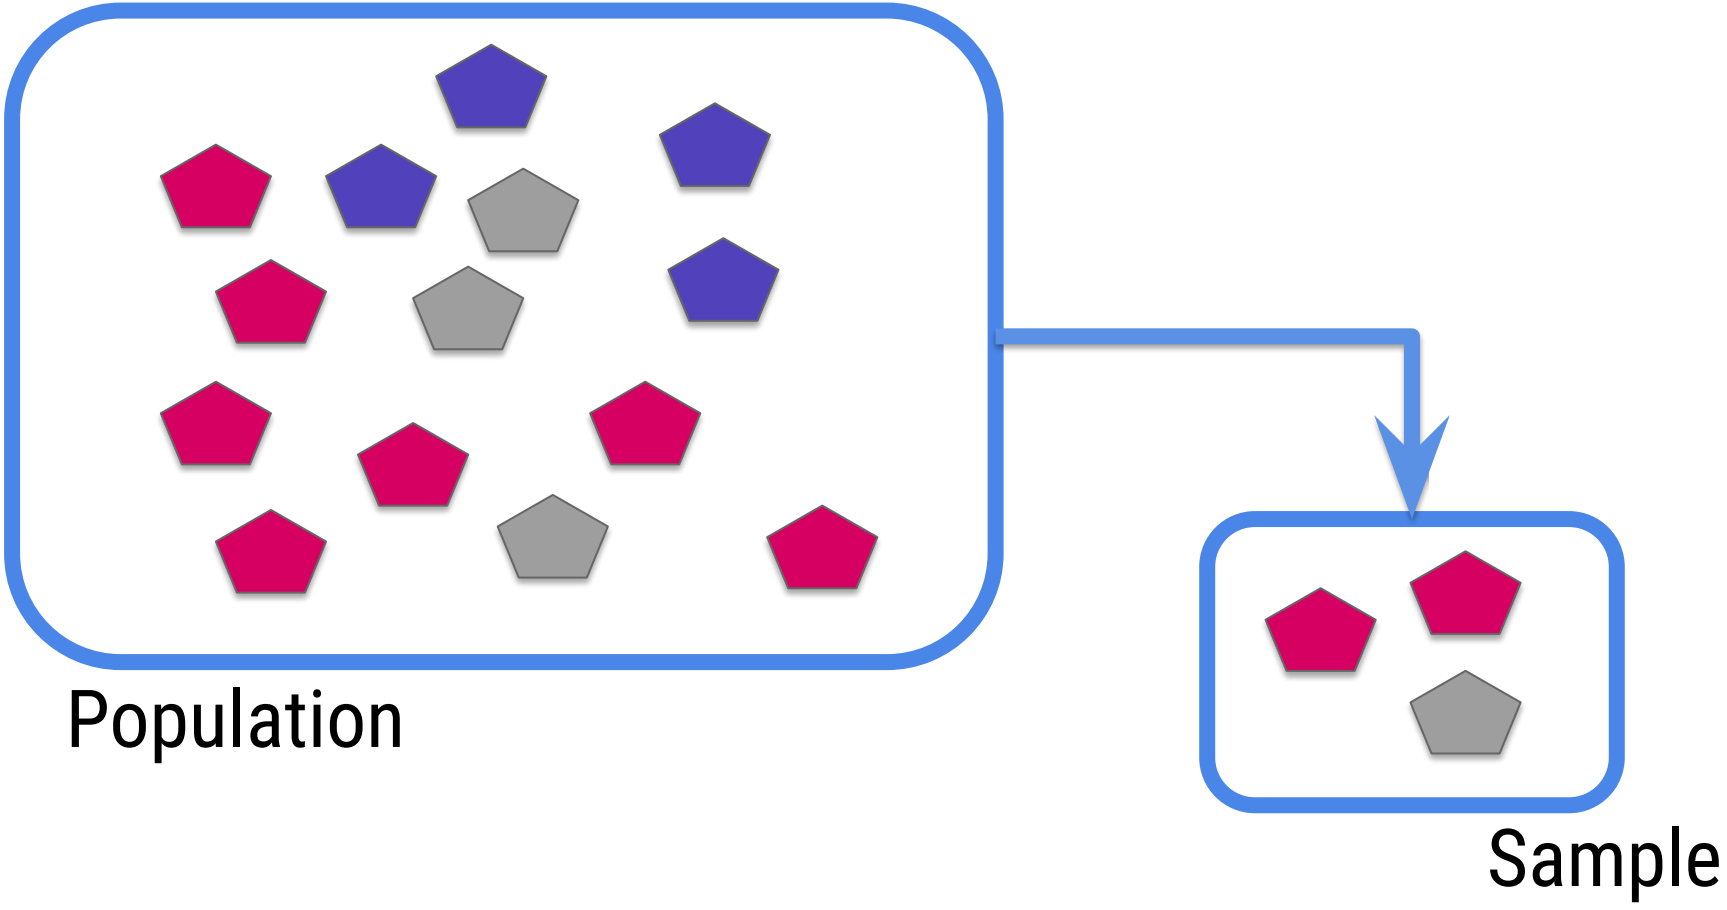
Population
Sample

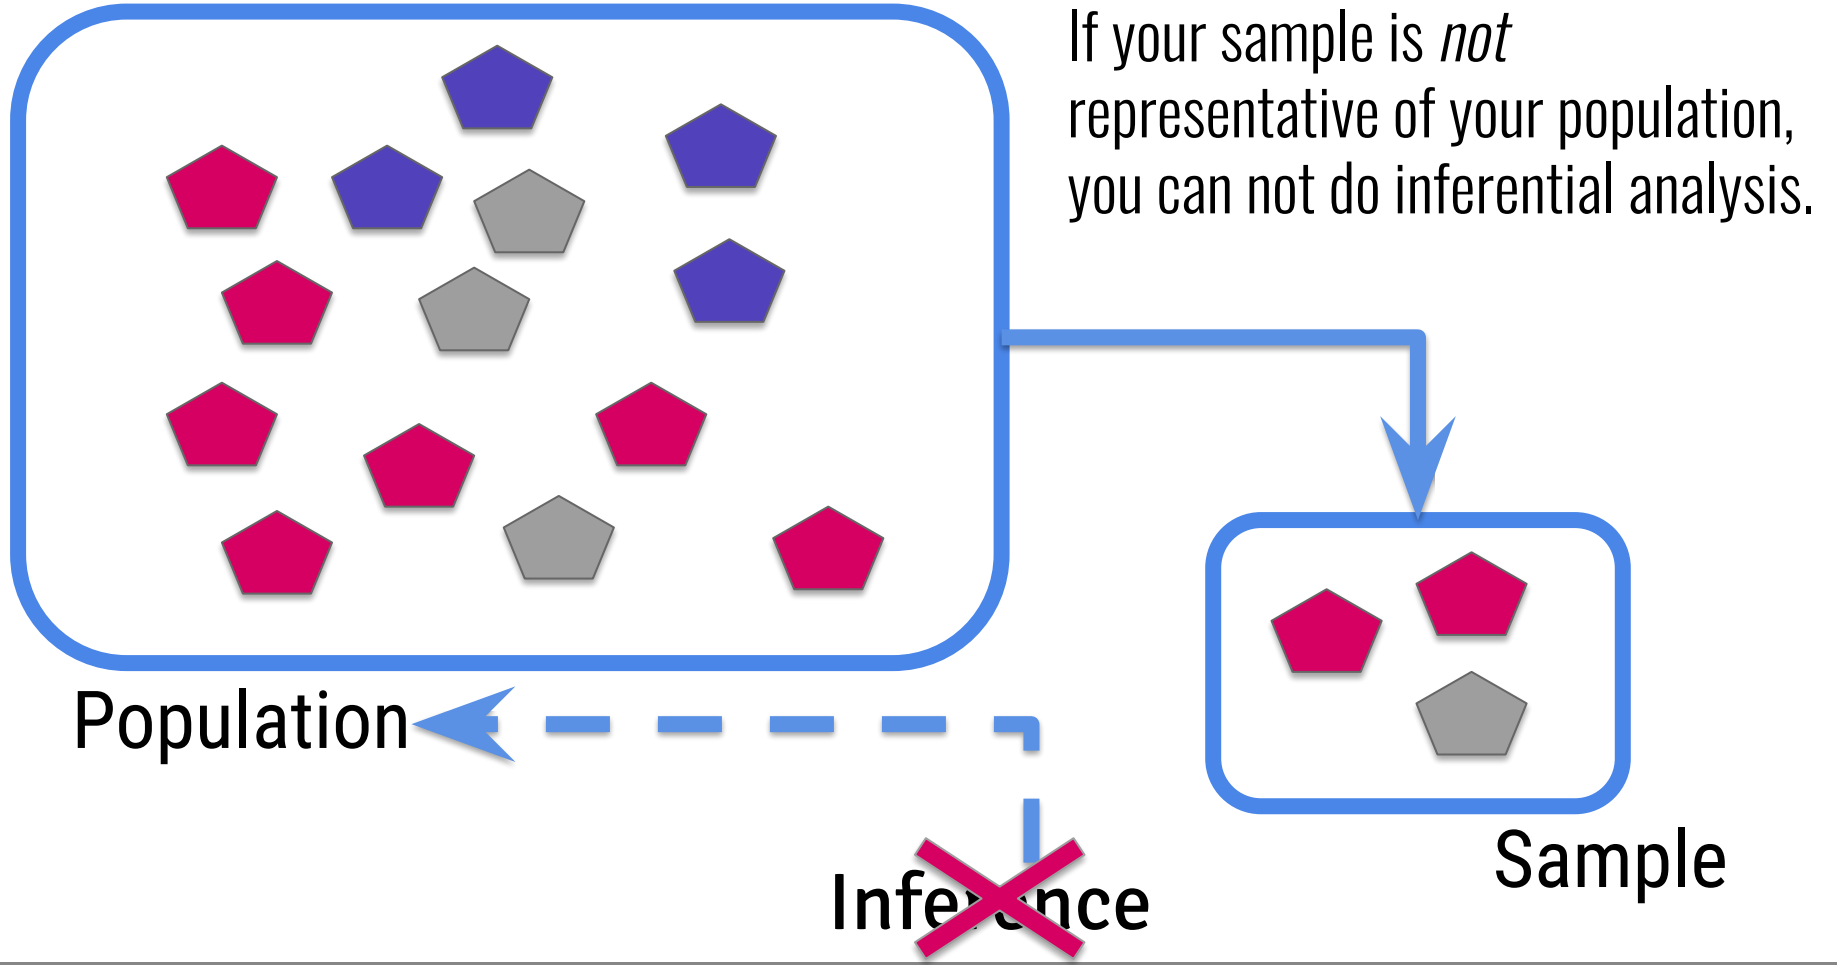
If your sample is *not* representative of your population, you can not do inferential analysis.
Population
Sample
Inference

Approaches to Inference

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN VARIABLES

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN VARIABLES

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

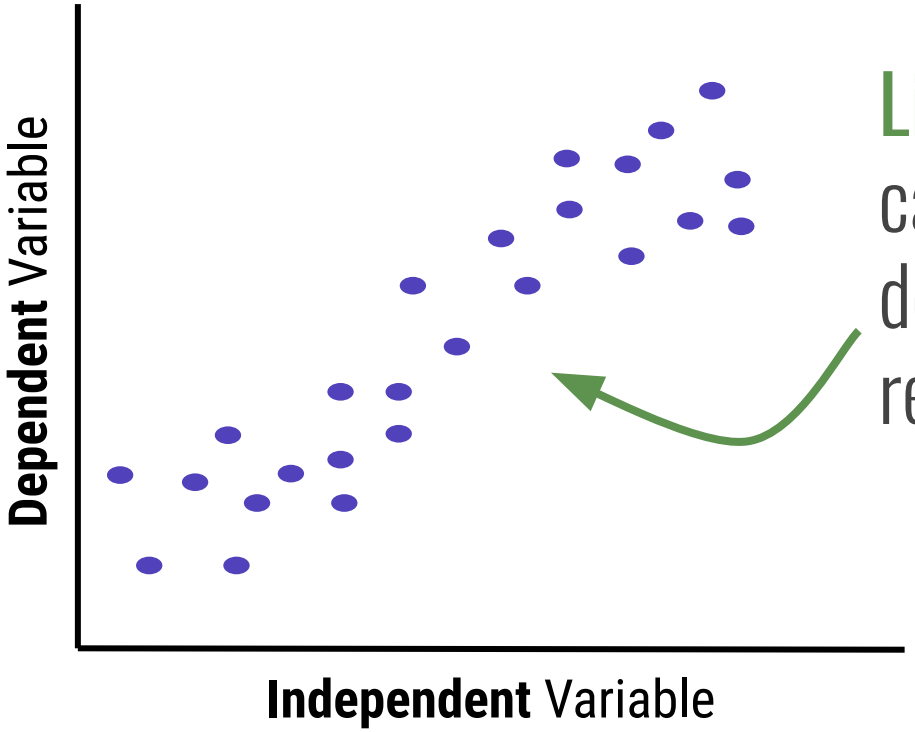

Linear regression can be used to describe this relationship

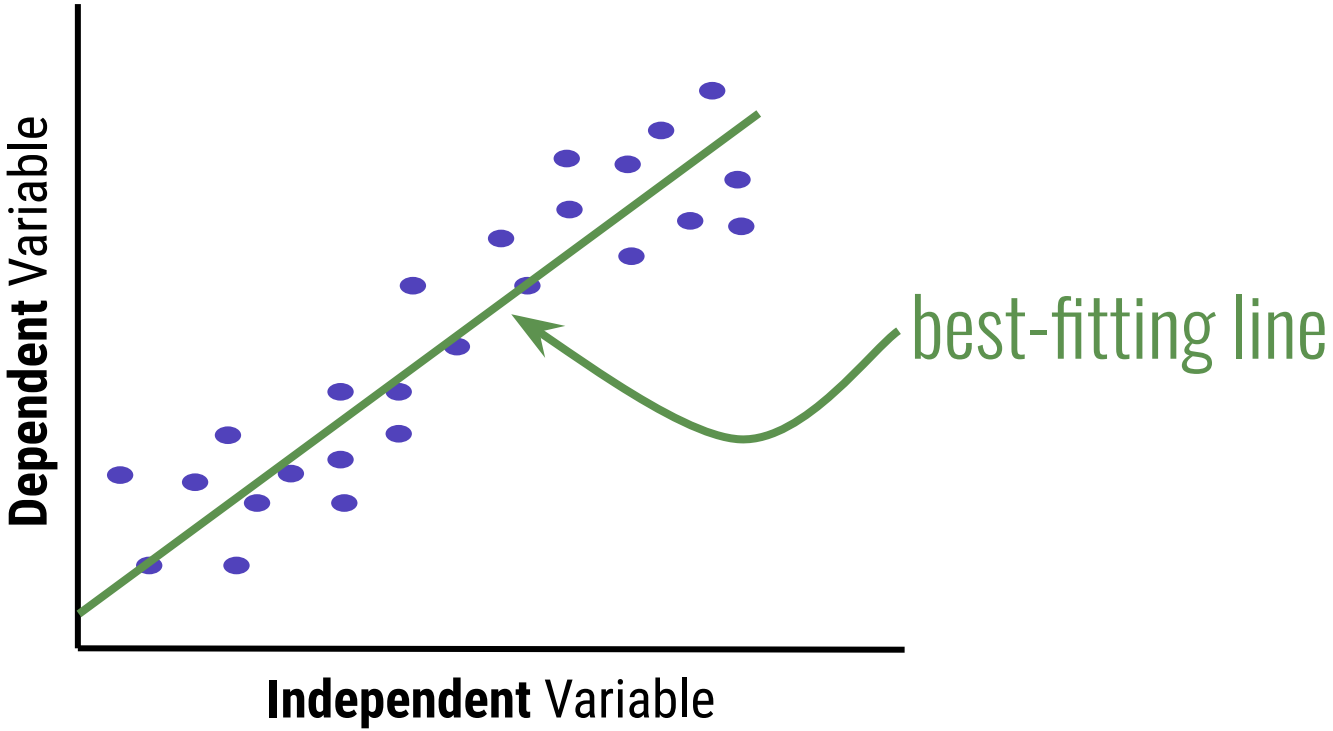
Dependent Variable
Independent Variable
best-fitting line

Best-fitting line

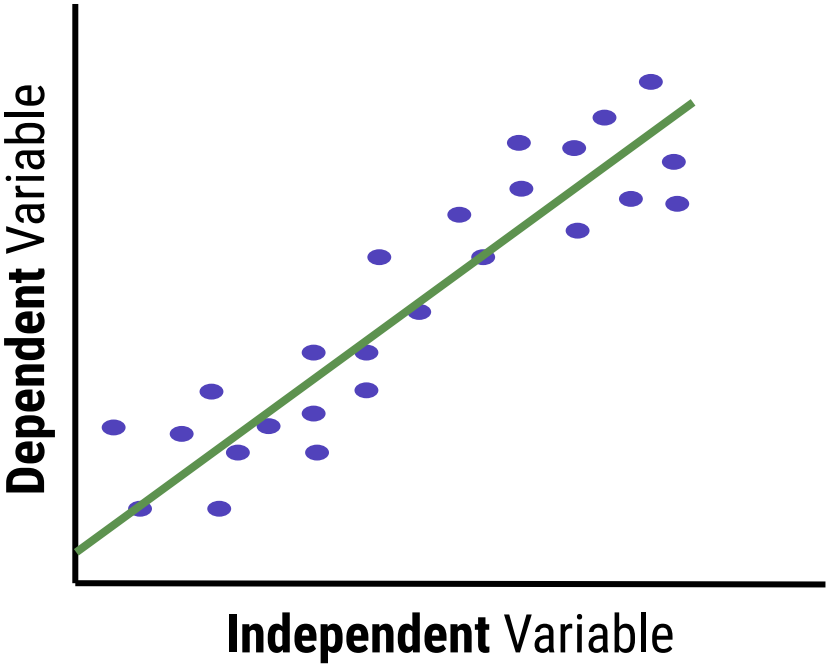

NOT a best-fitting line

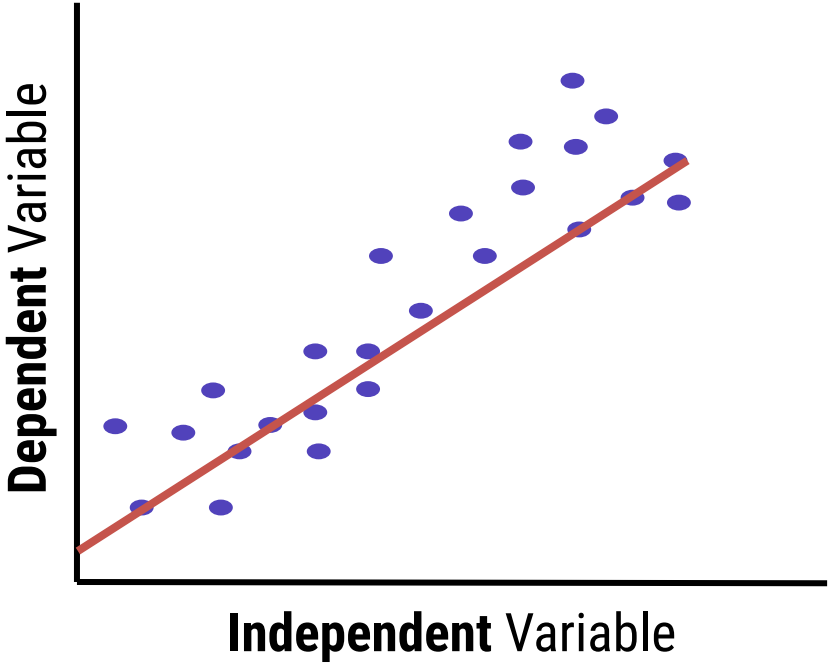

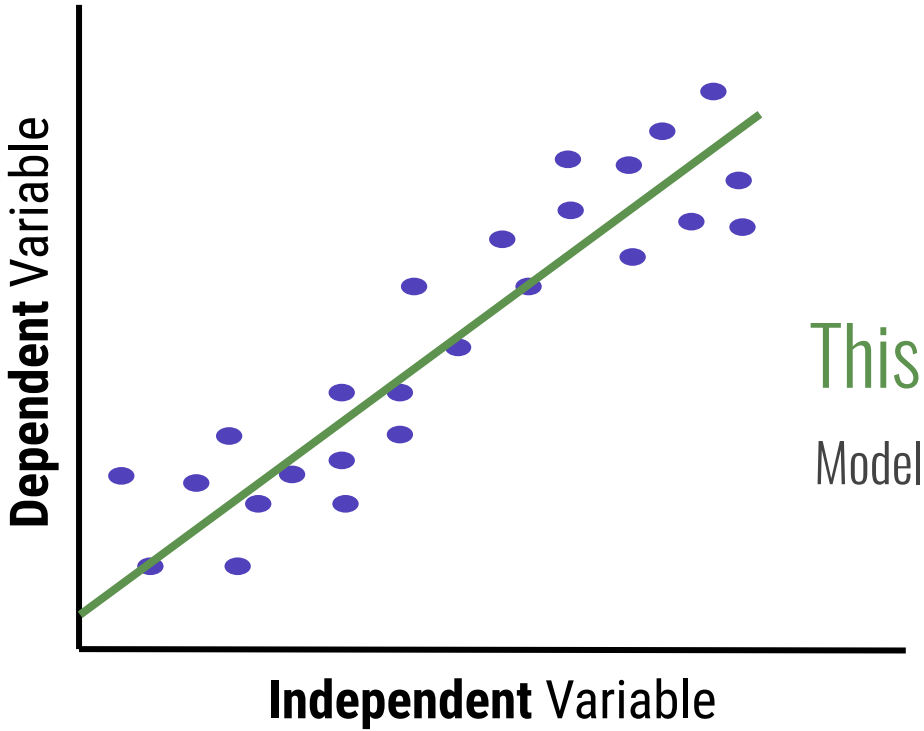

This line is a **model** of the data

Models are mathematical equations generated to *represent* the real life situation

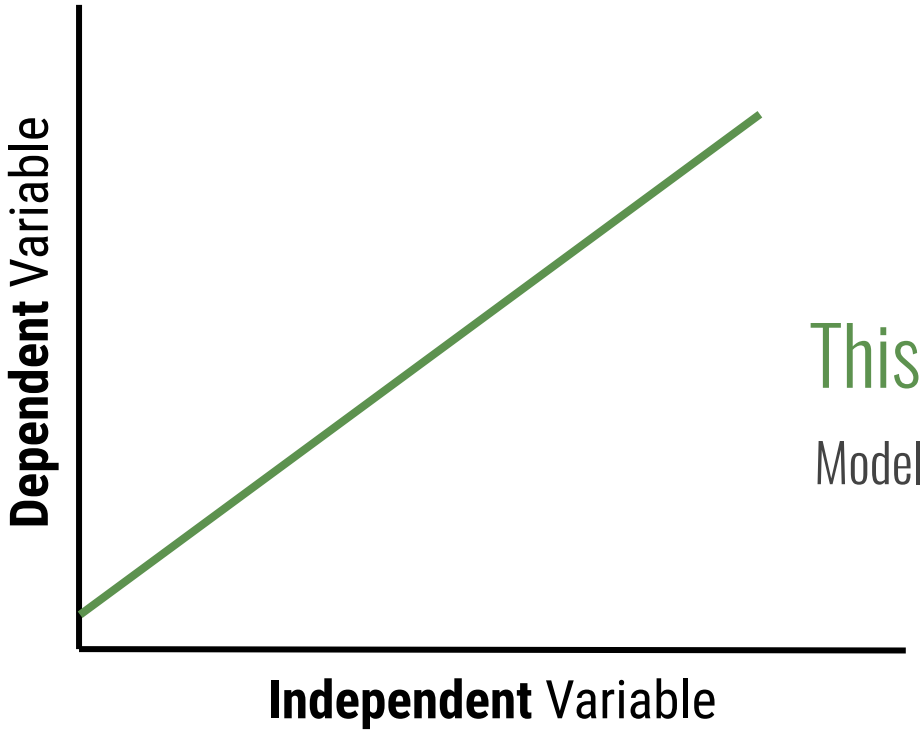

This line is a **model** of the data

Models are mathematical equations generated to *represent* the real life situation

“*All models are wrong, but some are useful*”

-George Box (British Statistician, *JASA* 1976)

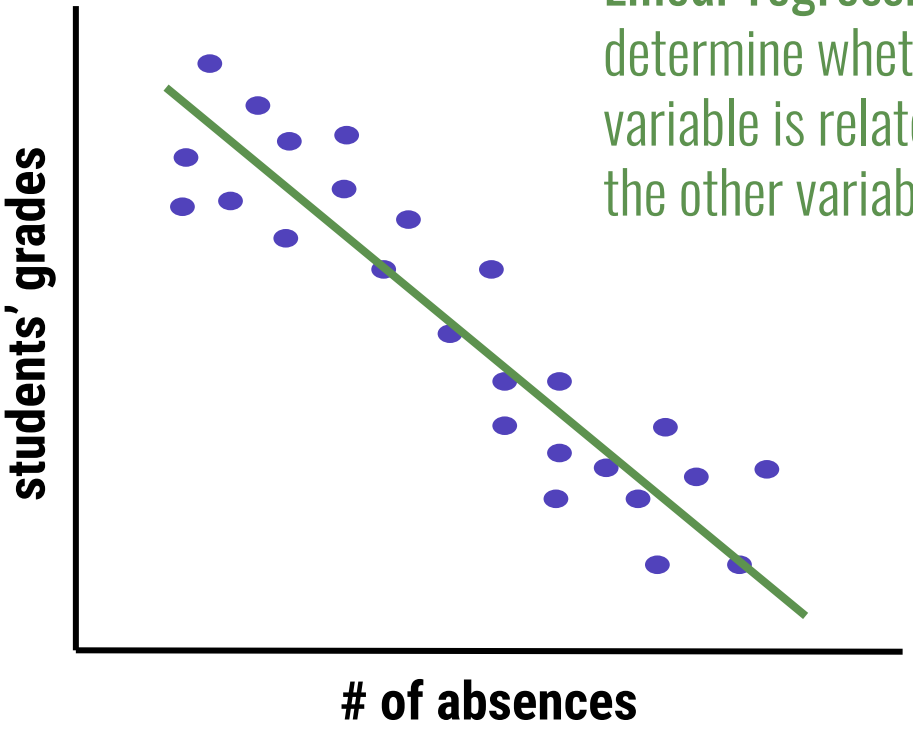

Linear regression can be used to determine whether a change in one variable is related to the change in the other variable

Linear regression can be used to determine whether a change in one variable is related to the change in the other variable

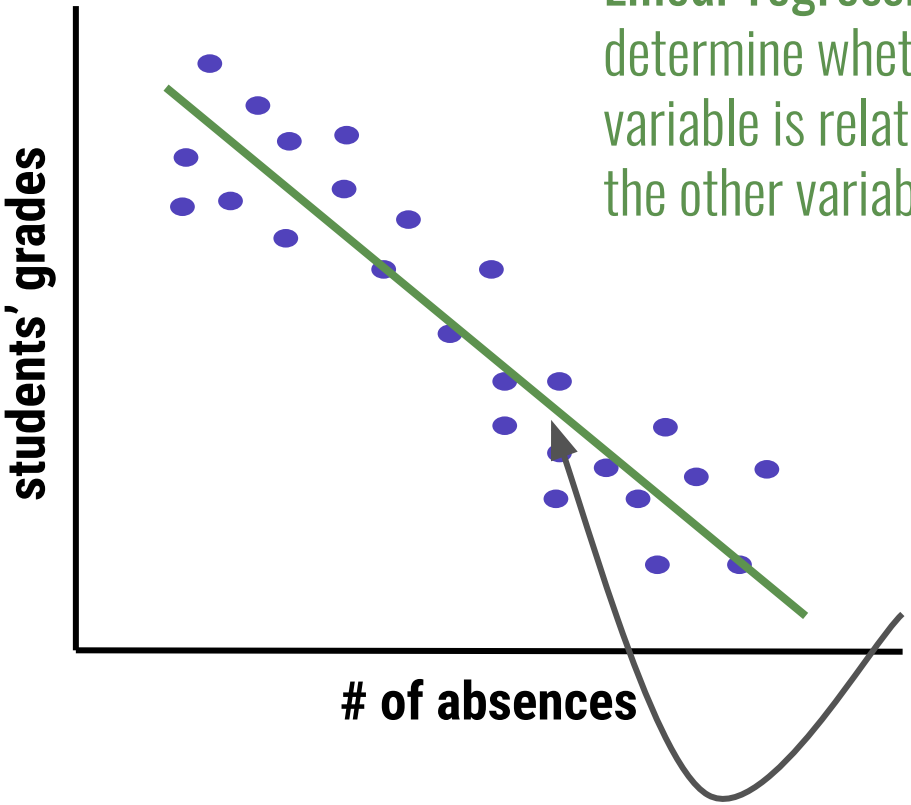

The magnitude of the relationship is measured by the slope of the line

Linear regression can be used to determine whether a change in one variable is related to the change in the other variable

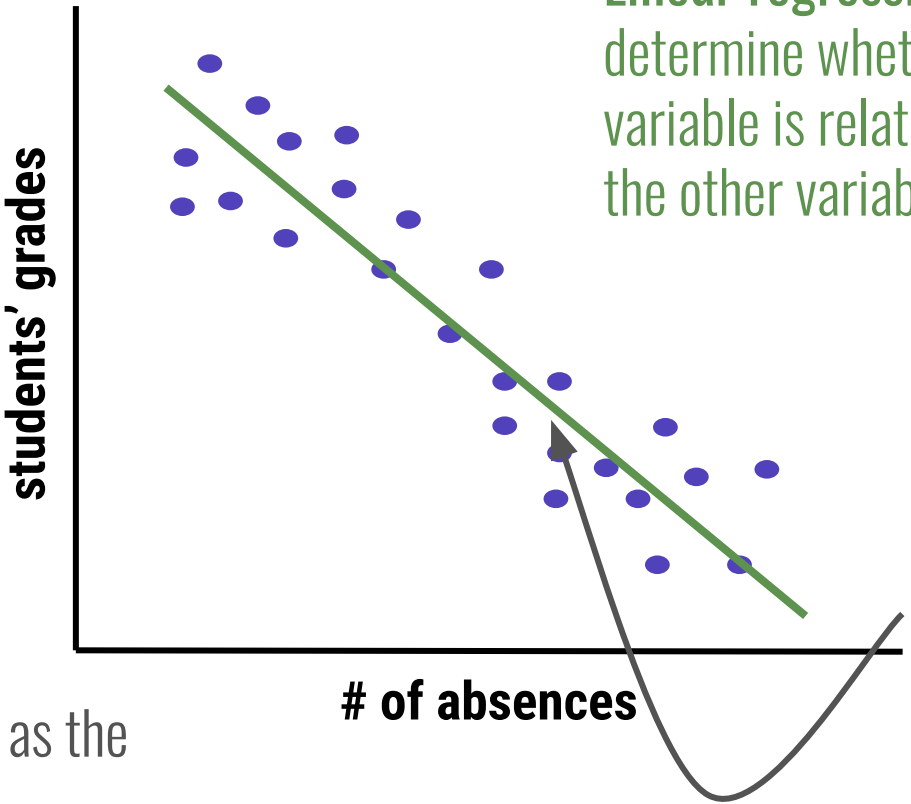

The magnitude of the relationship is measured by the slope of the line

This is also referred to as the model's effect size (β)

Effect size (β) can be estimated using the slope of the line

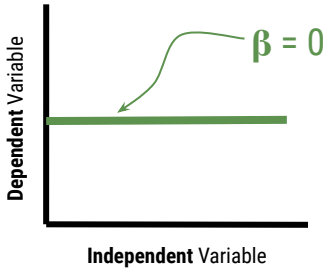

Effect size (β) can be estimated using the slope of the line

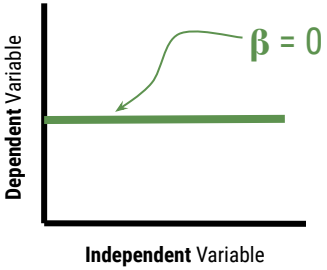

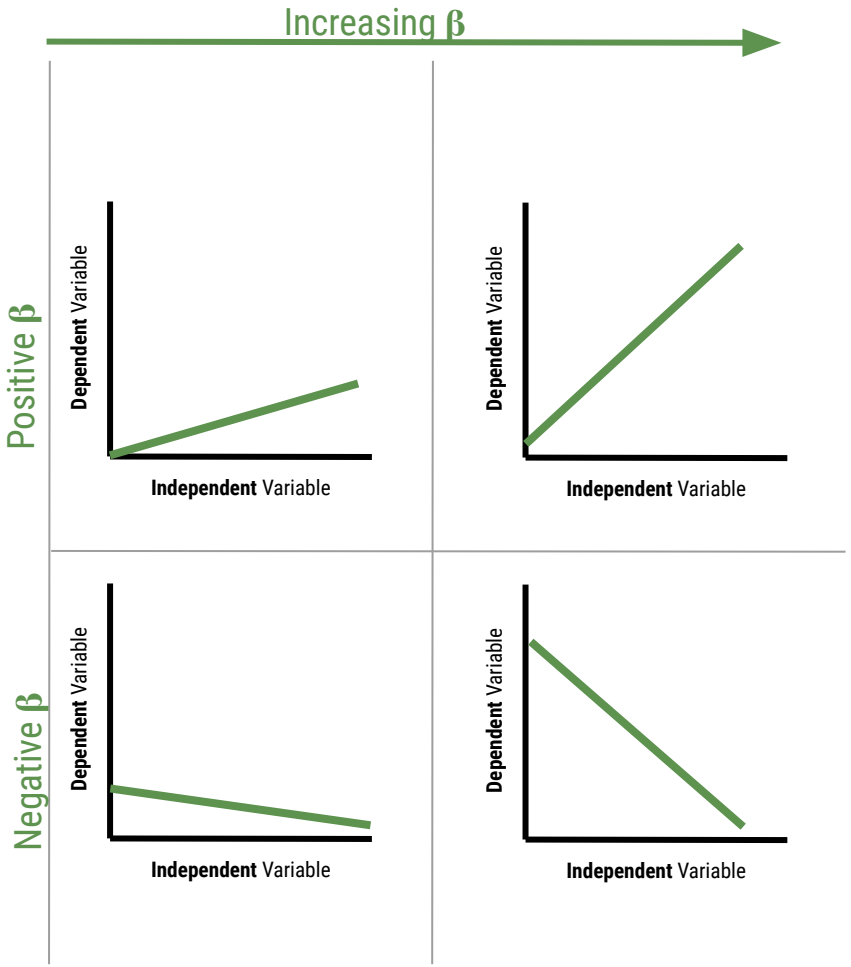

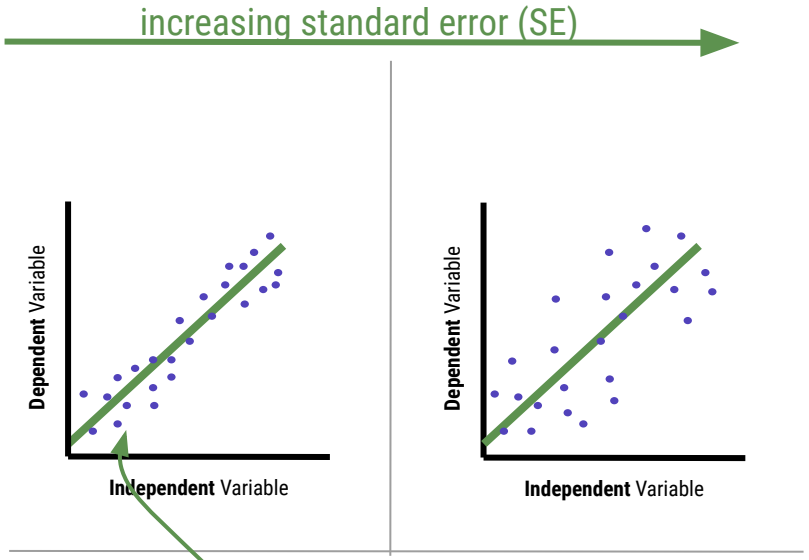

The *closer* the points are to the regression line, the *less uncertain* we are in our estimate

Assumptions of linear regression

1. Linear relationship
2. Multivariate normality
3. No multicollinearity
4. No auto-correlation
5. Homoscedasticity

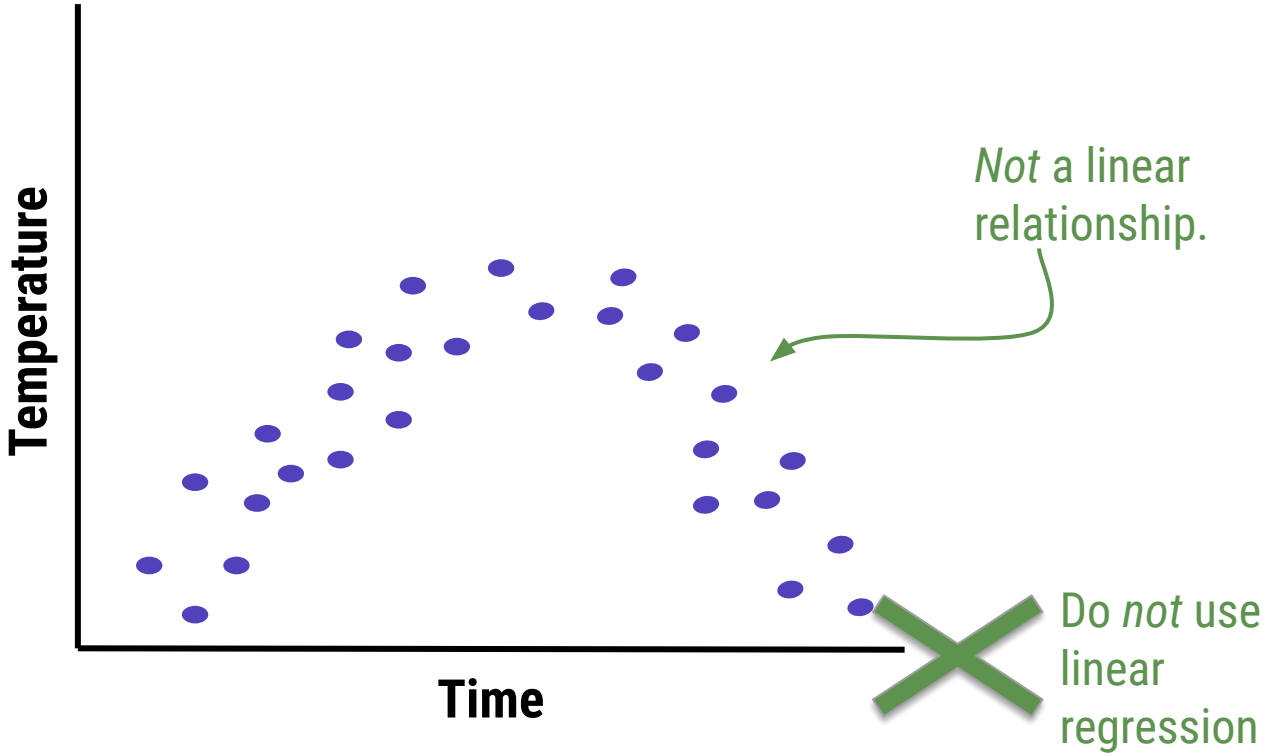
Temperature
Time
Not a linear relationship.
Do *not* use linear regression

A multivariate normal probability distribution (joint normal)

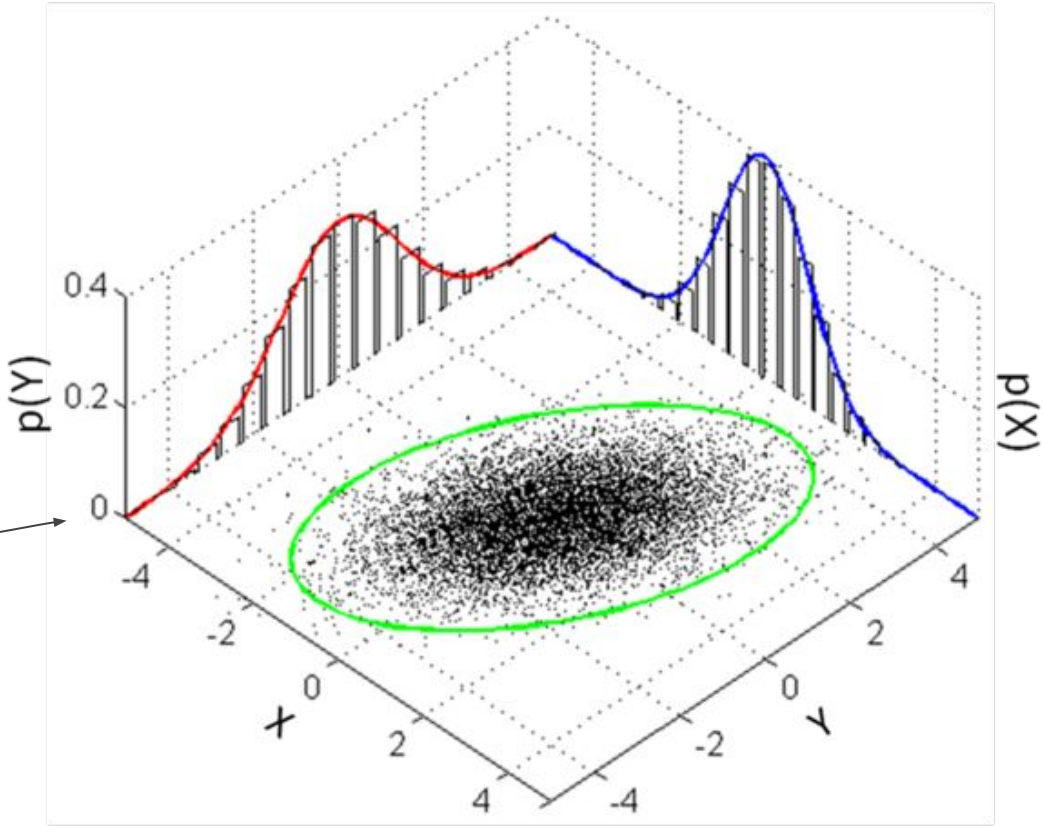

Linear regression assumes no multicollinearity. **Multicollinearity** occurs when the independent variables (in multiple linear regression) are too highly correlated with each other.

Autocorrelation occurs when the observations are *not* independent of one another (i.e. stock prices)

https://www.statisticssolutions.com/assumptions-of-linear-regression/

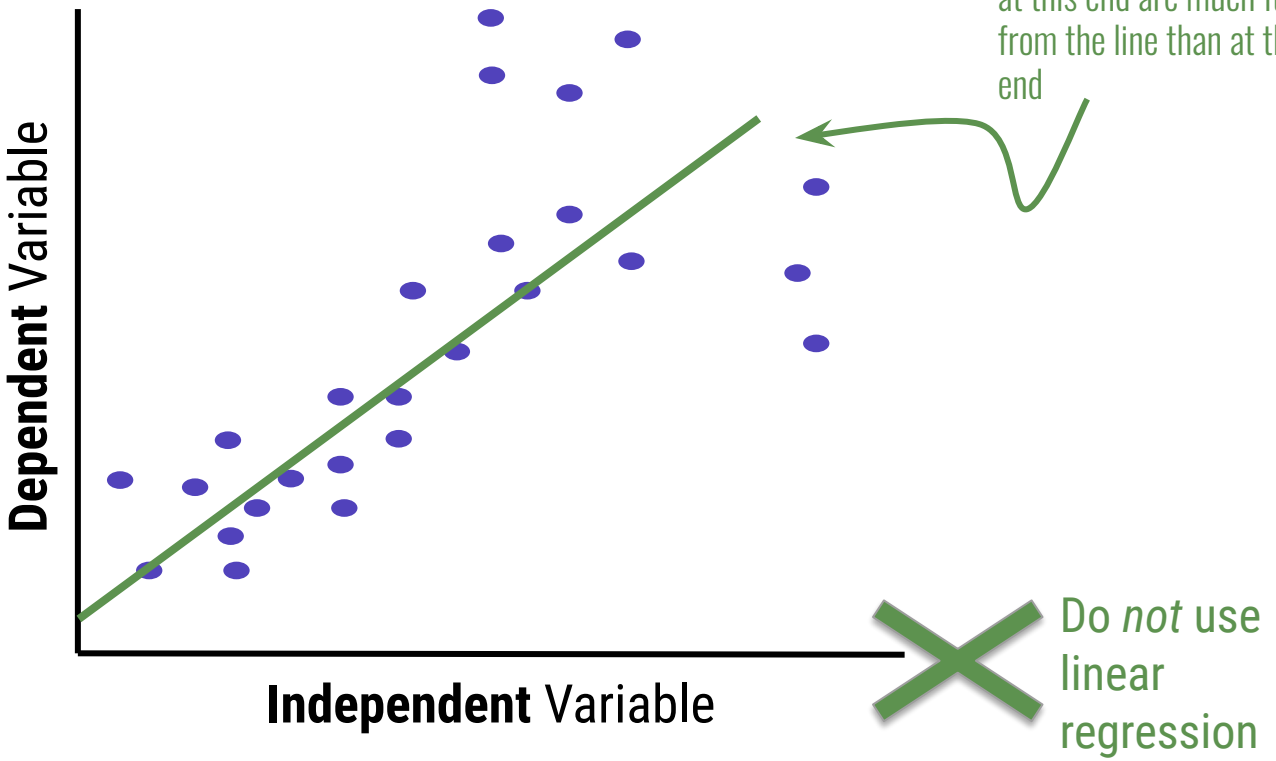
Not homoscedastic: points at this end are much further from the line than at the other end
Dependent Variable
Independent Variable
Do *not* use linear regression

Does Poverty Percentage affect Teen Birth Rate?

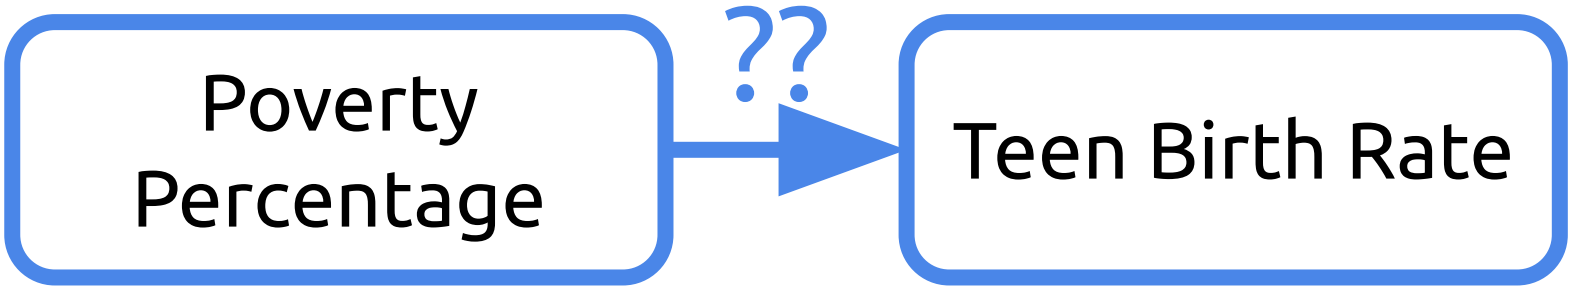

Null Hypothesis:

H_0: Poverty Rate does not affect Teen Birth Rate ($\beta=0$)

Alternative Hypothesis:

H_a: Poverty Rate affects Teen Birth Rate ($\beta\neq0$)

What is the relationship between Poverty Percentage & Teen Birth Rate?

What's your hypothesis?

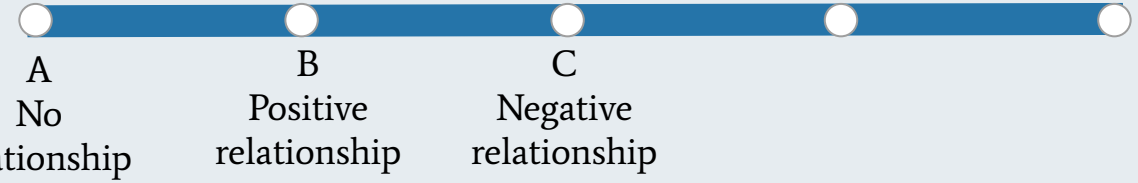

	Location	PovPct	Brth15to17	Brth18to19	ViolCrime	TeenBrth
1	Alabama	20.1	31.5	88.7	11.2	54.5
2	Alaska	7.1	18.9	73.7	9.1	39.5
3	Arizona	16.1	35.0	102.5	10.4	61.2
4	Arkansas	14.9	31.6	101.7	10.4	59.9
5	California	16.7	22.6	69.1	11.2	41.1
6	Colorado	8.8	26.2	79.1	5.8	47.0
7	Connecticut	9.7	14.1	45.1	4.6	25.8
8	Delaware	10.3	24.7	77.8	3.5	46.3
9	District_of_Columbia	22.0	44.8	101.5	65.0	69.1
10	Florida	16.2	23.2	78.4	7.3	44.5
11	Georgia	12.1	31.4	92.8	9.5	55.7
12	Hawaii	10.3	17.7	66.4	4.7	38.2
13	Idaho	14.5	18.4	69.1	4.1	39.1
14	Illinois	12.4	23.4	70.5	10.3	42.2
15	Indiana	9.6	22.6	78.5	8.0	44.6
16	Iowa	12.2	16.4	55.4	1.8	32.5
17	Kansas	10.8	21.4	74.2	6.2	43.0

Normal(ish) distributions

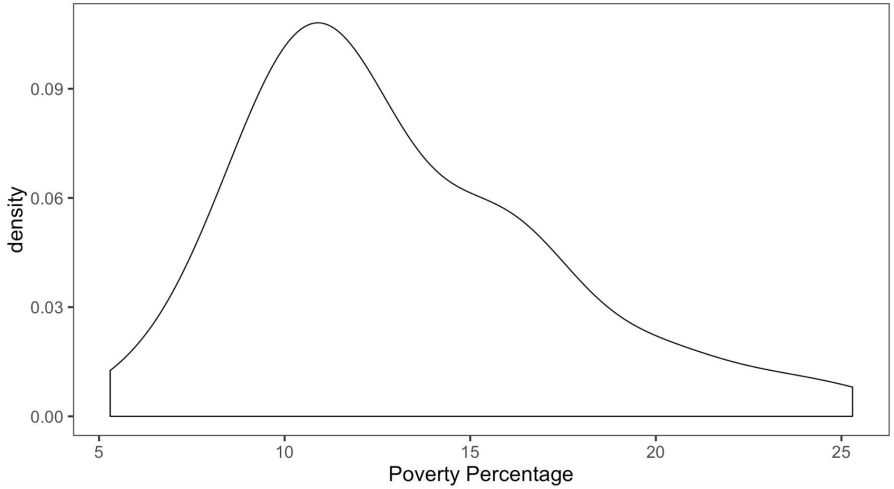

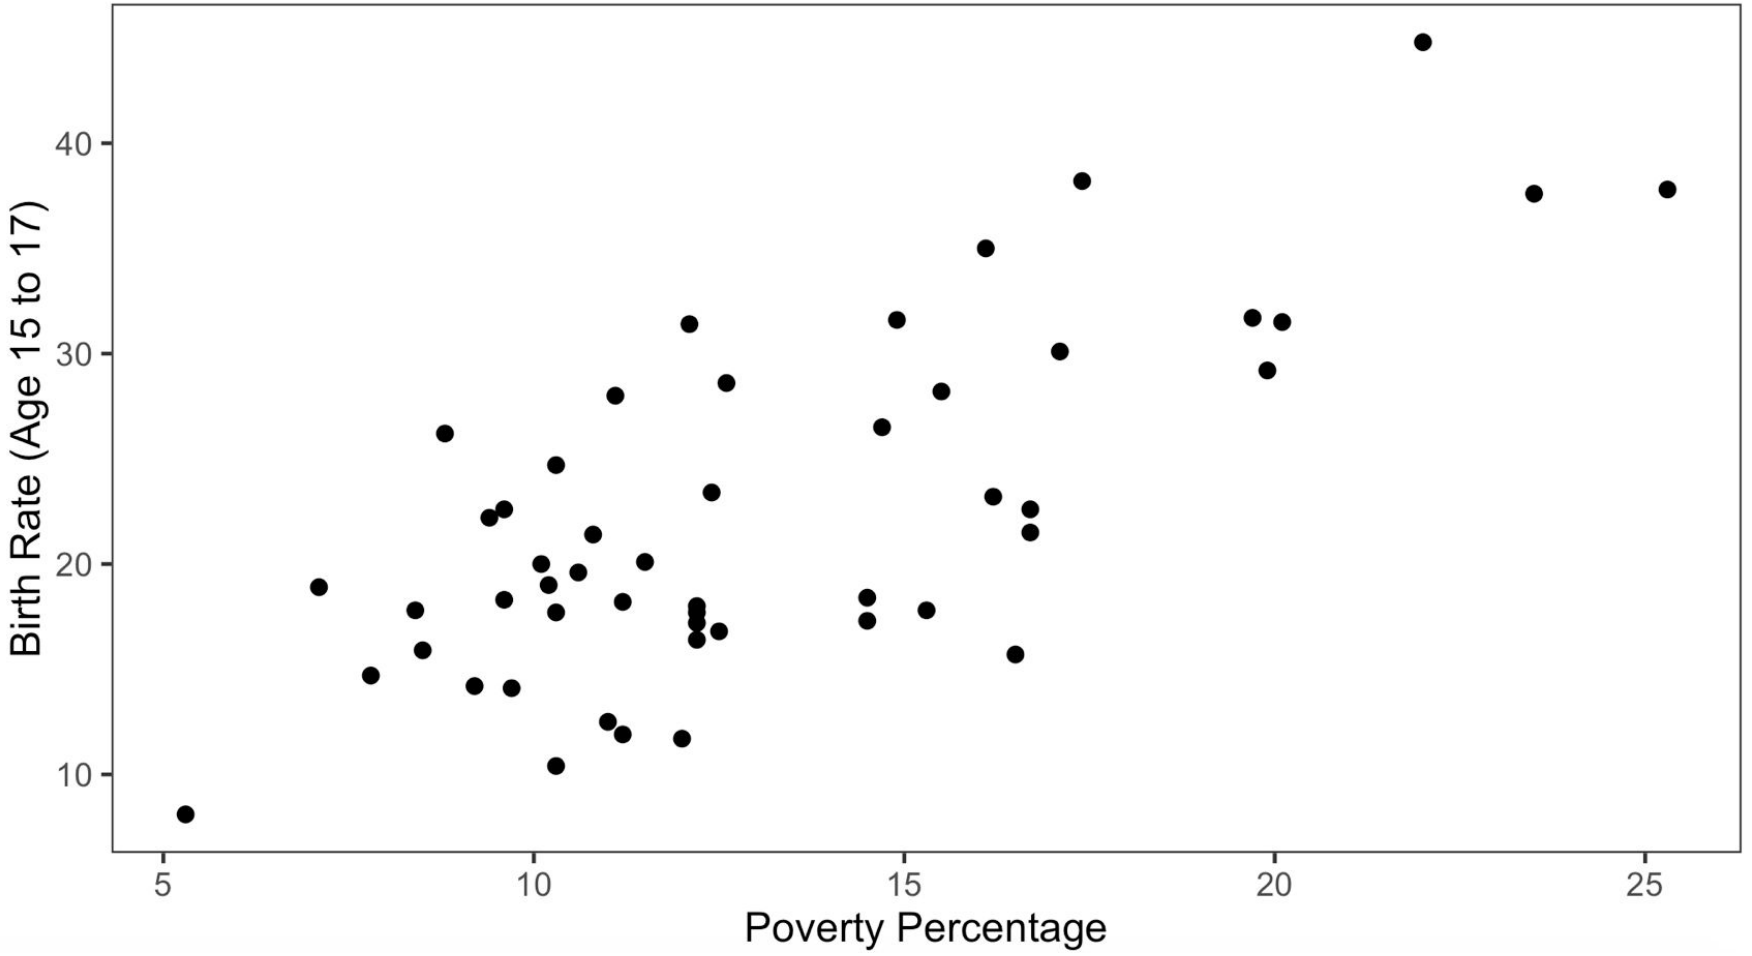

Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard.

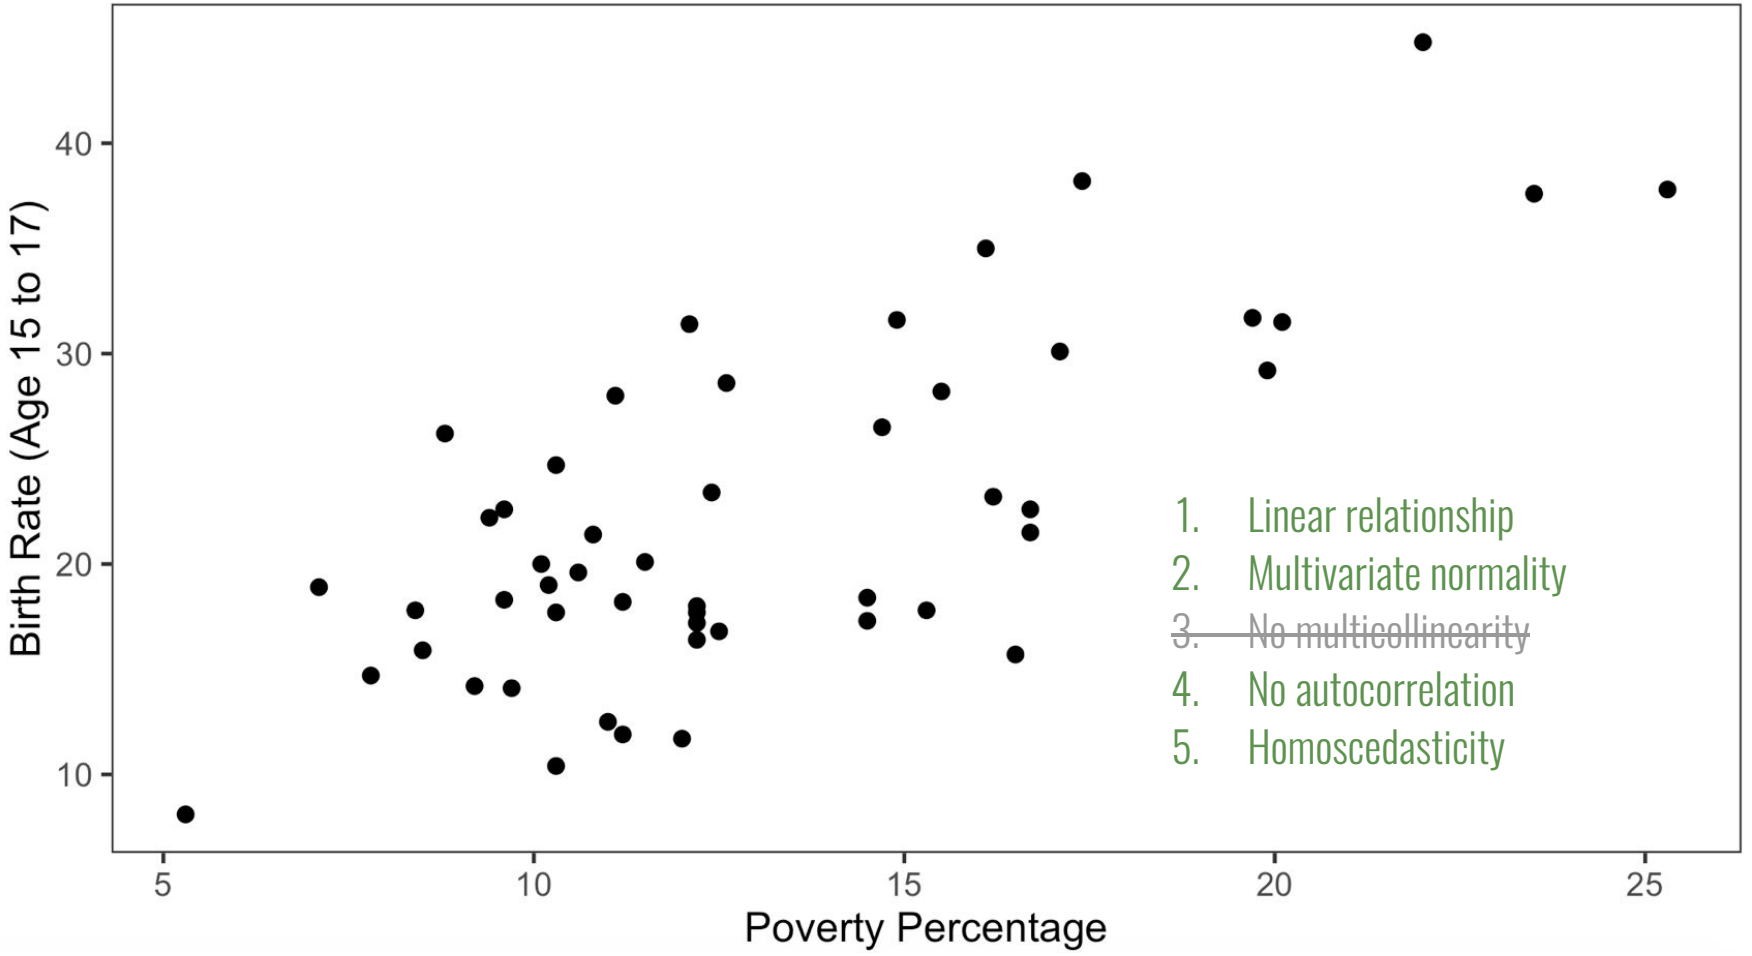

Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard.

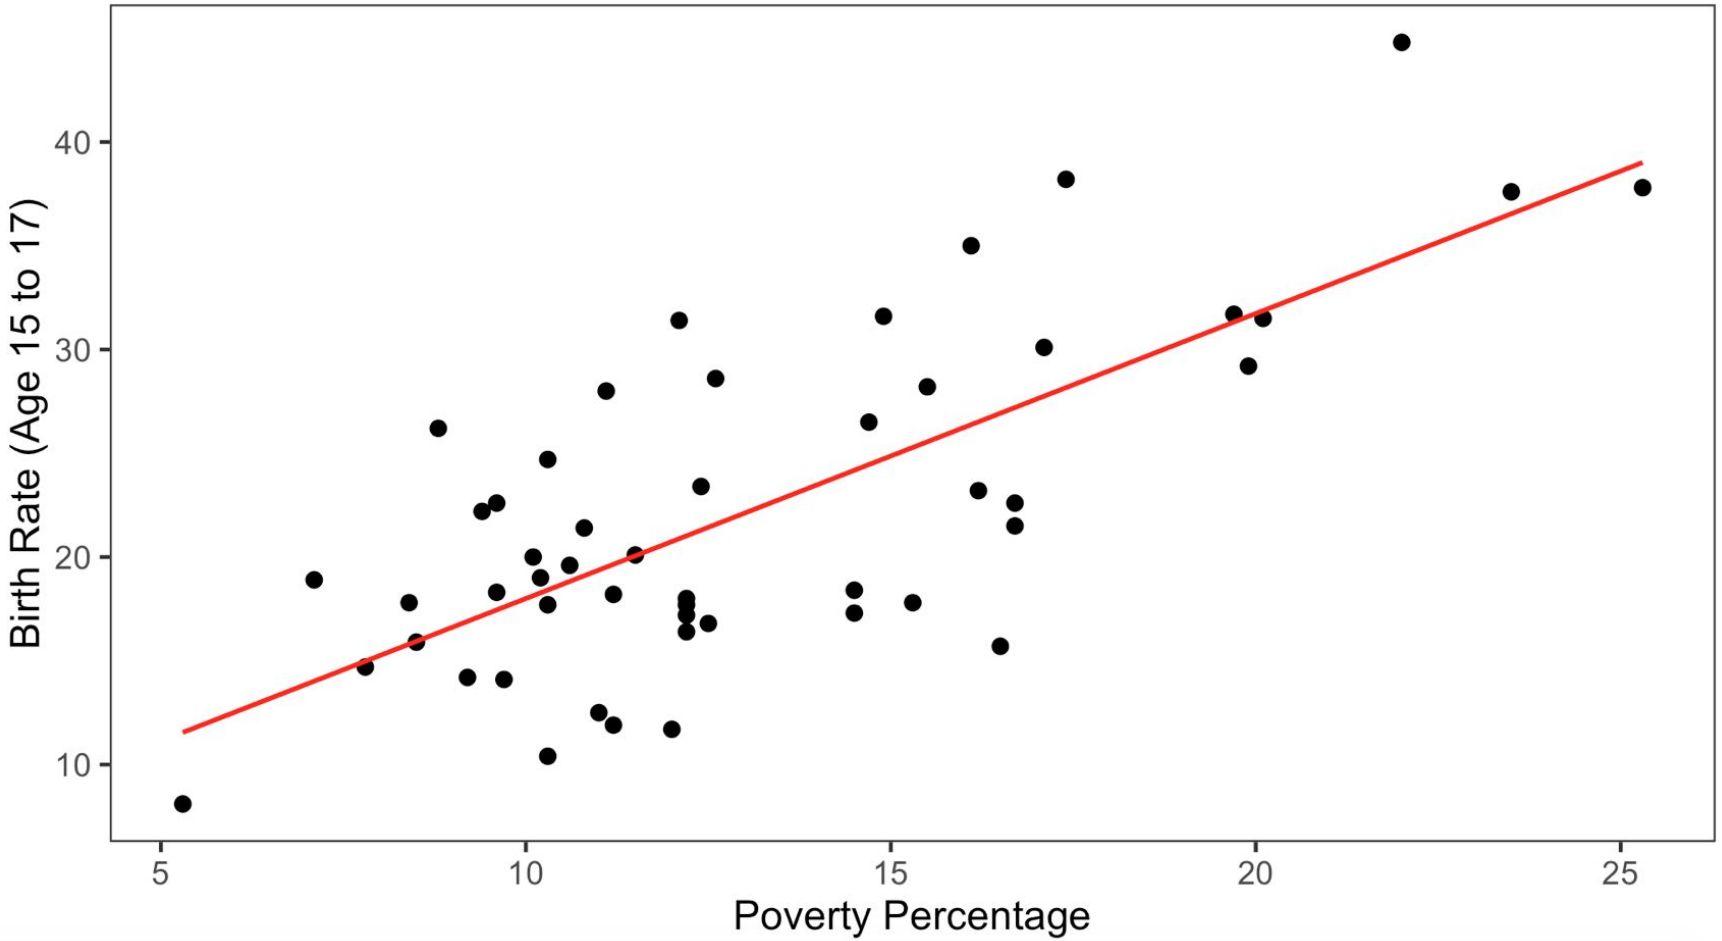
Birth Rate (Age 15 to 17)
40
30
20
10
5
10
15
20
25
Poverty Percentage

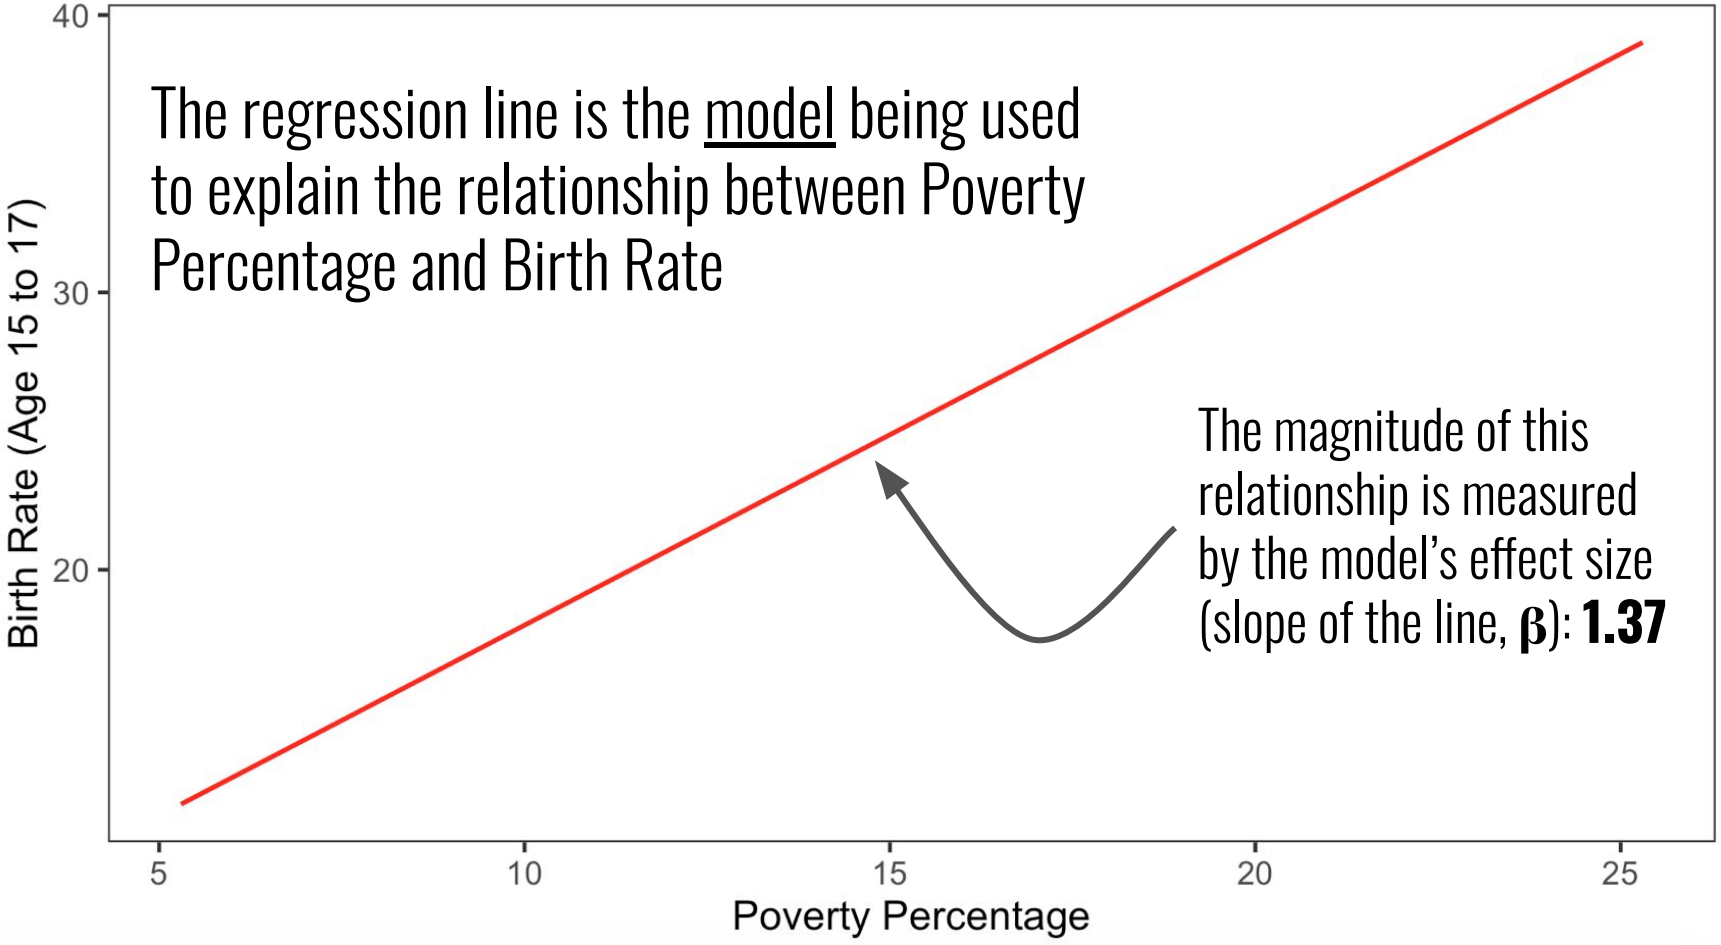
The regression line is the model being used to explain the relationship between Poverty Percentage and Birth Rate
The magnitude of this relationship is measured by the model's effect size (slope of the line, β): 1.37
Birth Rate (Age 15 to 17)
40
30
20
5
10
15
20
25
Poverty Percentage

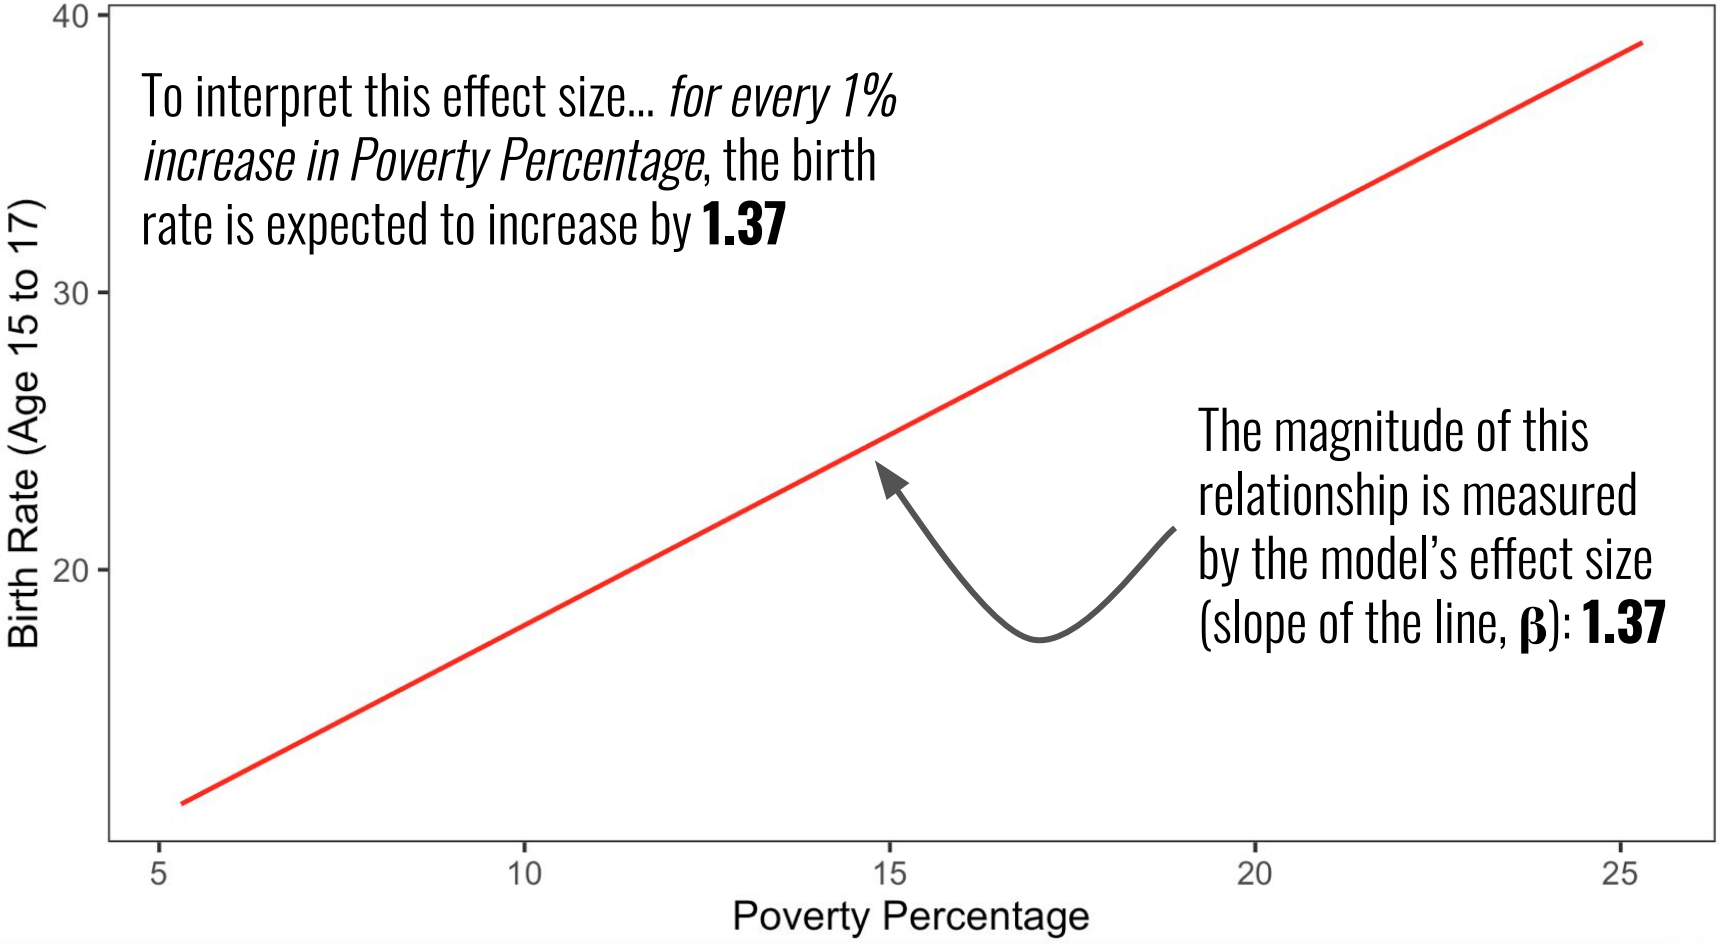
To interpret this effect size... *for every 1% increase in Poverty Percentage*, the birth rate is expected to increase by **1.37**
The magnitude of this relationship is measured by the model's effect size (slope of the line, β): **1.37**
Birth Rate (Age 15 to 17)
40
30
20
5
10
15
20
25
Poverty Percentage

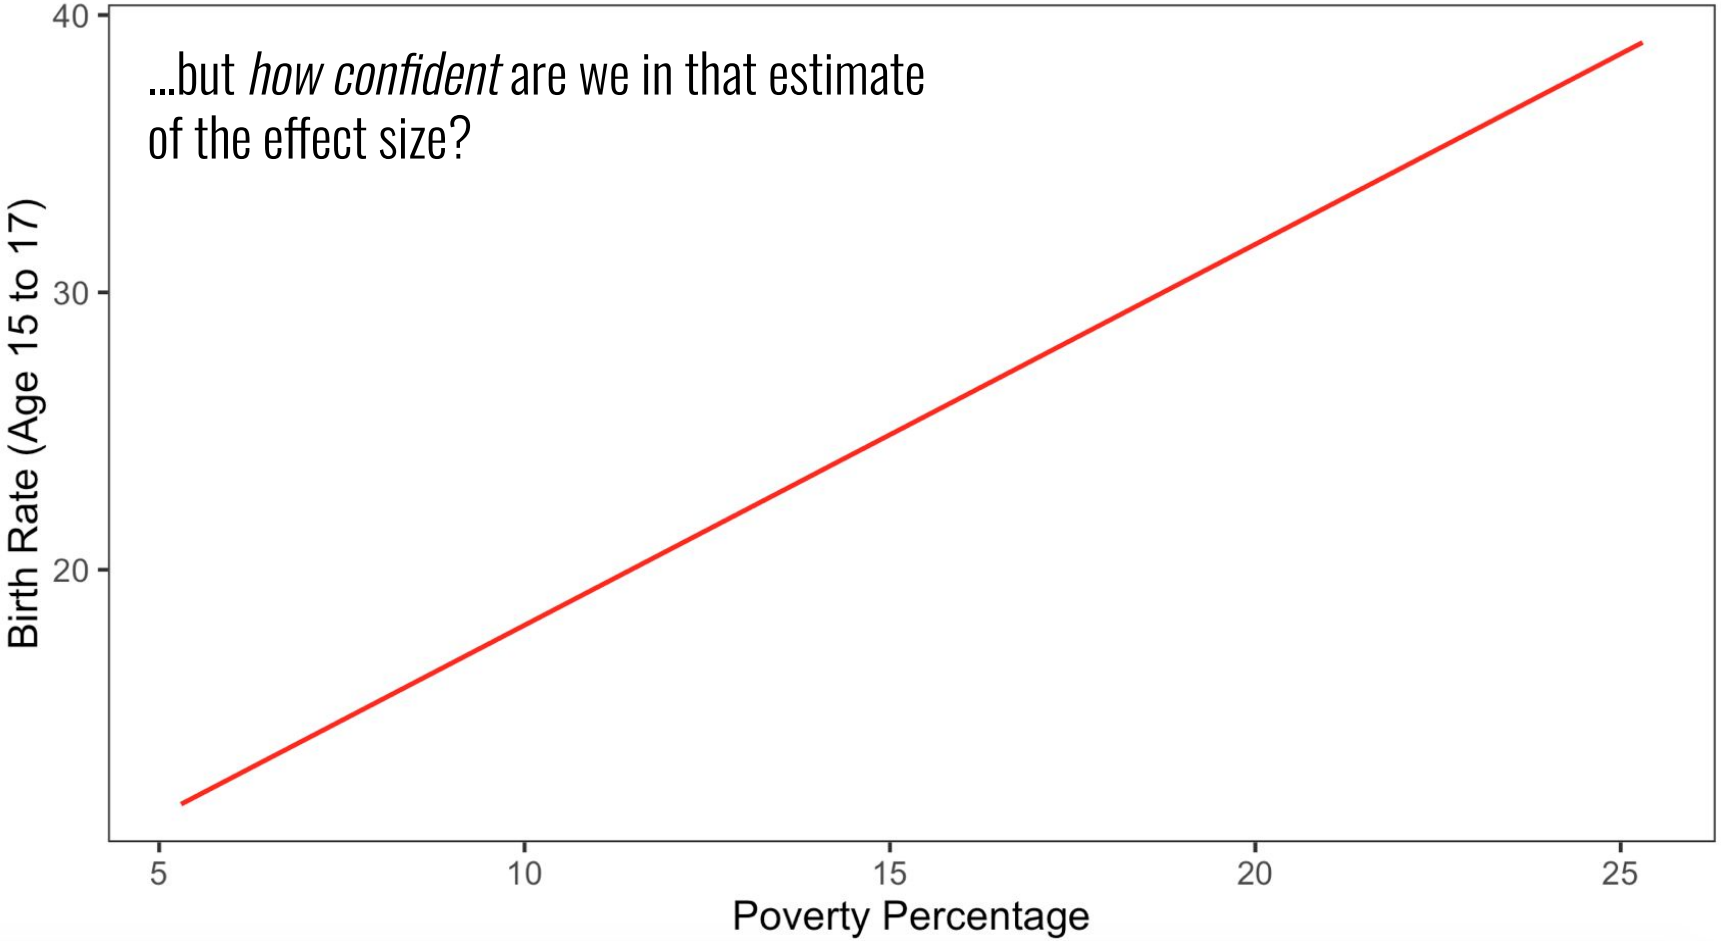
...but *how confident* are we in that estimate of the effect size?
Birth Rate (Age 15 to 17)
40
30
20
5
10
15
20
25
Poverty Percentage

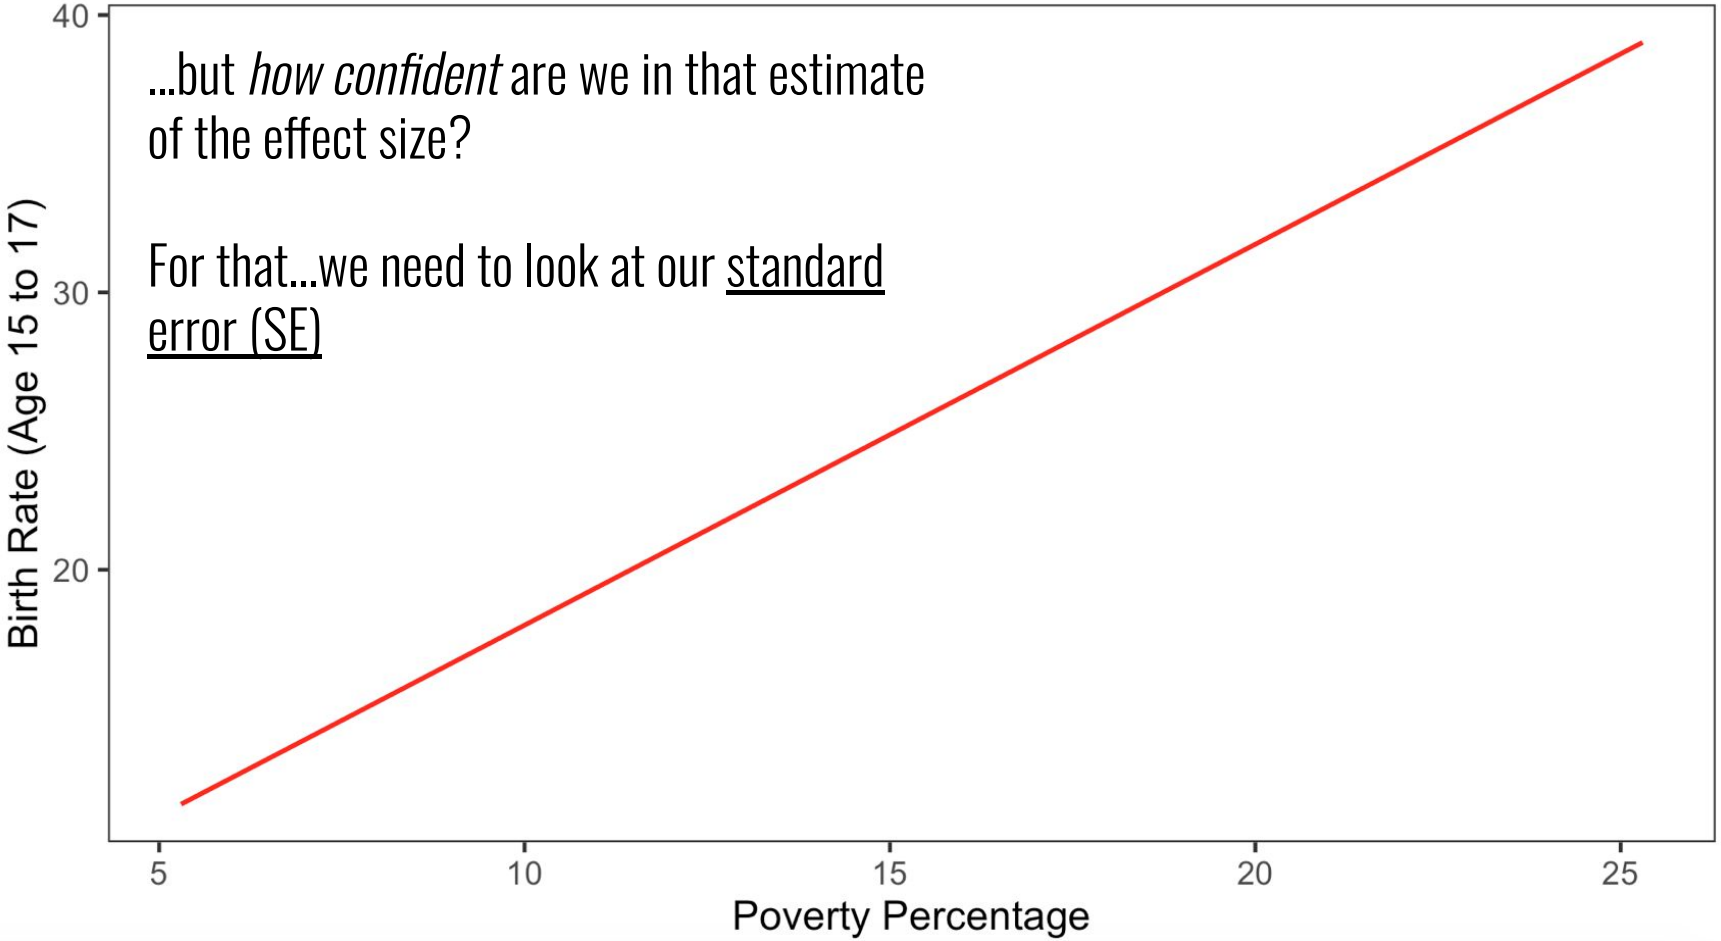
...but *how confident* are we in that estimate of the effect size?
For that...we need to look at our standard error (SE)
Birth Rate (Age 15 to 17)
40
30
20
5
10
15
20
25
Poverty Percentage

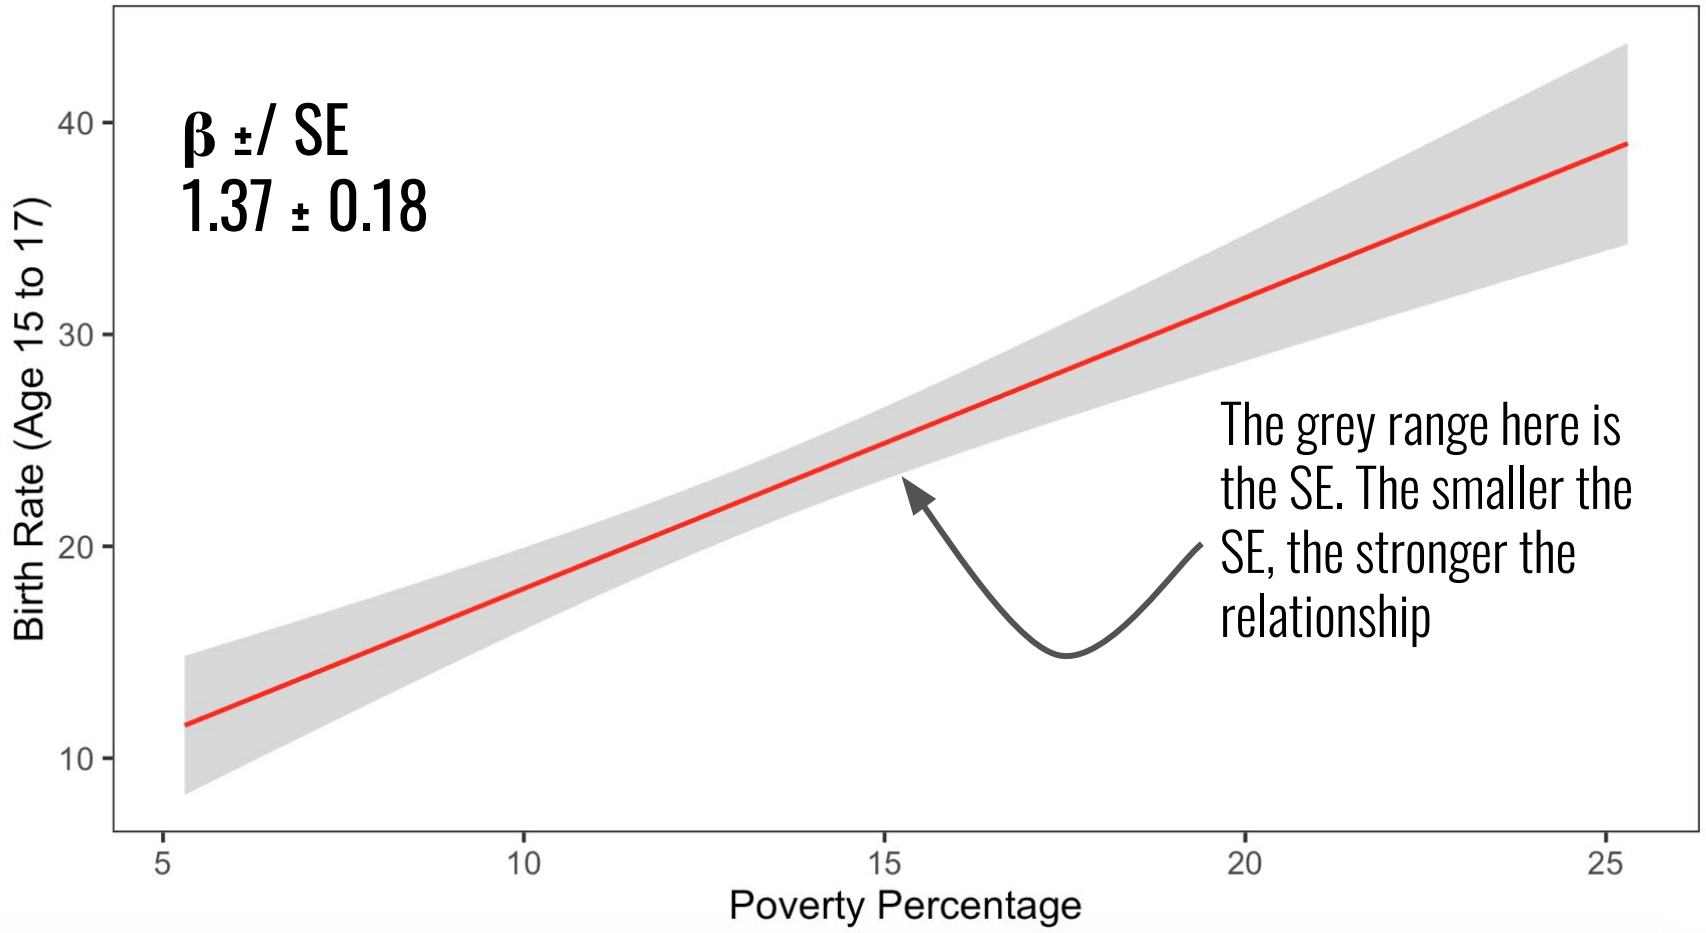
β ±/ SE
1.37 ± 0.18
Birth Rate (Age 15 to 17)
40
30
20
10
5
10
15
20
25
Poverty Percentage
The grey range here is the SE. The smaller the SE, the stronger the relationship

Normal(ish) distributions

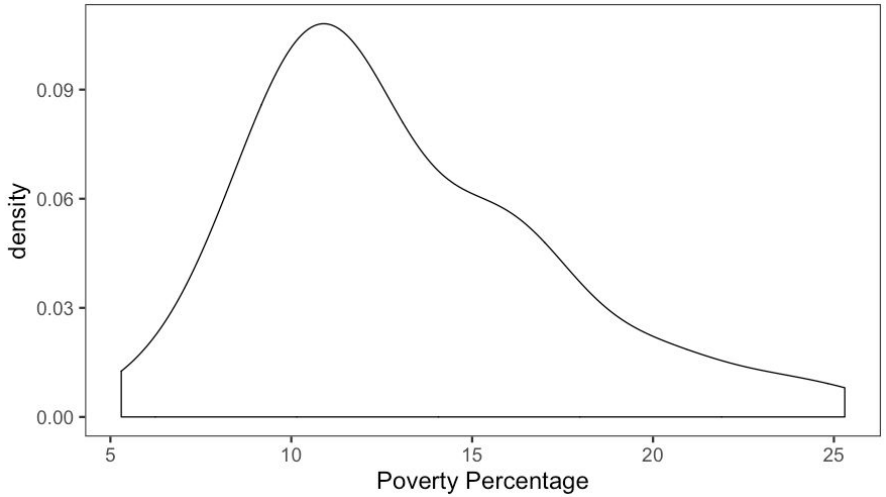

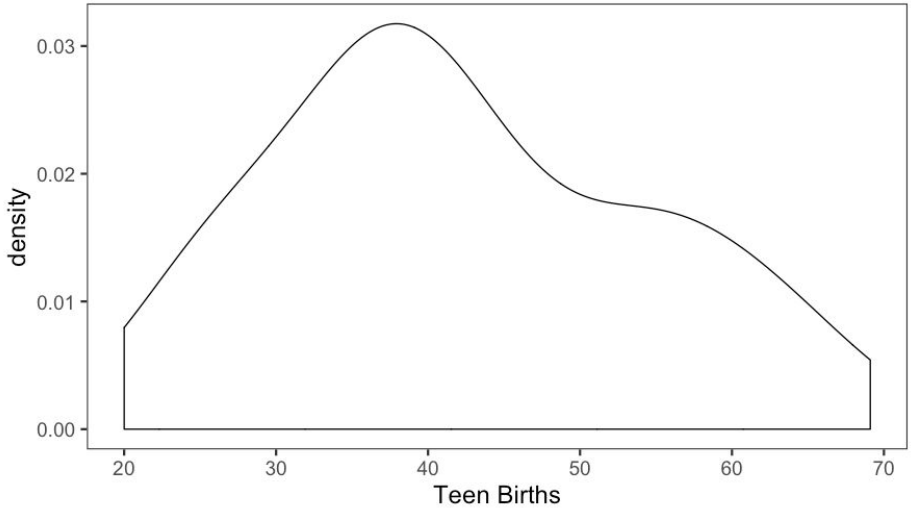

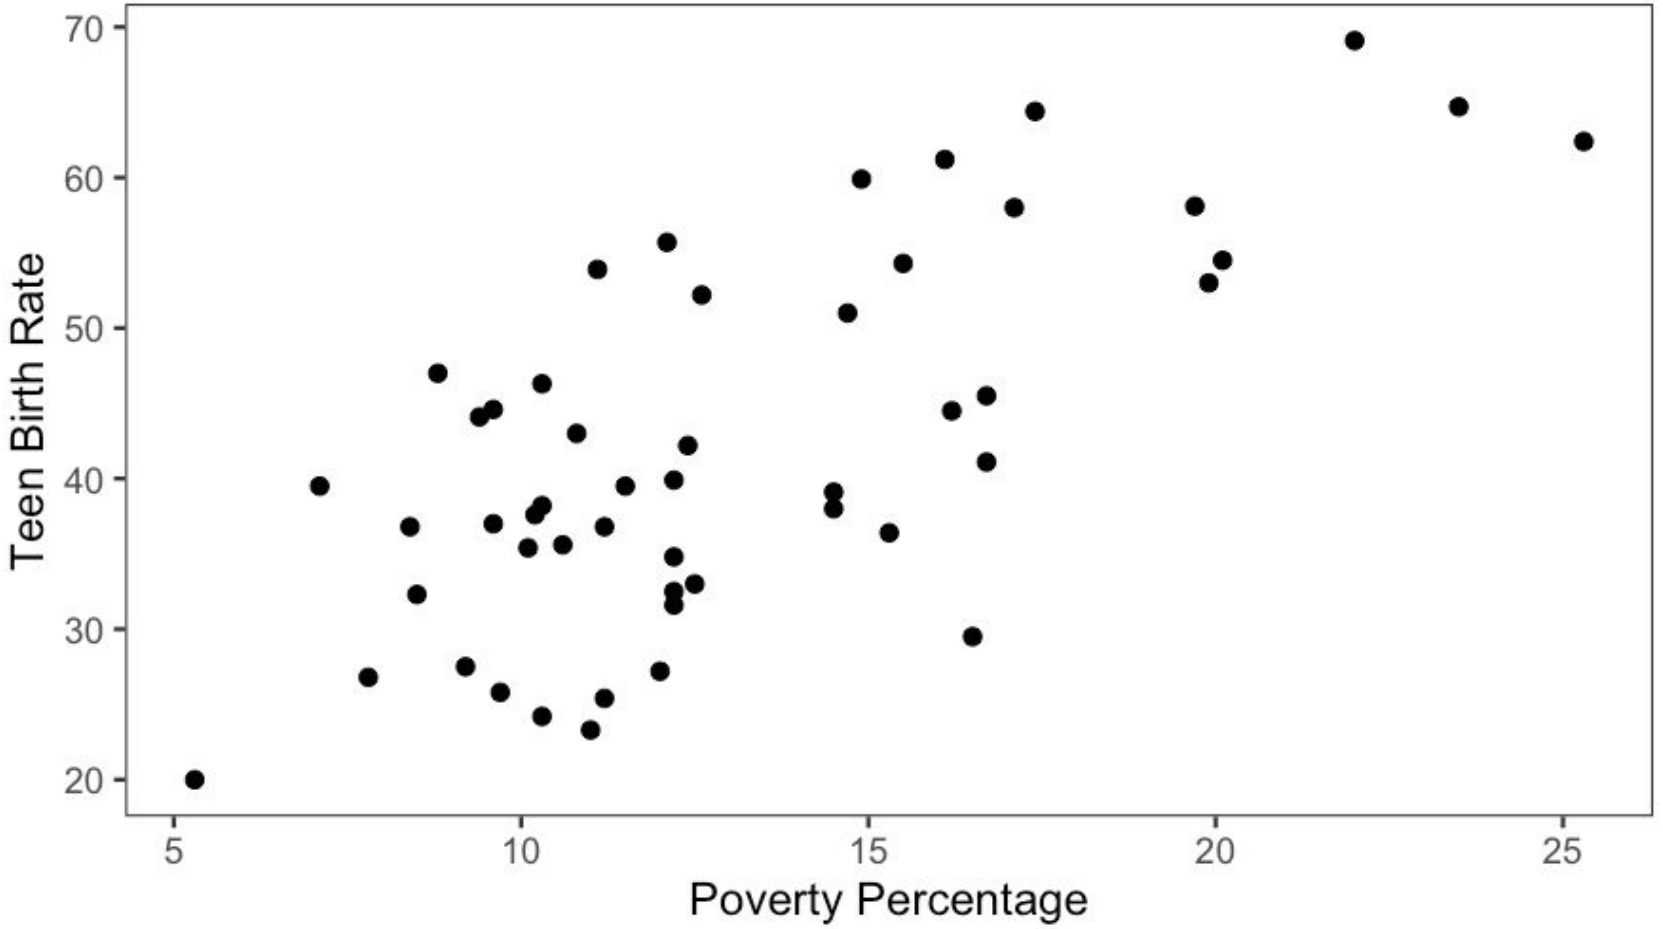

Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard.

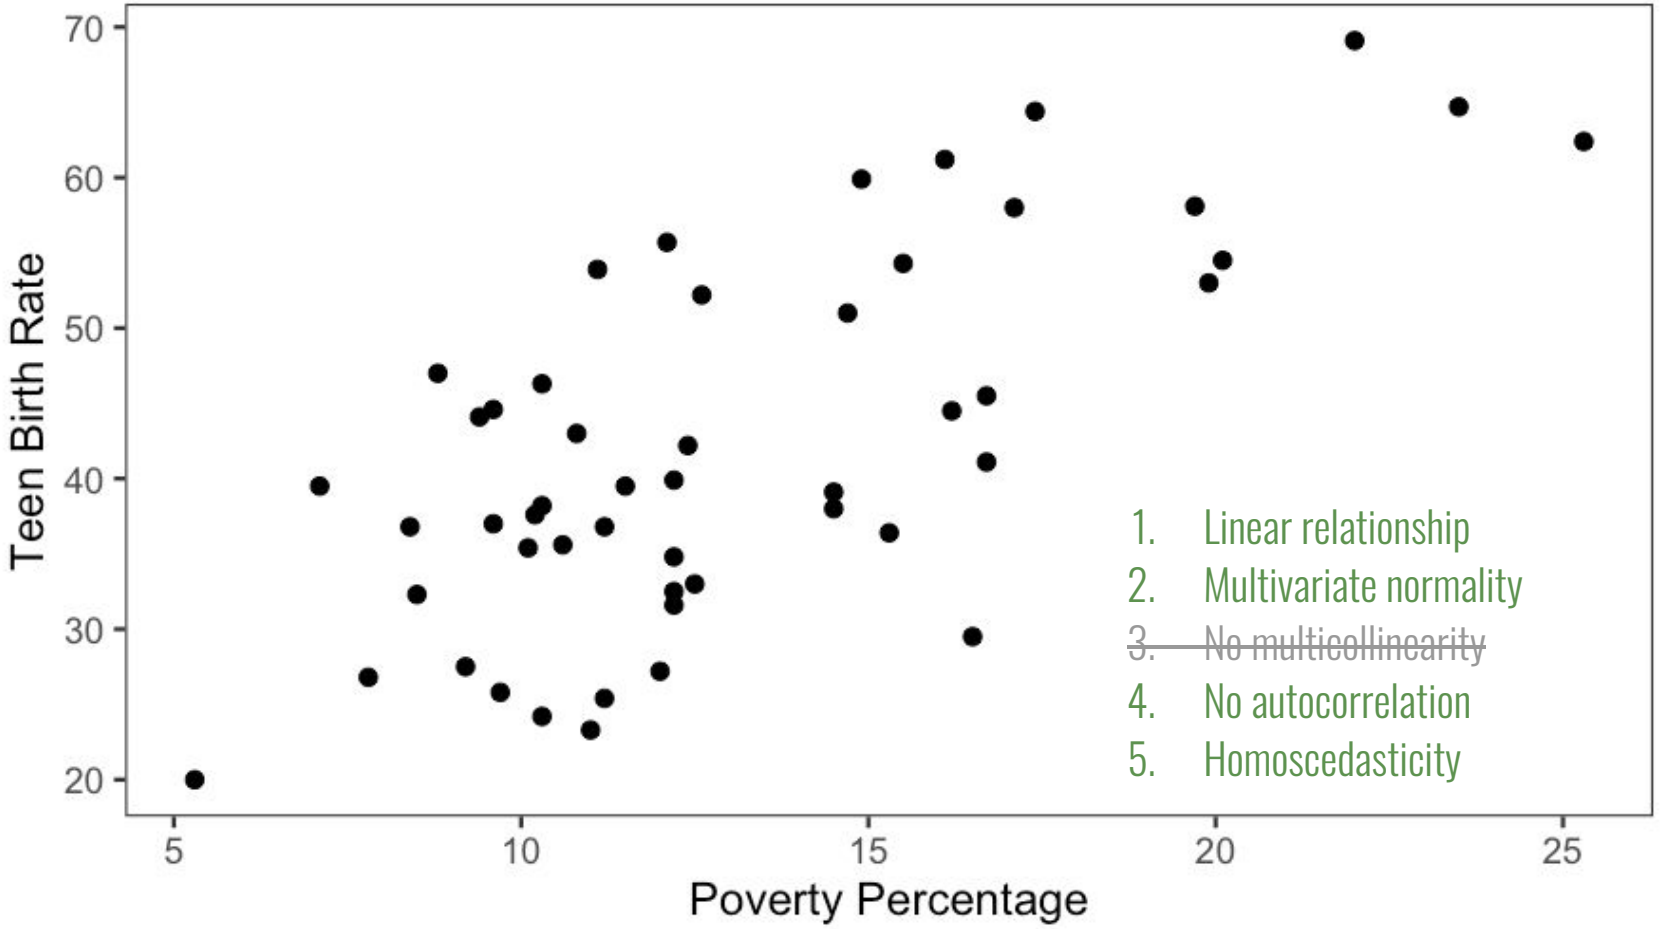

Data source: *Mind On Statistics*, 3rd edition, Utts and Heckard.

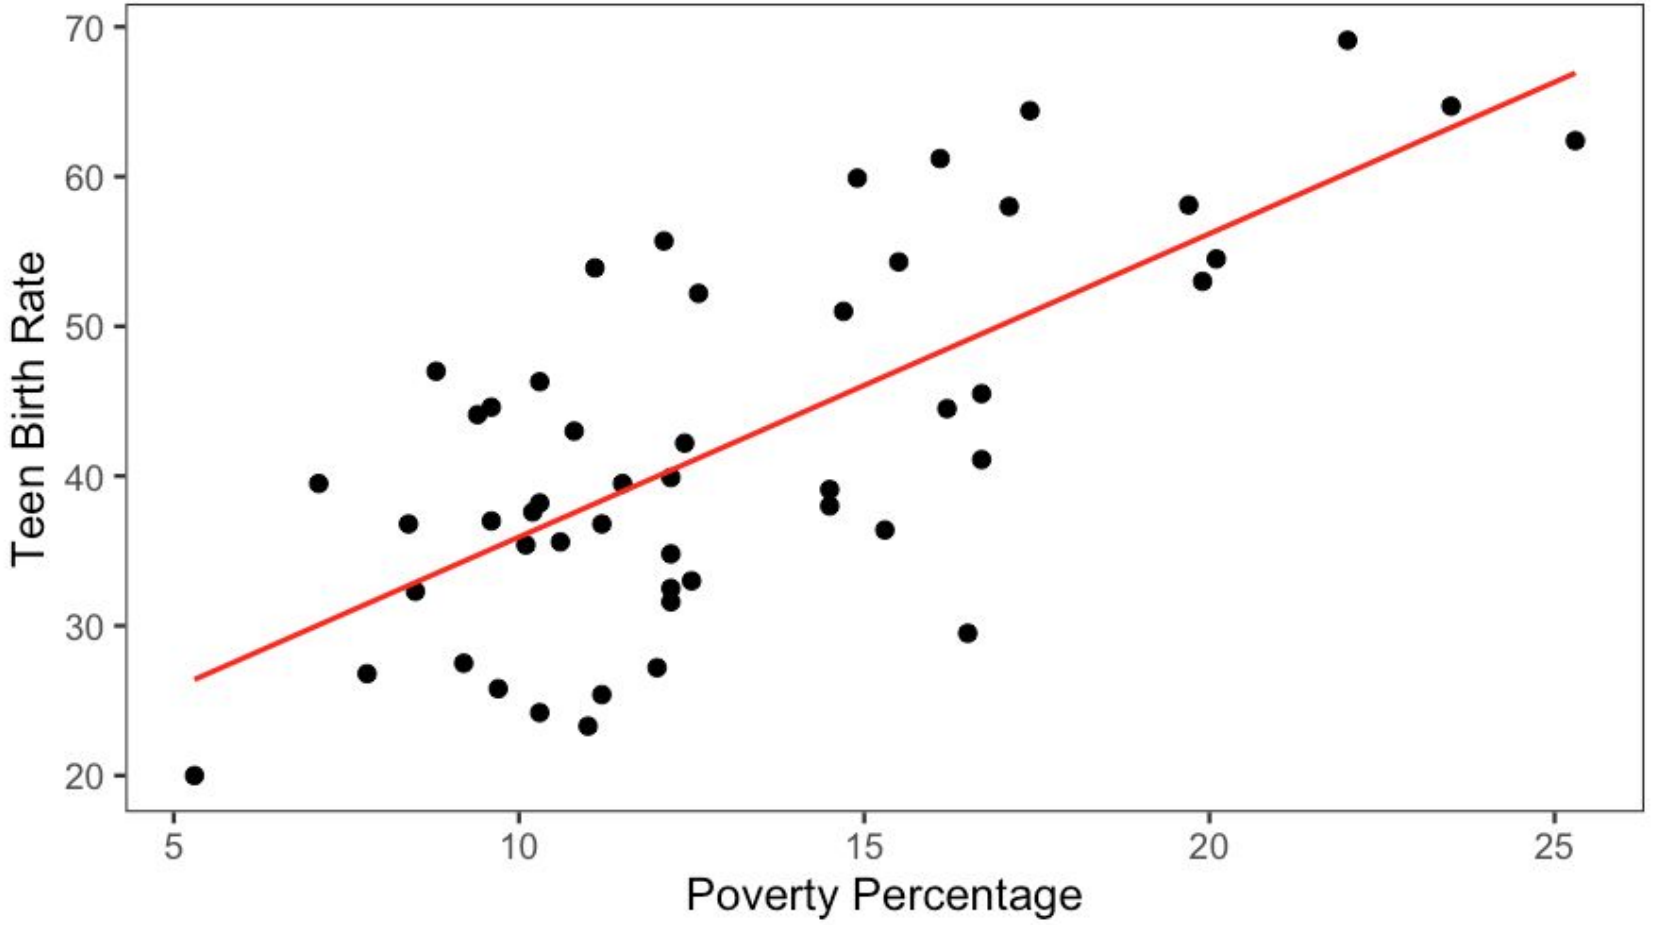
Teen Birth Rate
70
60
50
40
30
20
5
10
15
20
25
Poverty Percentage

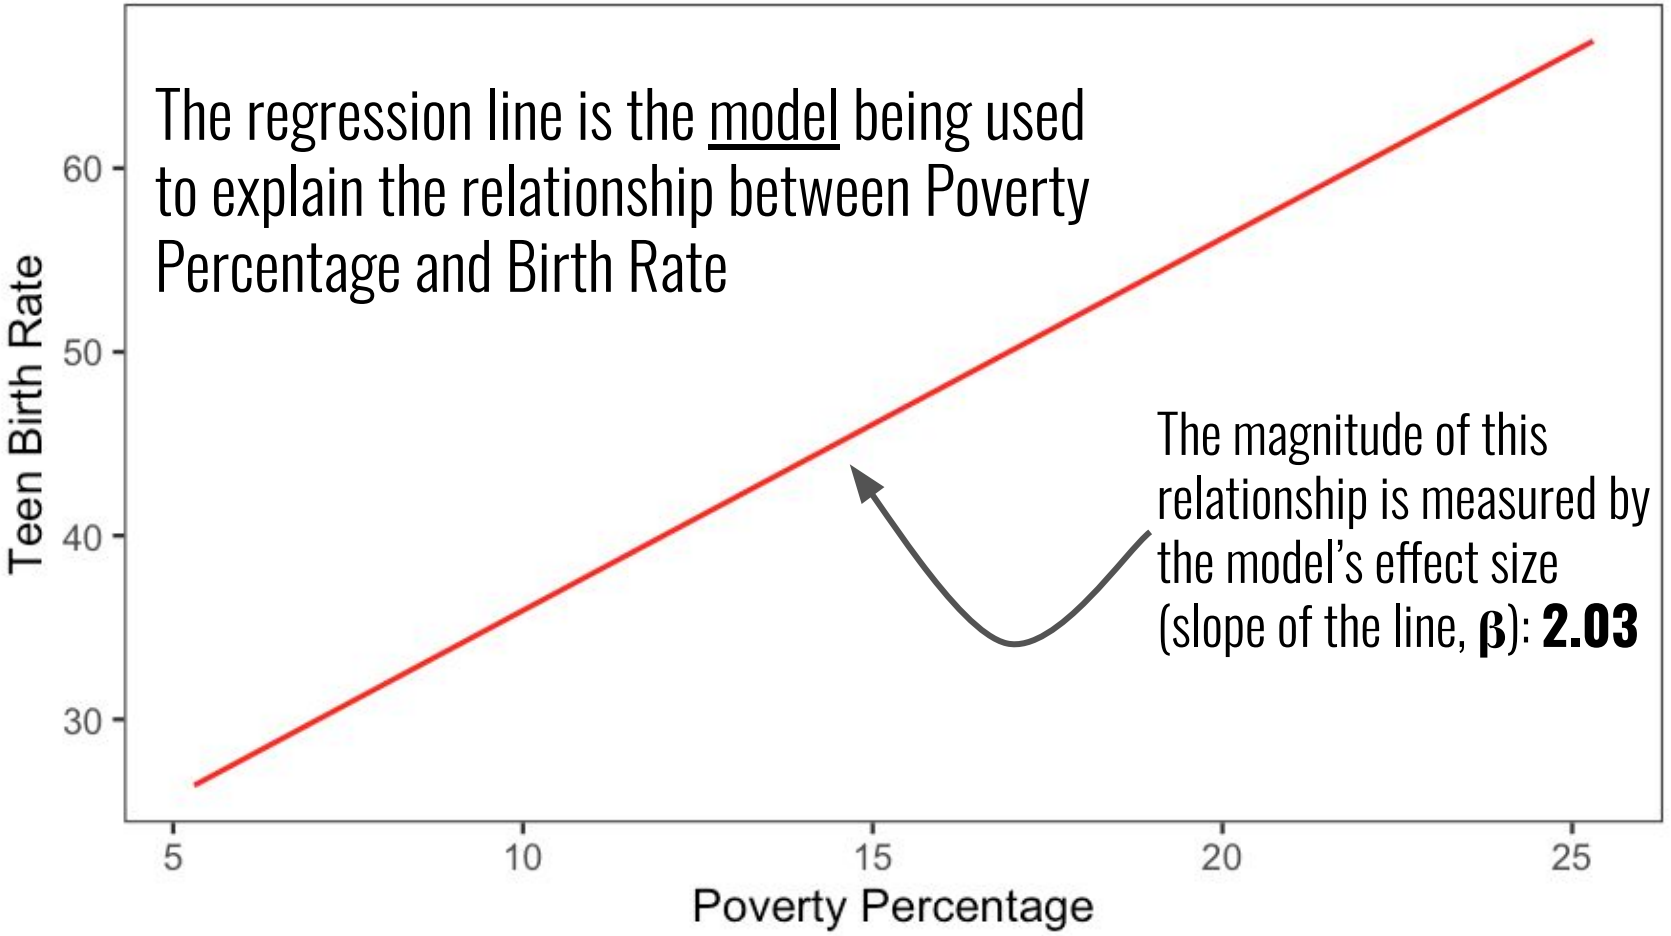
The regression line is the model being used to explain the relationship between Poverty Percentage and Birth Rate
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage
The magnitude of this relationship is measured by the model's effect size (slope of the line, β): 2.03

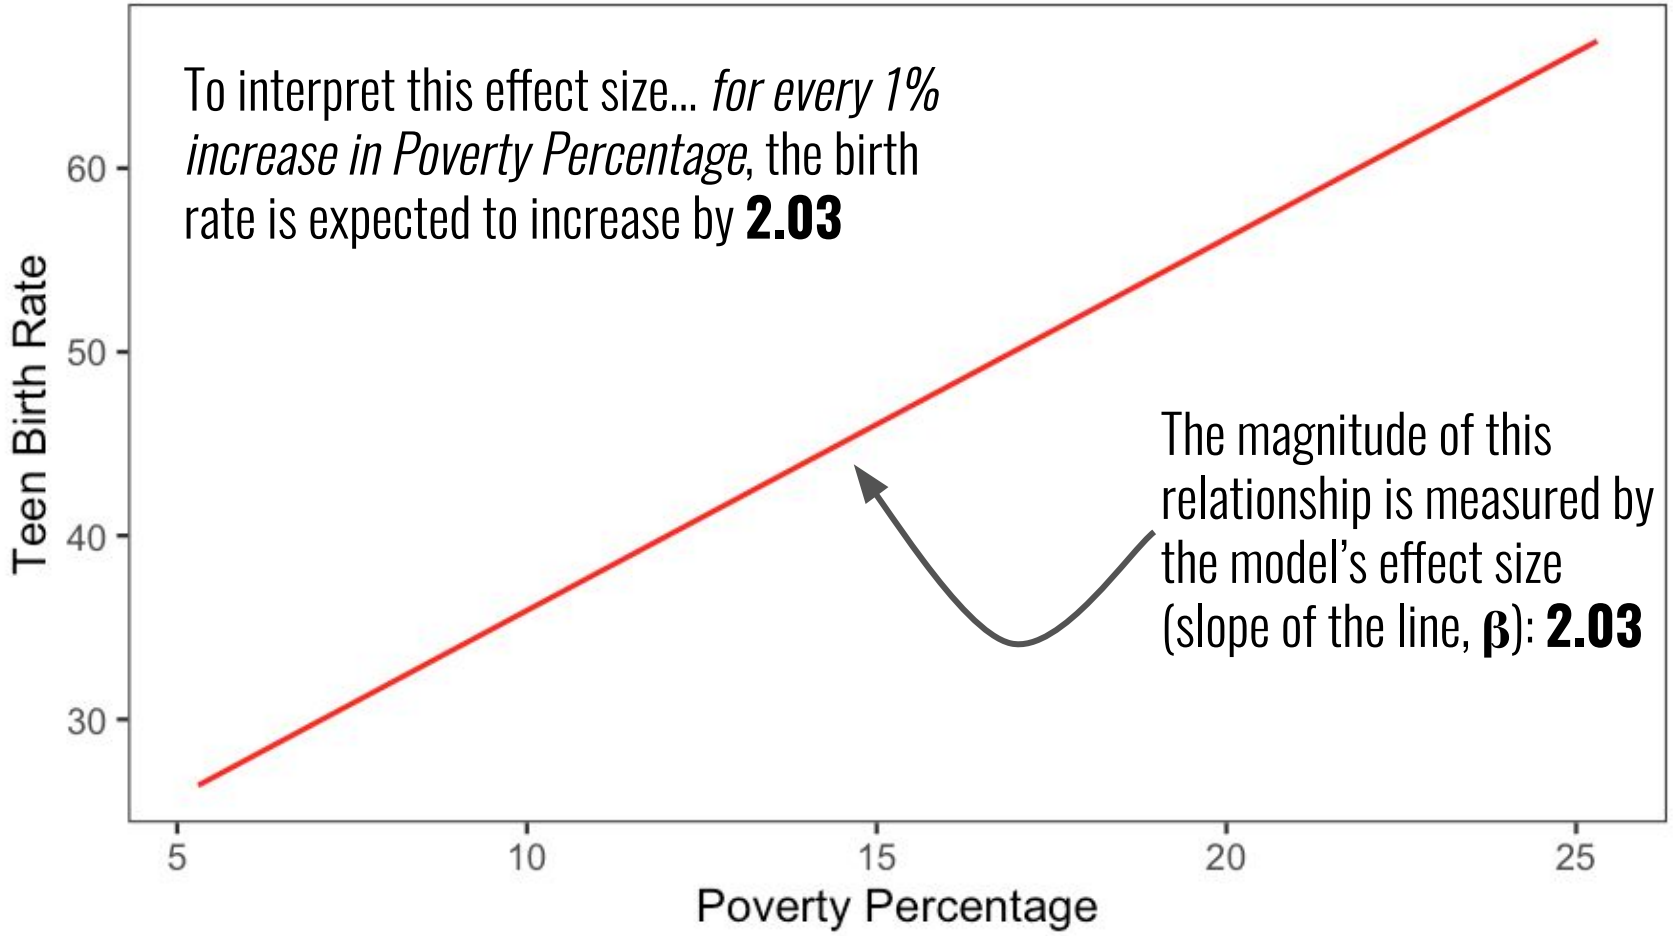
To interpret this effect size... *for every 1% increase in Poverty Percentage*, the birth rate is expected to increase by **2.03**
The magnitude of this relationship is measured by the model's effect size (slope of the line, β): **2.03**
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage

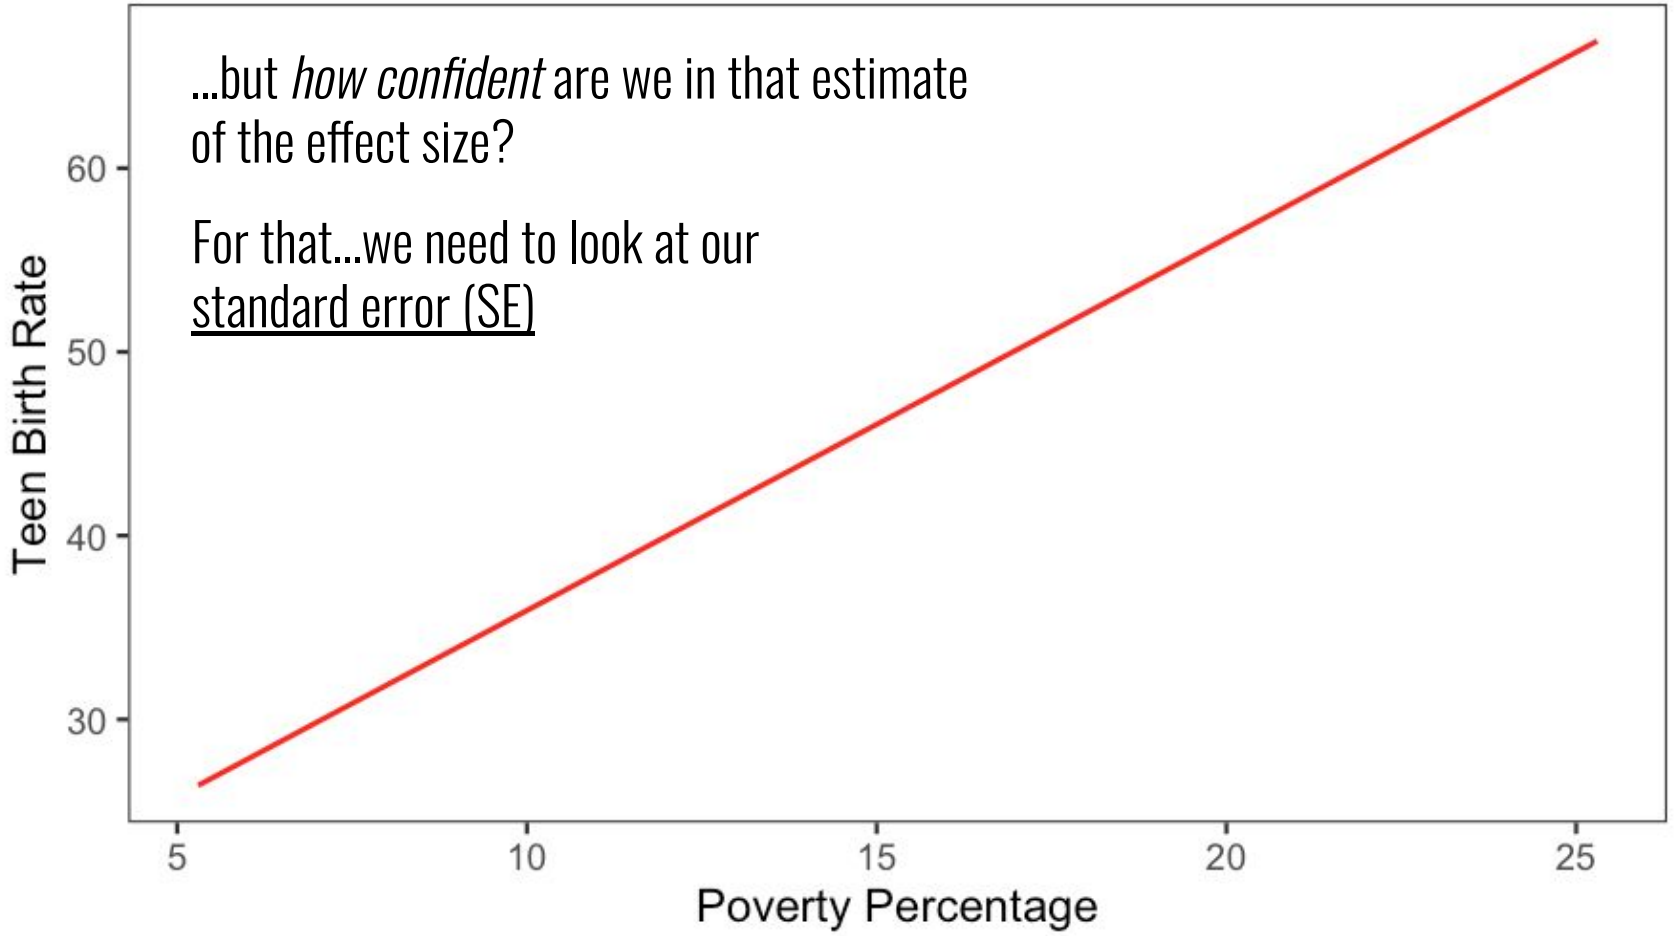
...but *how confident* are we in that estimate of the effect size?
For that...we need to look at our standard error (SE)
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage

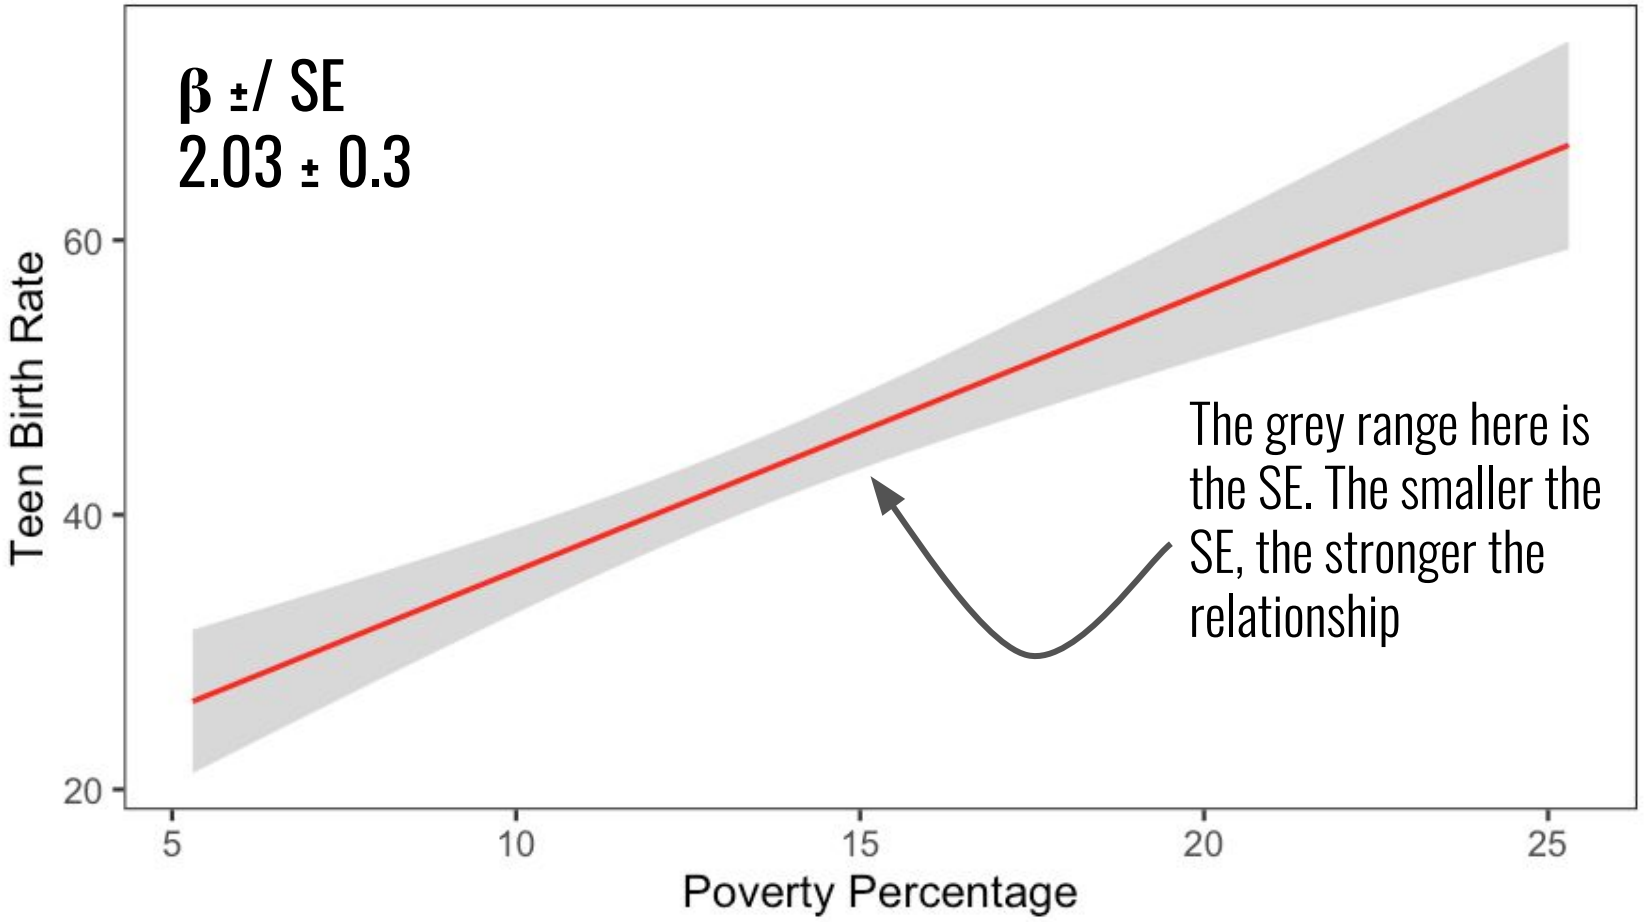
β ±/ SE
2.03 ± 0.3
Teen Birth Rate
60
40
20
5
10
15
20
25
Poverty Percentage
The grey range here is the SE. The smaller the SE, the stronger the relationship

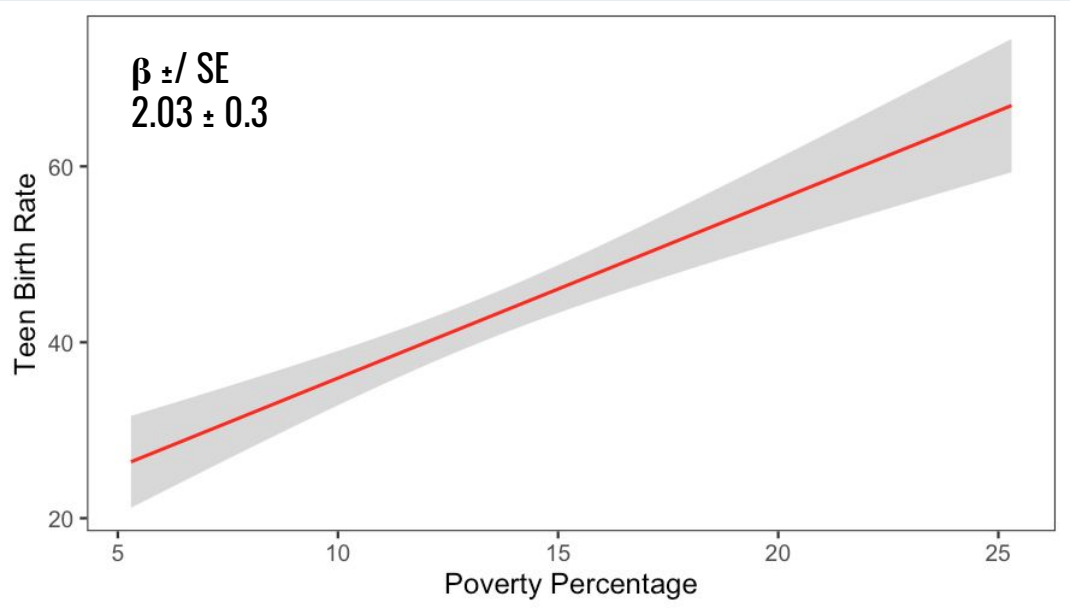

If there were a stronger effect of Poverty on Birth rate, what would β be?

A
< 2.03

B
> 2.03

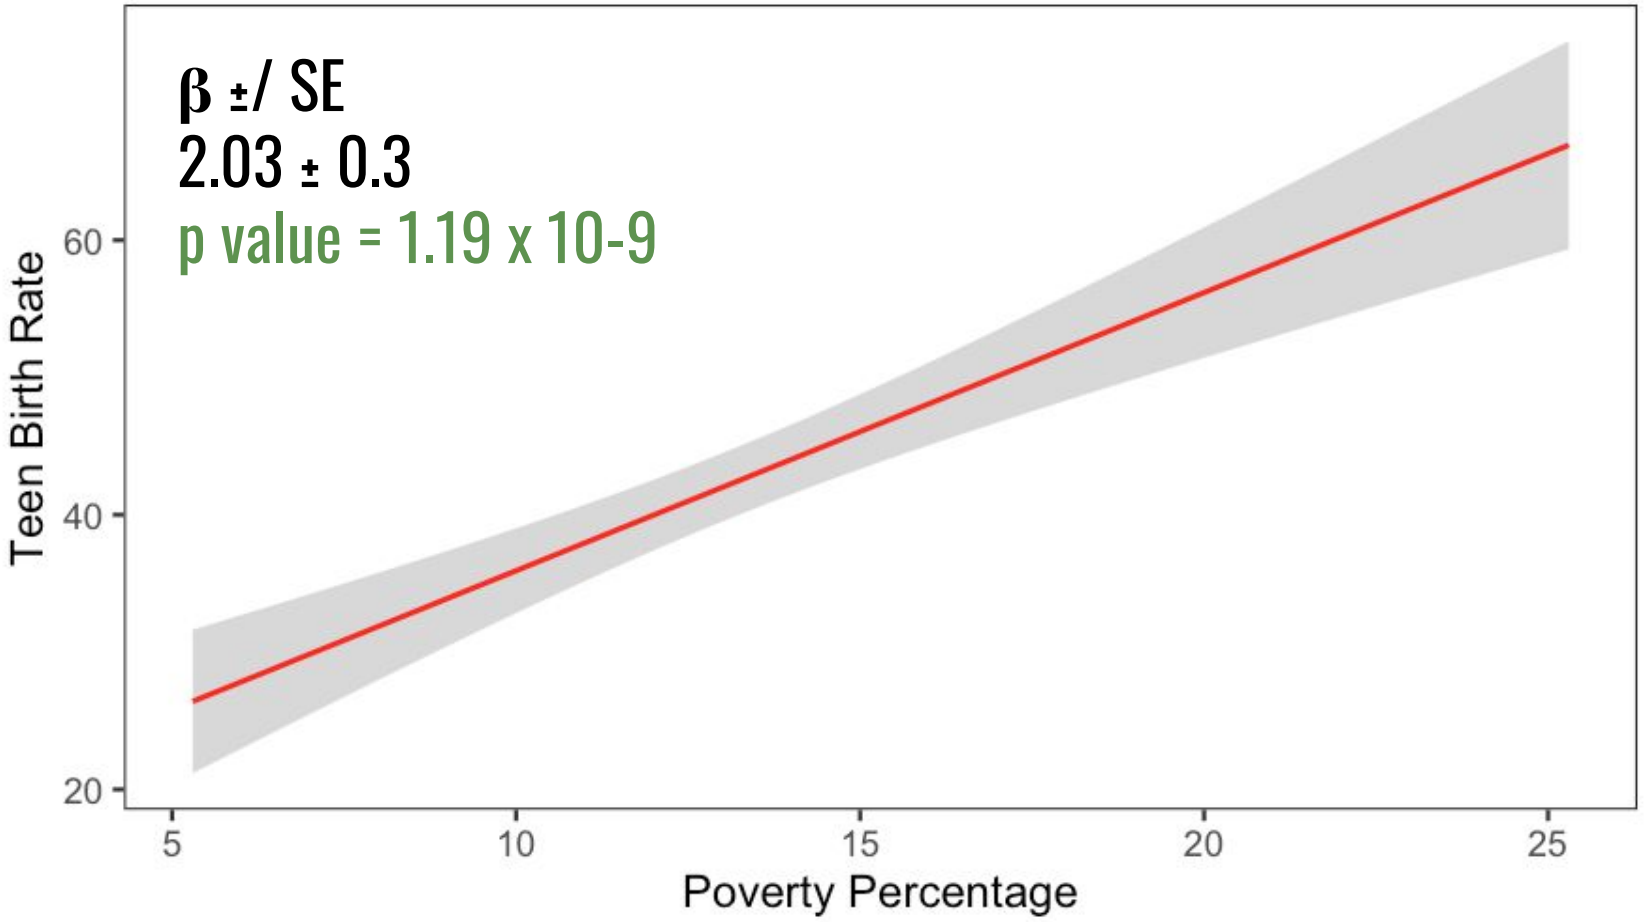
β ±/ SE
2.03 ± 0.3
p value = 1.19 x 10-9
Teen Birth Rate
60
40
20
5
10
15
20
25
Poverty Percentage

p-value : the probability of getting the observed results (or results more extreme) by chance alone

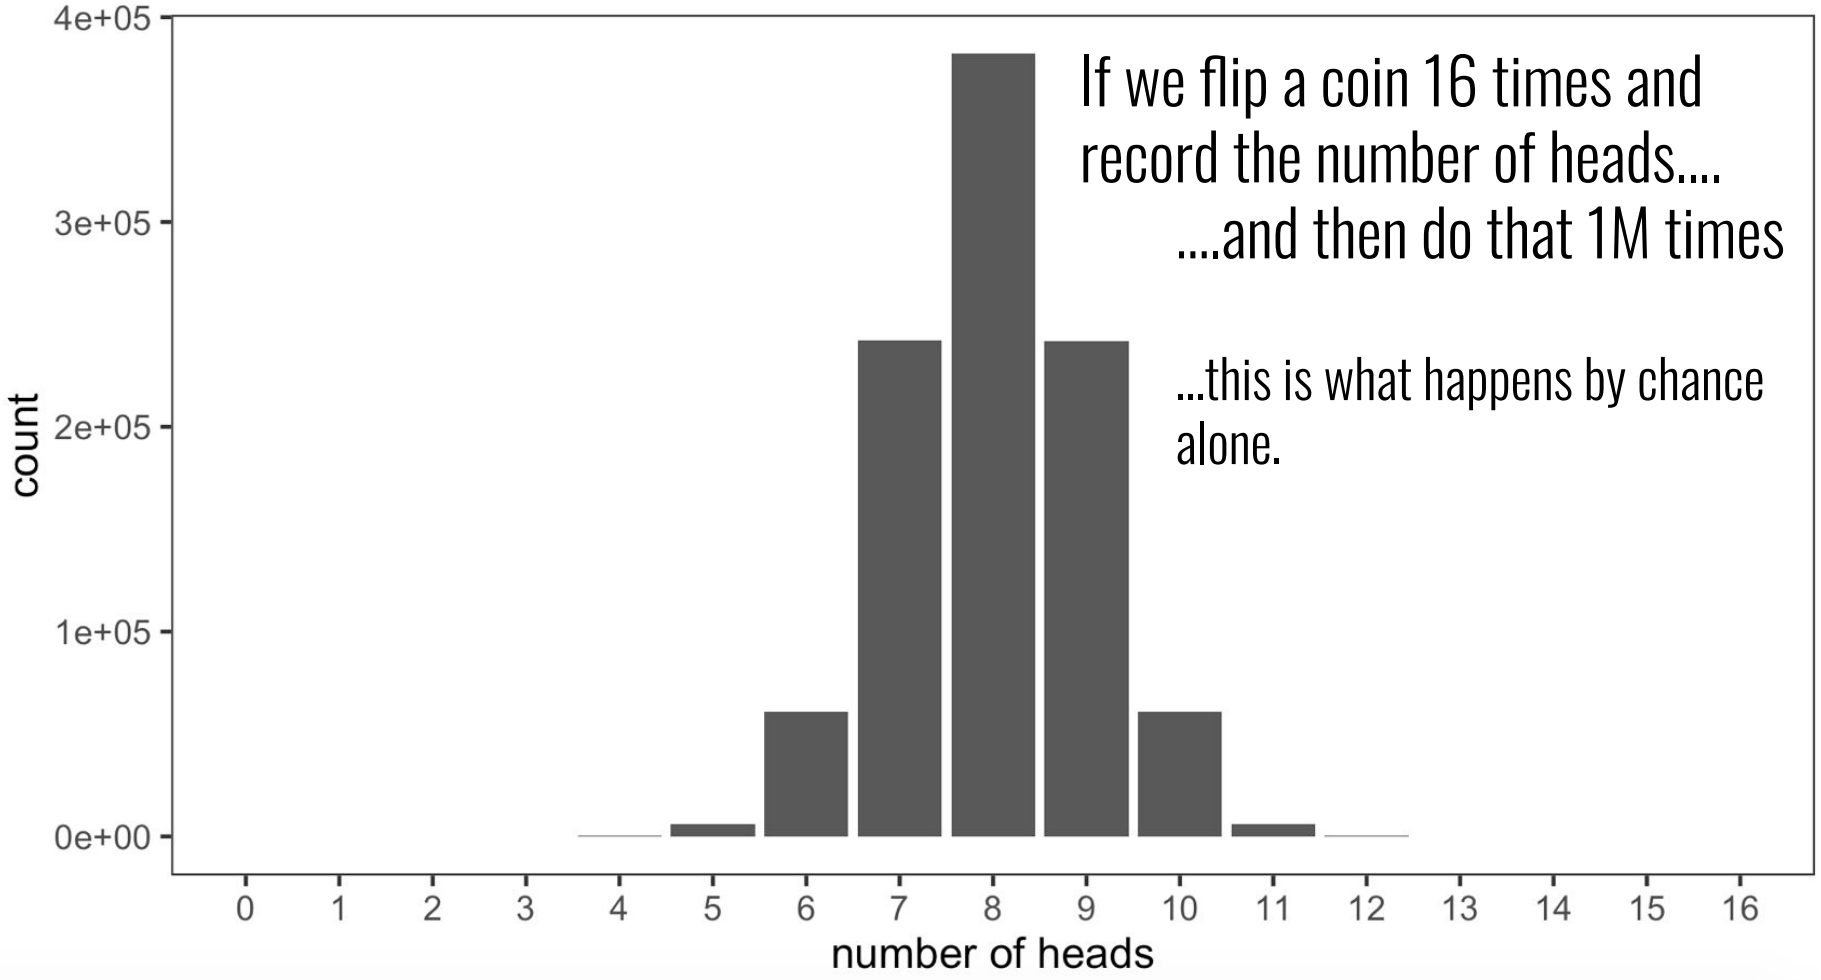
4e+05
3e+05
2e+05
1e+05
0e+00
count
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
If we flip a coin 16 times and record the number of heads....
....and then do that 1M times
...this is what happens by chance alone.

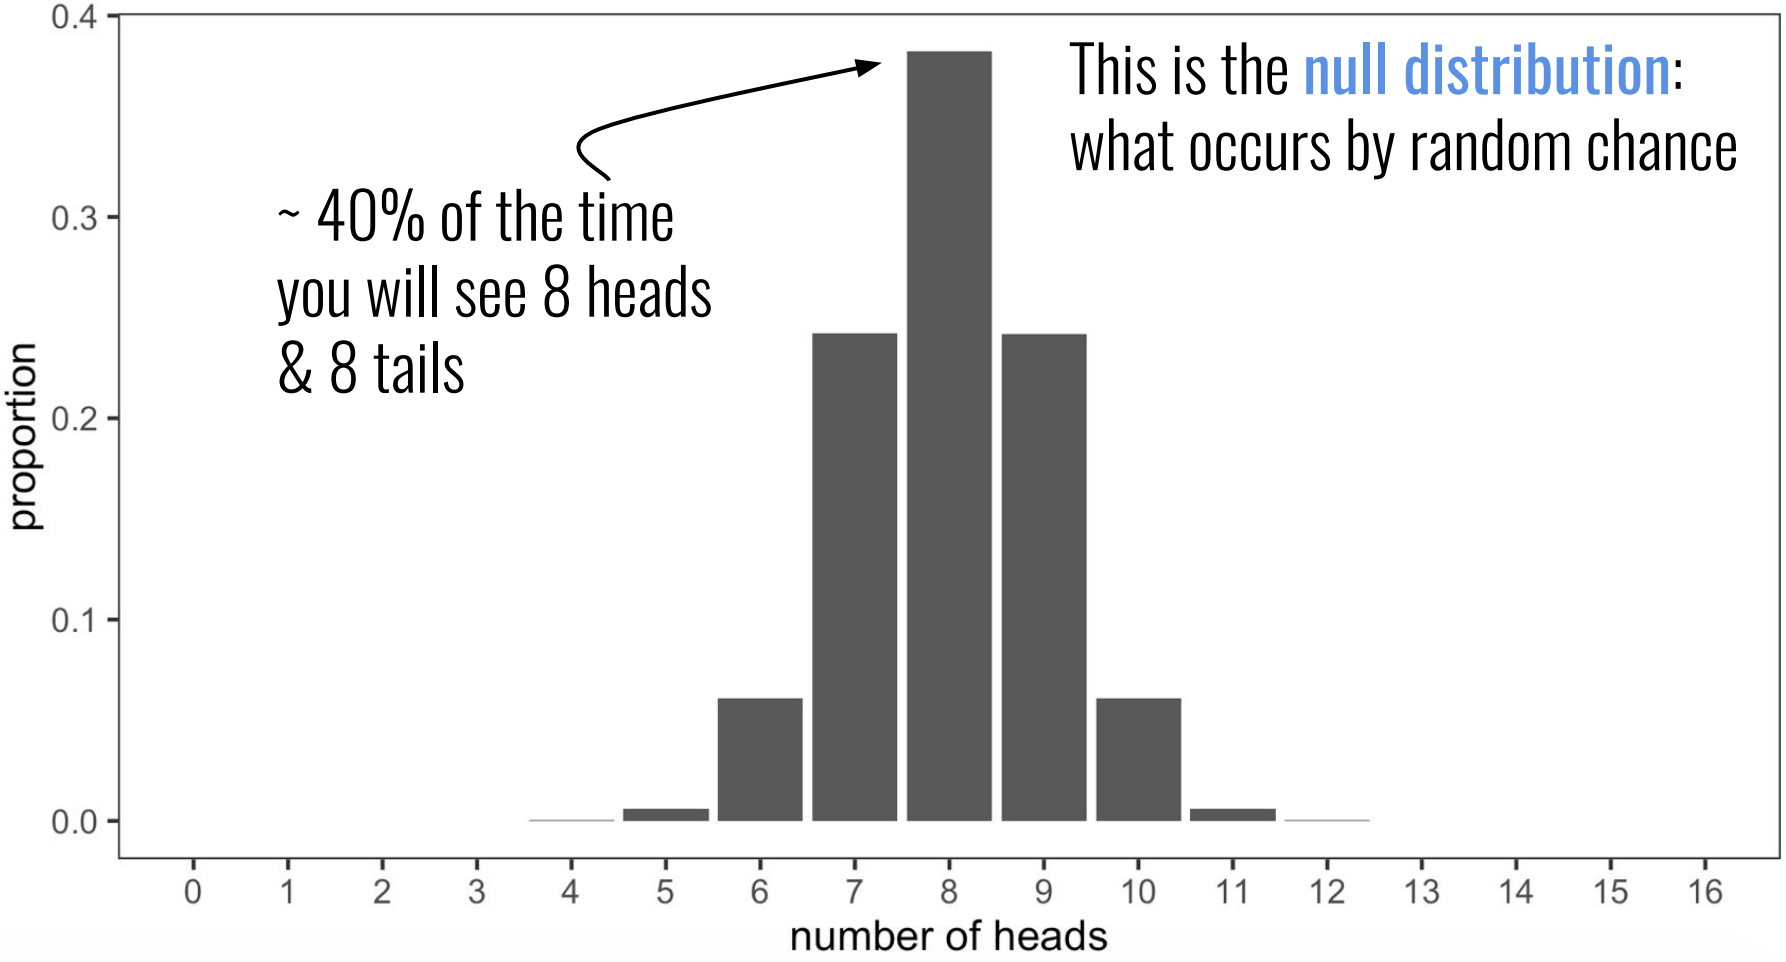
This is the null distribution:
what occurs by random chance
~ 40% of the time
you will see 8 heads
& 8 tails
proportion
0.4
0.3
0.2
0.1
0.0
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads

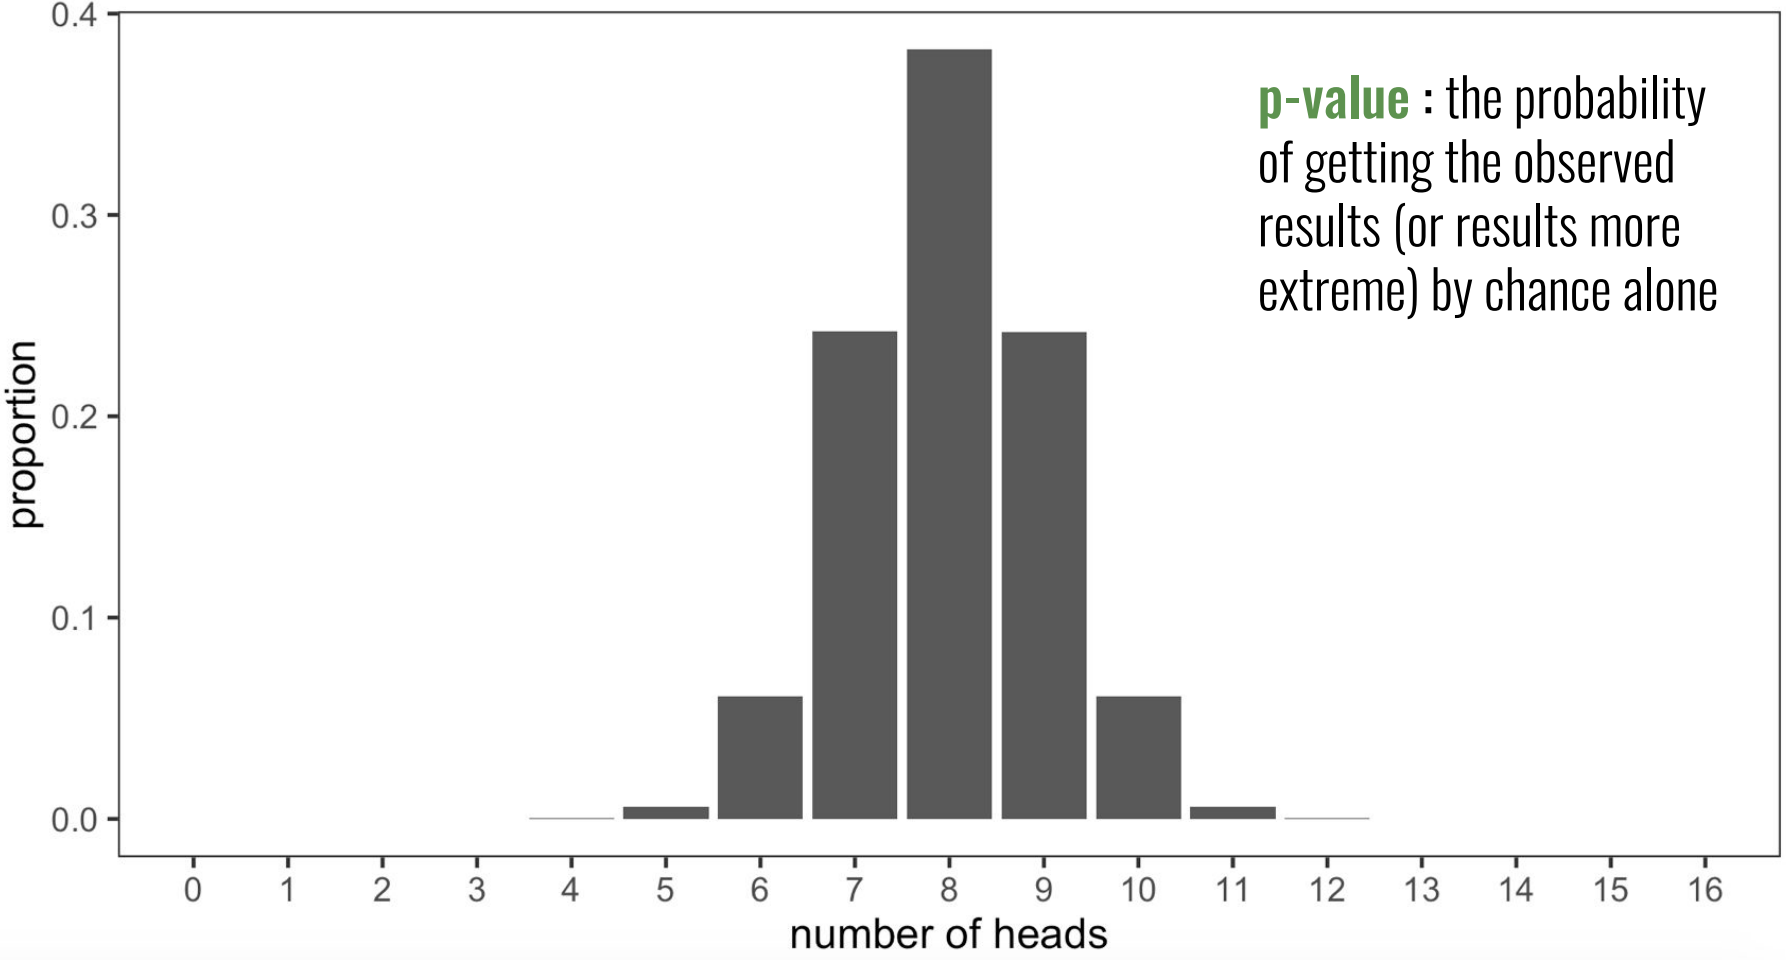

p-value : the probability of getting the observed results (or results more extreme) by chance alone

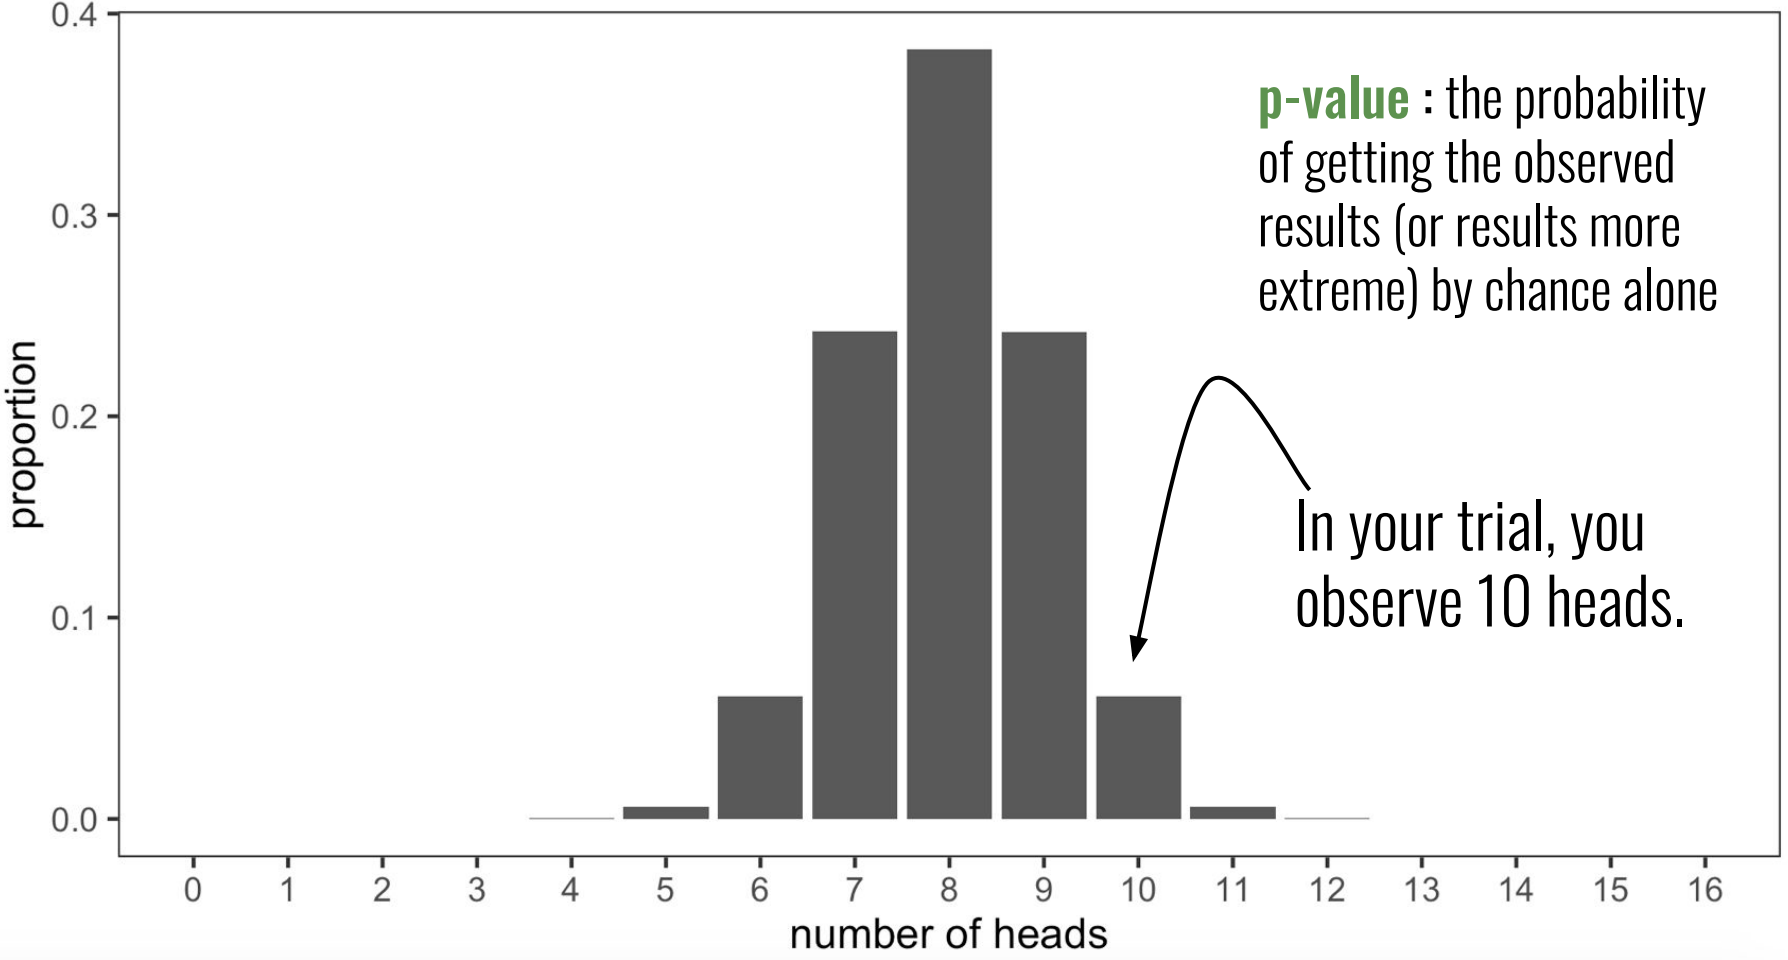
0.4
0.3
0.2
0.1
0.0
proportion
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
p-value : the probability of getting the observed results (or results more extreme) by chance alone
In your trial, you observe 10 heads.

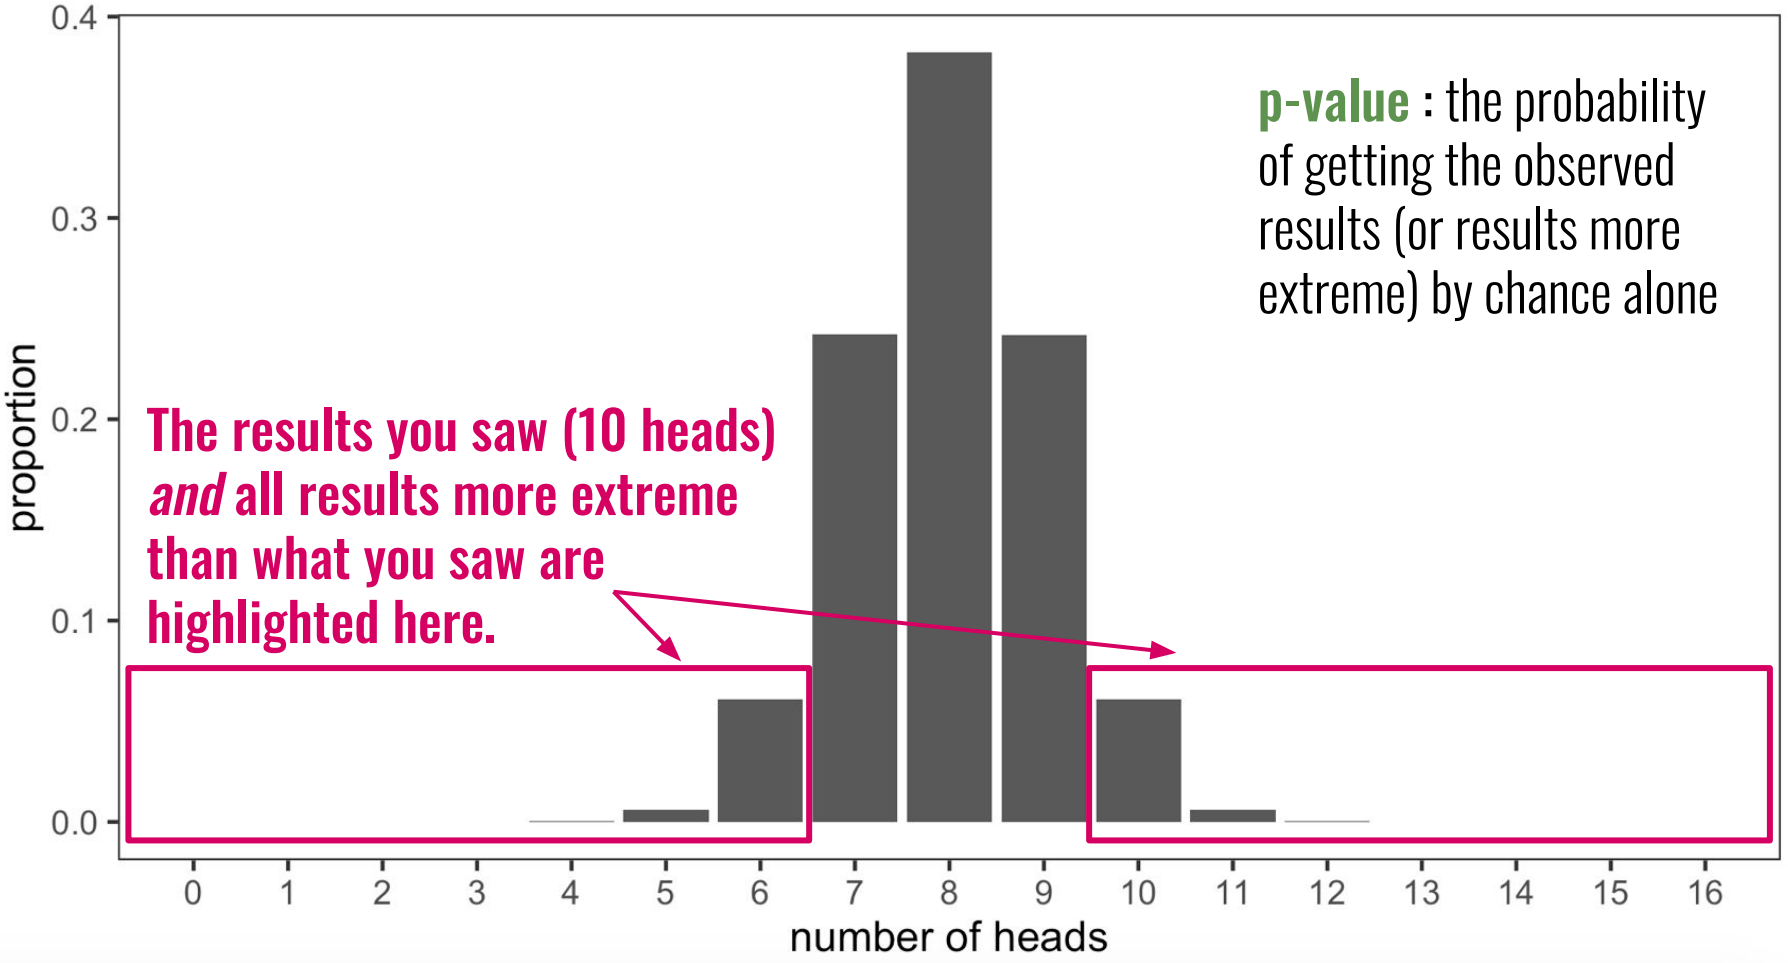
0.4
0.3
0.2
0.1
0.0
proportion
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
p-value : the probability of getting the observed results (or results more extreme) by chance alone
The results you saw (10 heads) *and* all results more extreme than what you saw are highlighted here.

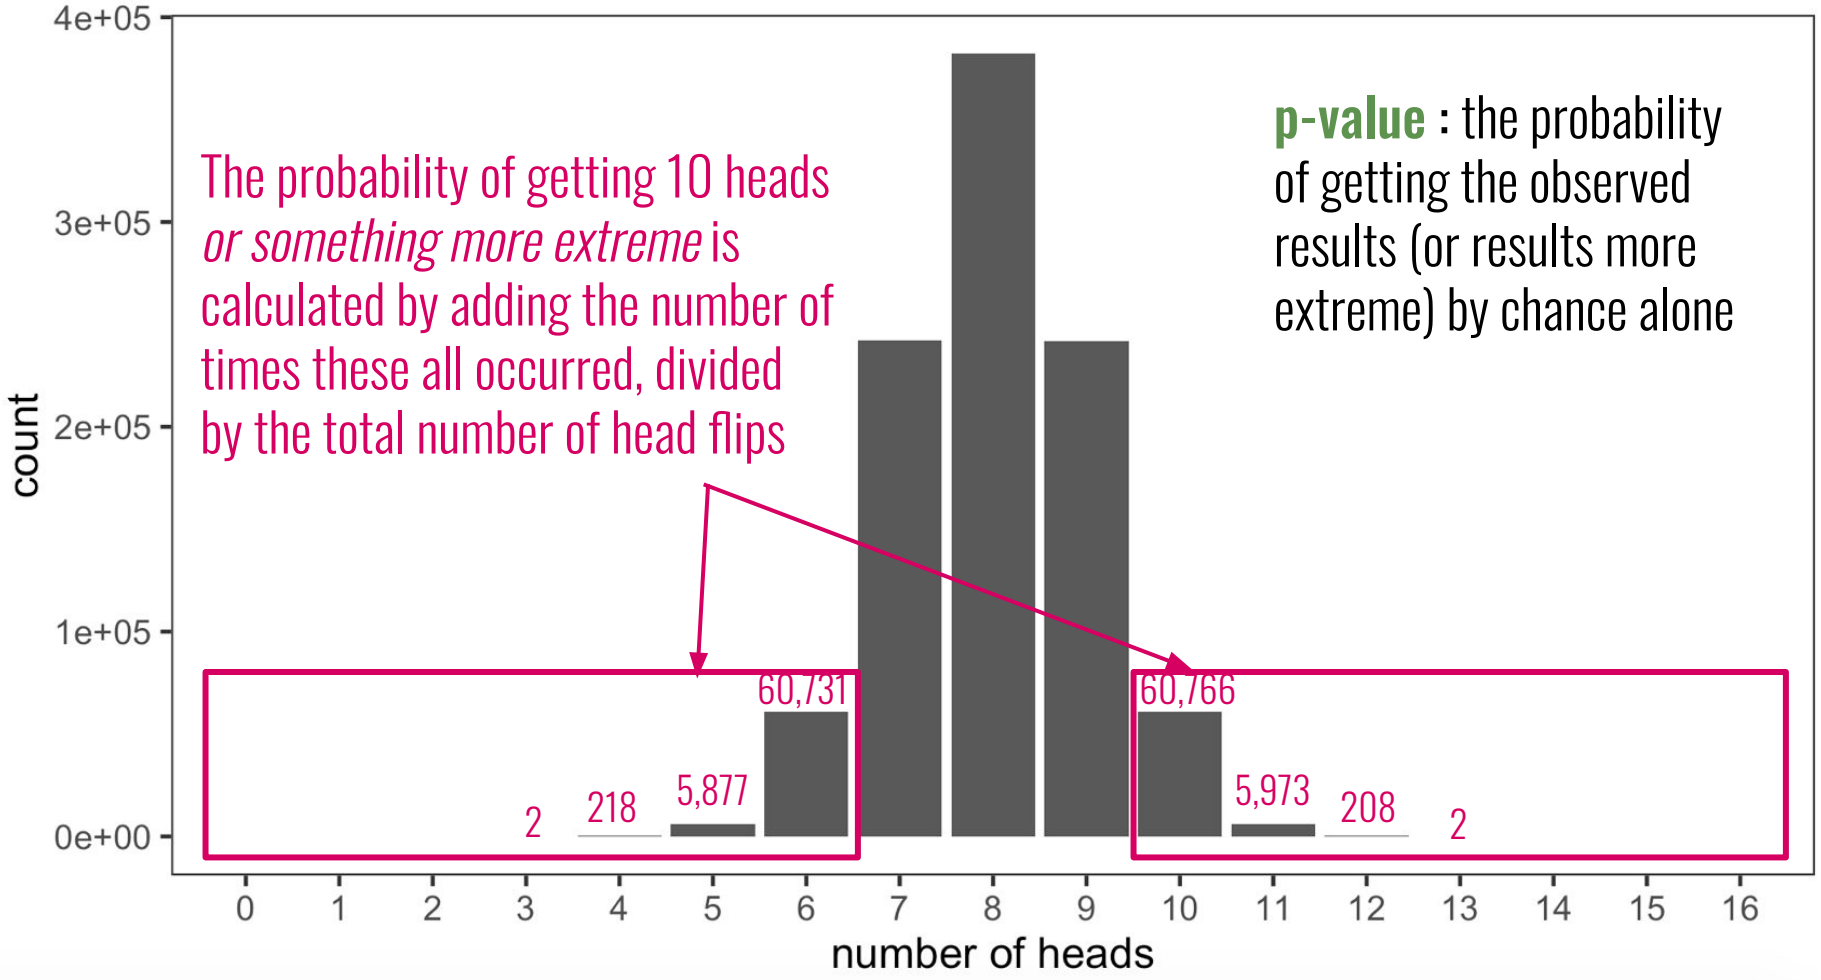
4e+05
3e+05
2e+05
1e+05
0e+00
count
The probability of getting 10 heads *or something more extreme* is calculated by adding the number of times these all occurred, divided by the total number of head flips
p-value : the probability of getting the observed results (or results more extreme) by chance alone
2
218
5,877
60,731
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

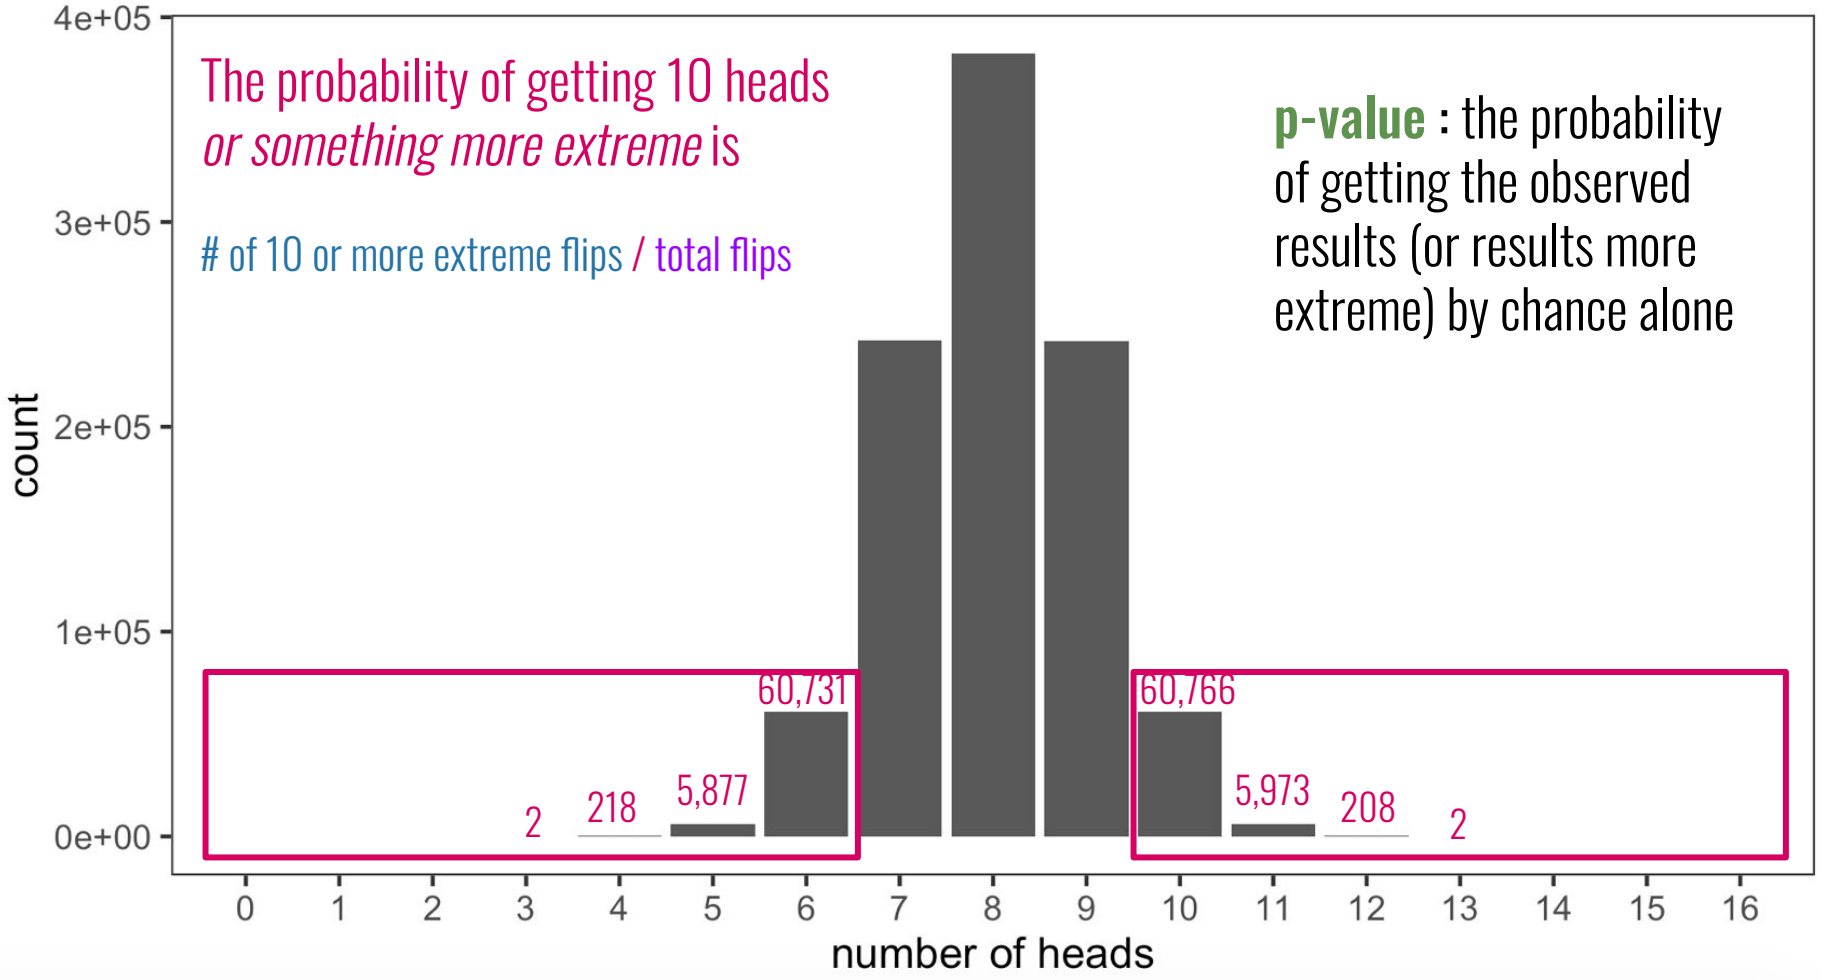
The probability of getting 10 heads
or something more extreme is
# of 10 or more extreme flips / total flips
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,877
60,731
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

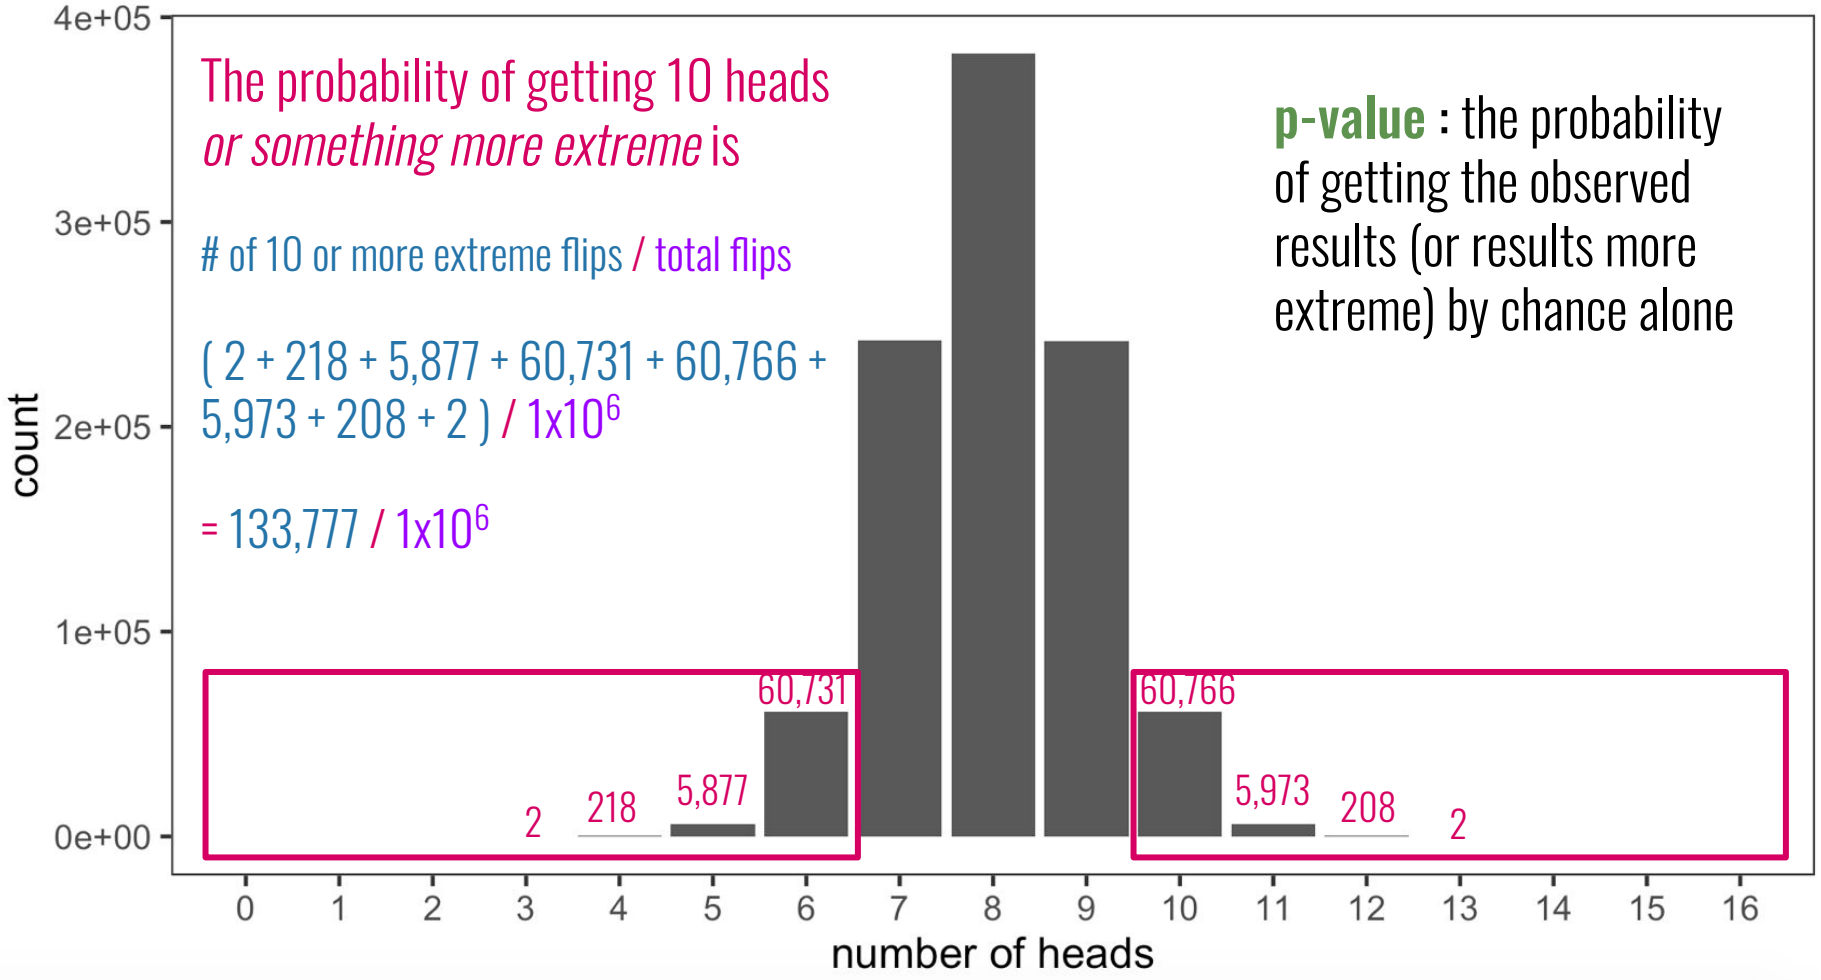
The probability of getting 10 heads
or something more extreme is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,877
60,731
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

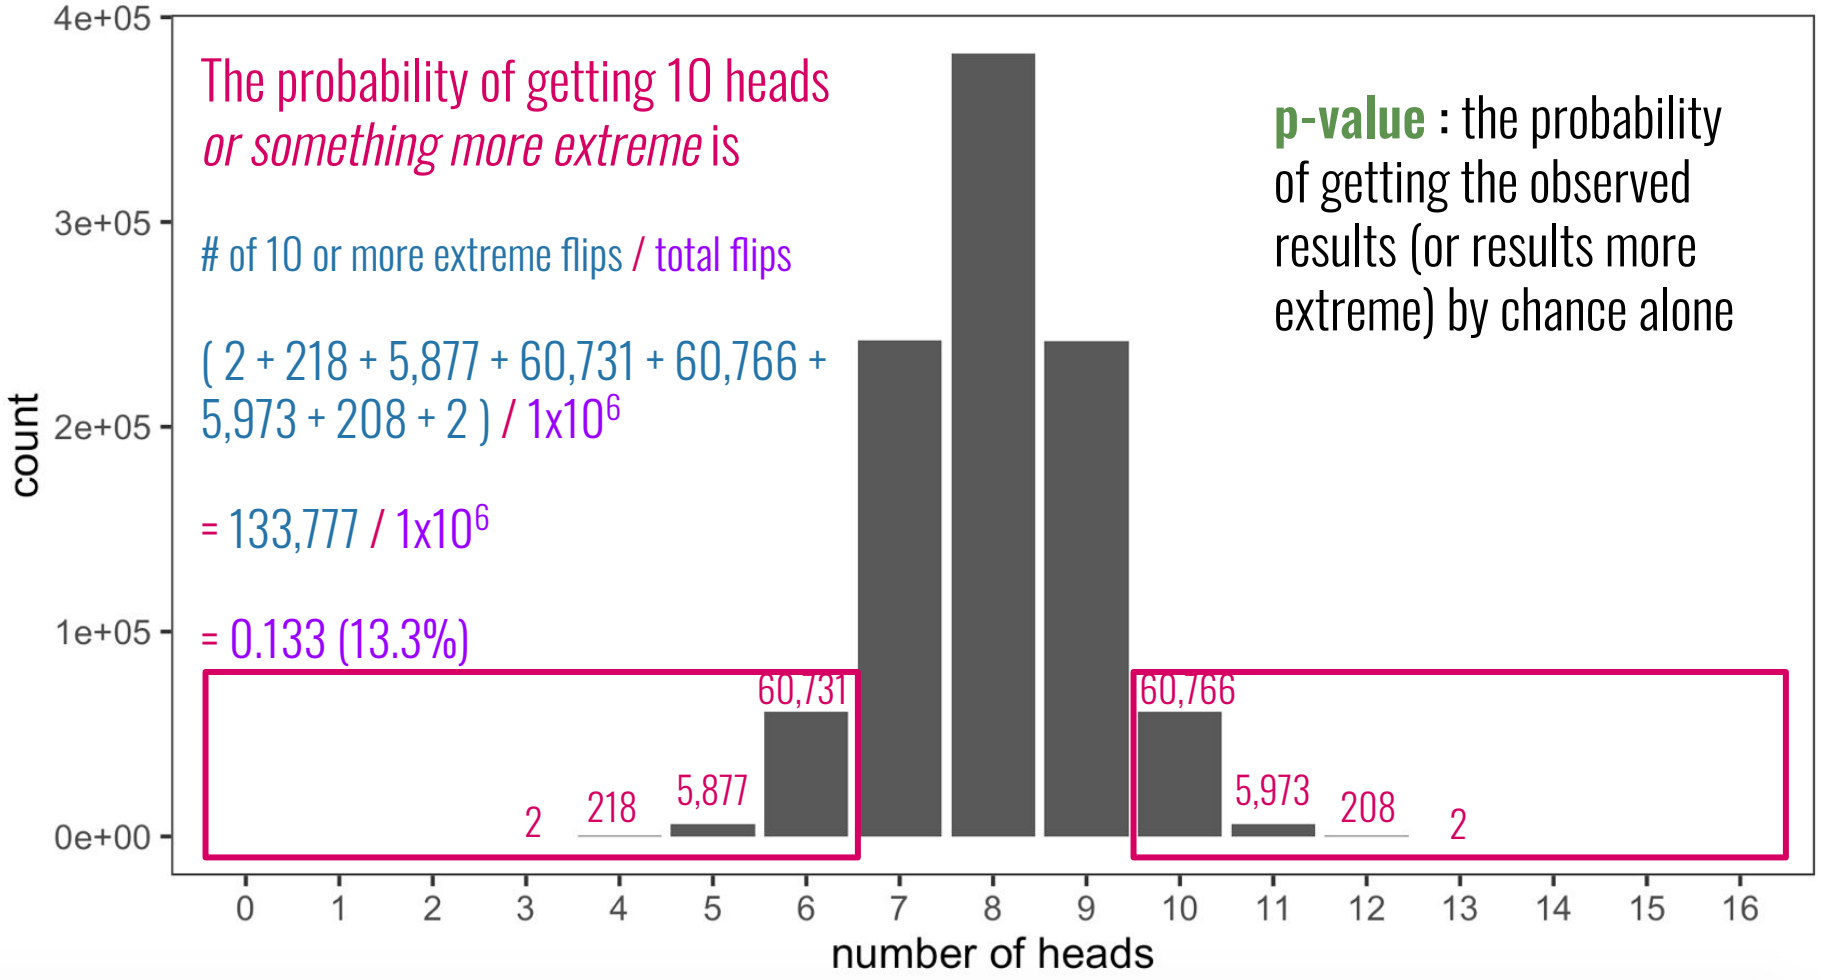
The probability of getting 10 heads
or something more extreme is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
= 0.133 (13.3%)
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,877
60,731
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

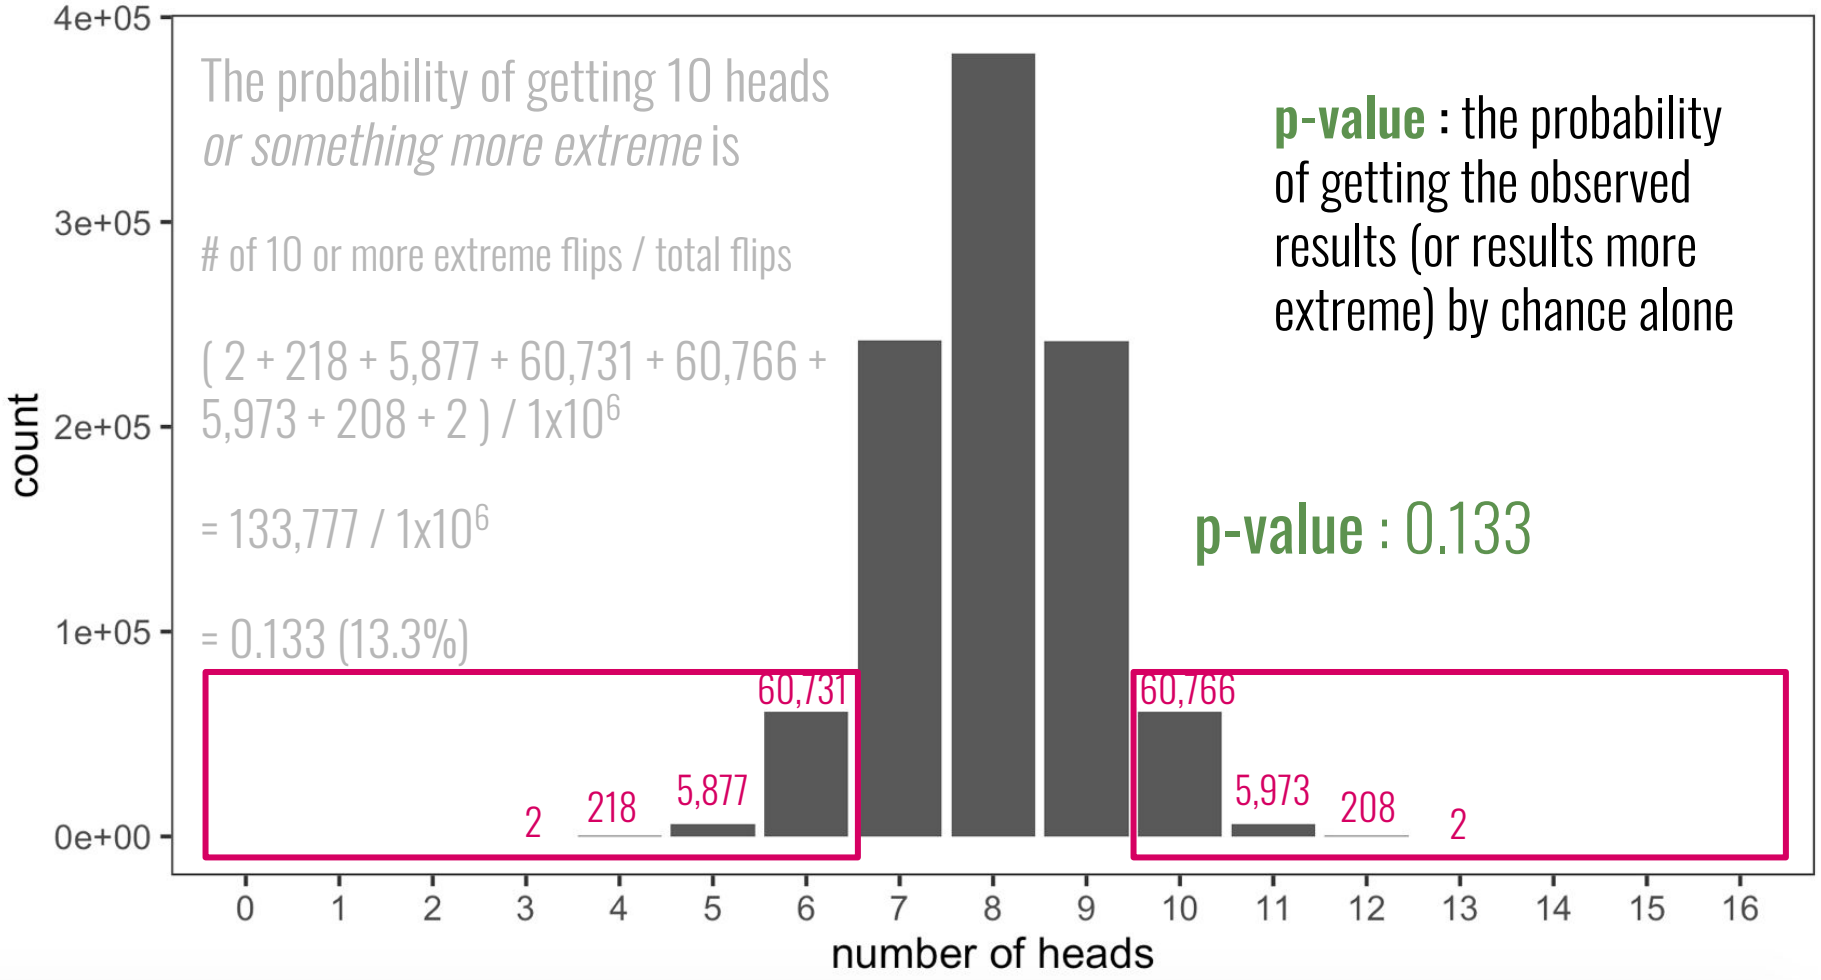
The probability of getting 10 heads
or something more extreme is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
= 0.133 (13.3%)
p-value : the probability of getting the observed results (or results more extreme) by chance alone
p-value : 0.133
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,877
60,731
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

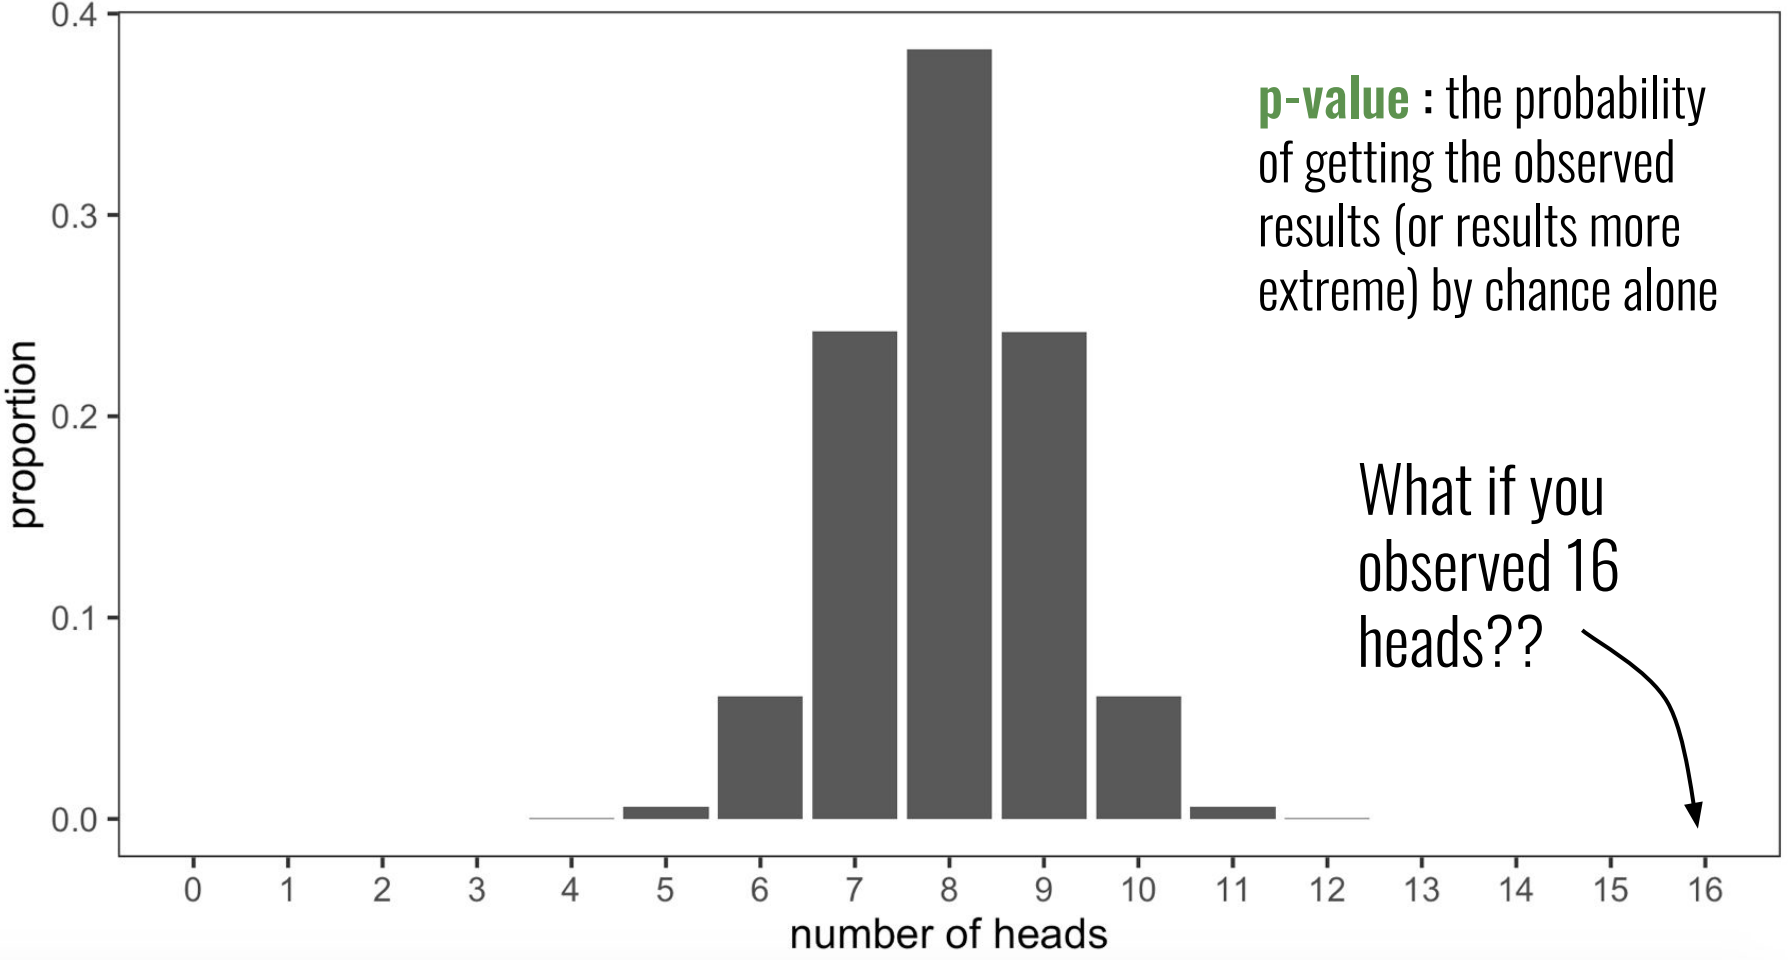

p-value : the probability of getting the observed results (or results more extreme) by chance alone

What if you observed 16 heads??

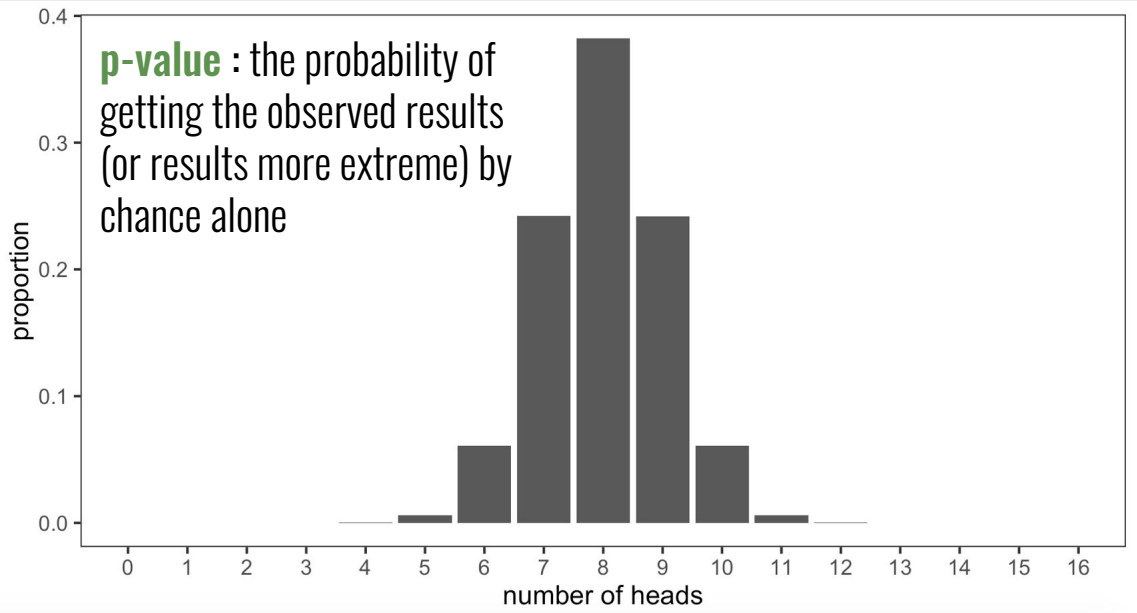

What would be the p-value of you flipping 16 heads?

A
< 0.13

B
> 0.13

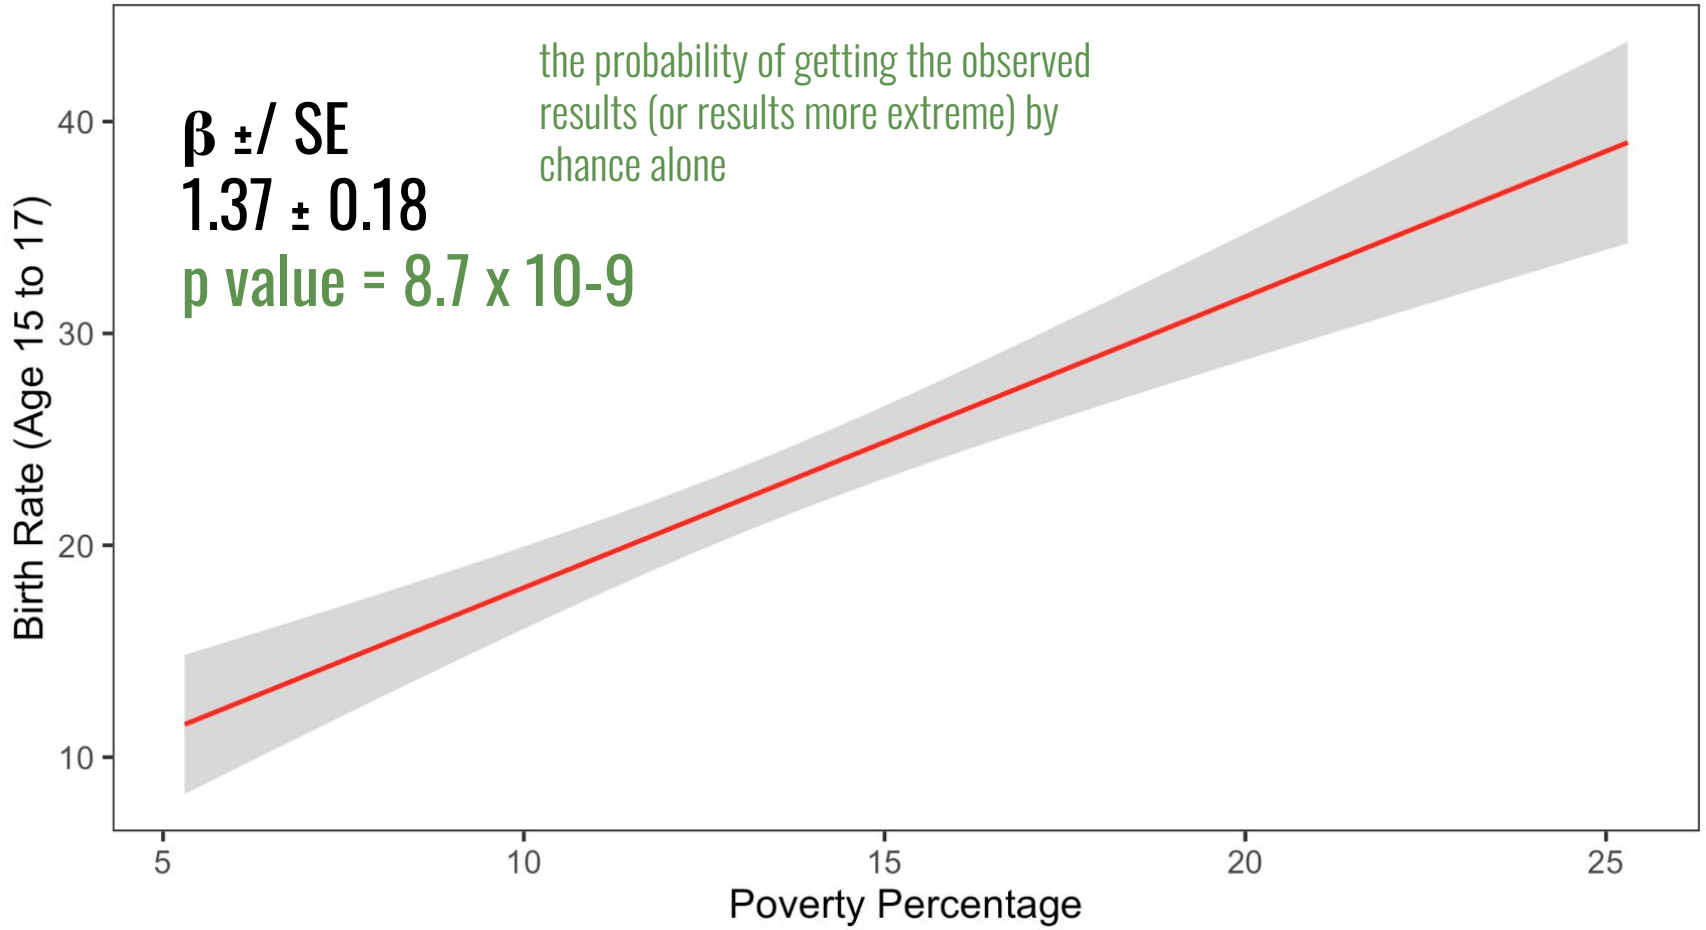
β ±/ SE
1.37 ± 0.18
p value = 8.7 x 10-9
the probability of getting the observed results (or results more extreme) by chance alone
Birth Rate (Age 15 to 17)
40
30
20
10
5
10
15
20
25
Poverty Percentage

Takes into account the effect size (β) and the SE

p-value : the probability of getting the observed results (or results more extreme) by chance alone

Confounding

Small shoes
Not literate
Big shoes
Literate

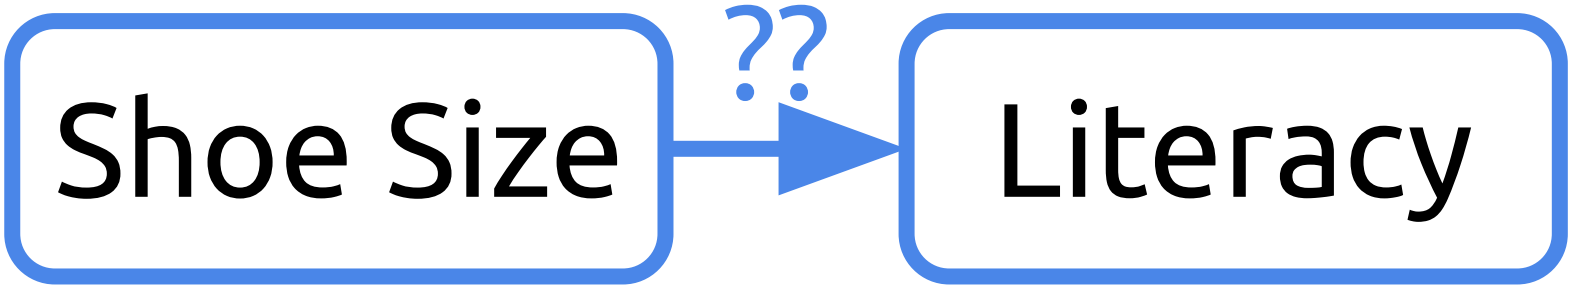
Shoe Size
??
Literacy

Small shoes
Not literate
Child
Big shoes
Literate
Adult

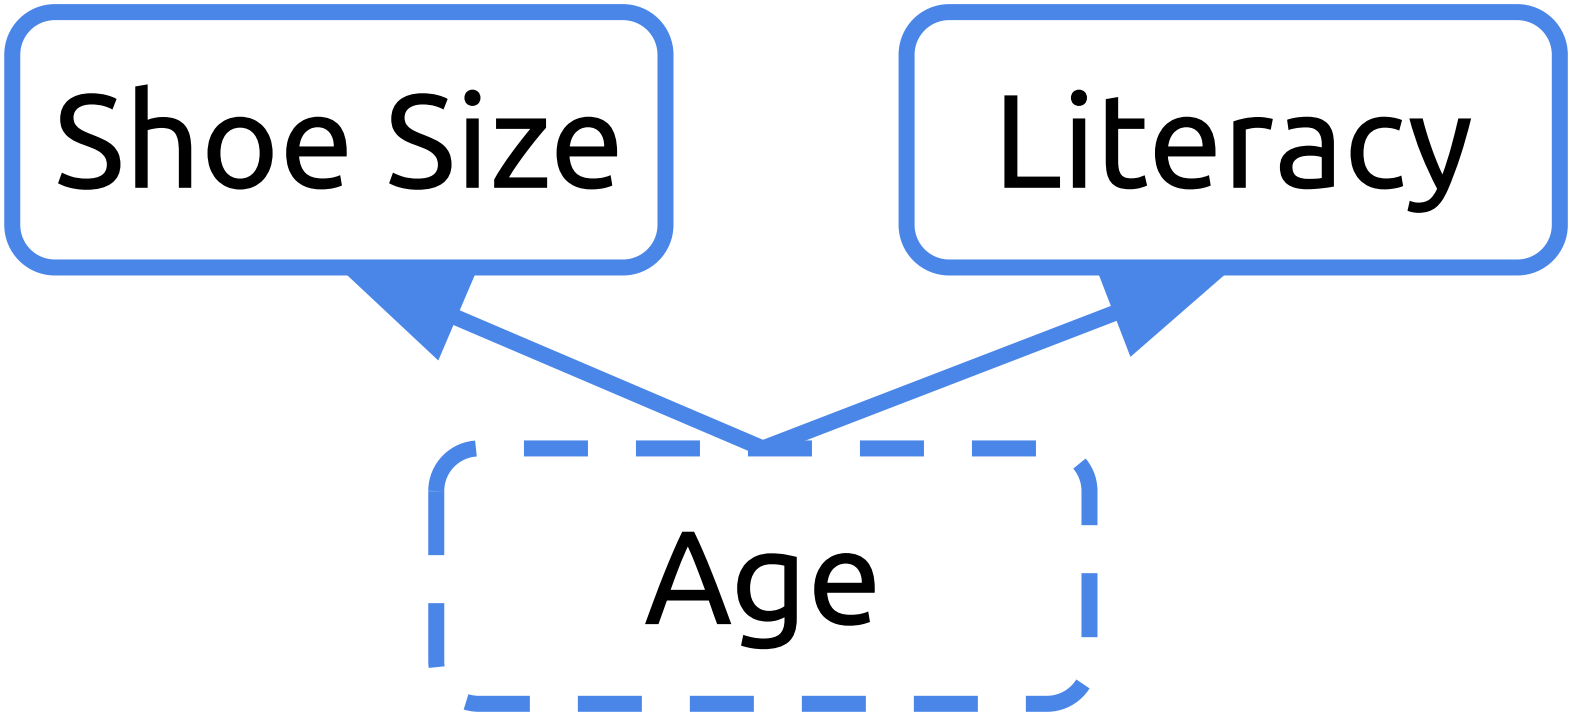
Shoe Size
Literacy
Age

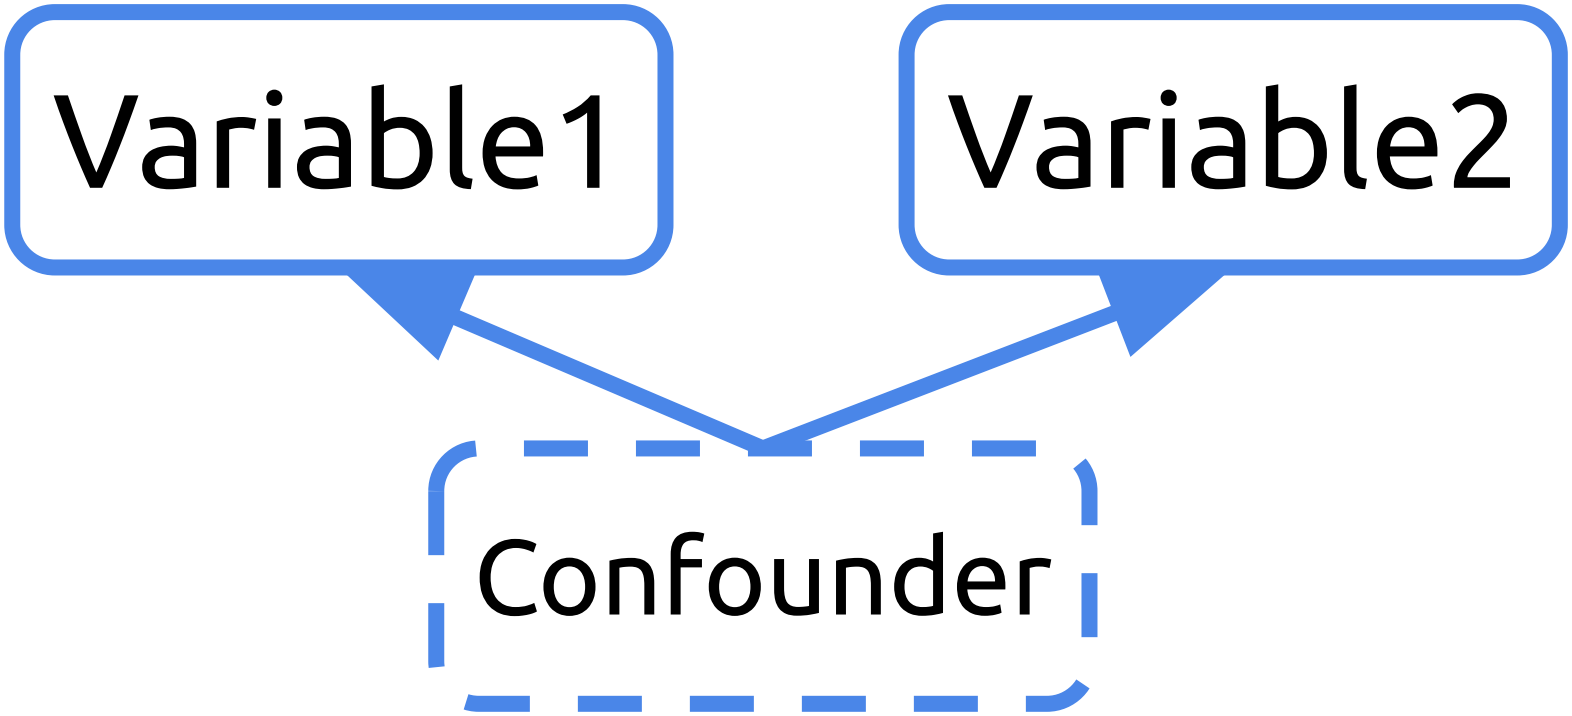
Variable1
Variable2
Confounder

Confounding

popsicles → crime rate

Your analysis sees an increase in crime rate whenever popsicle sales increase. What could confound this analysis?

A
popsicle preference

B
new gun laws

C
temperature

D
changes in popsicle prices

E
new law enforcement officers

Confounding

What are possible confounders for our analysis of the effect of poverty on teen birth rate?

A
I have some ideas

B
Not sure

We'll discuss additional approaches of how to account for confounding in your analysis in the next lecture.

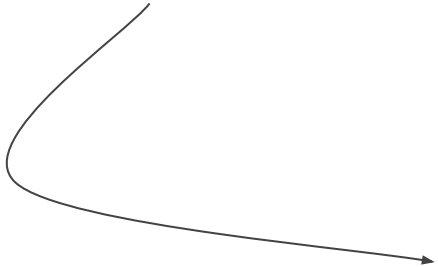

Ignoring confounders will lead you to draw incorrect conclusions from your analyses

Spine Surgery Results

Sample: 400 patients with index vertebral fractures

Vertebroplasty	Conservative care	Relative risk (95% confidence interval)
30/200 (15%)	15/200 (7.5%)	2.0 (1.1–3.6)

subsequent fractures

Eek....looks like vertebroplasty was *way* worse for patients!

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/

But wait...at time of initial fracture...

	Vertebroplasty N = 200	Conservative care N = 200
Age, y, mean ± SD	78.2 ± 4.1	79.0 ± 5.2
Weight, kg, mean ± SD	54.4 ± 2.3	53.9 ± 2.1
Smoking status, No. (%)	110 (55)	16 (8)

Age and weight are similar between groups. **Smoking Status** differs vastly.

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/

So...let's stratify those results real quick

Smoke			**No smoke**		
Vertebroplasty	Conservative	RR (95% confidence interval)	Vertebroplasty	Conservative	RR (95% confidence interval)
23/110 (21%)	3/16 (19%)	1.1 (0.4, 3.3)	7/90 (8%)	12/184(7%)	1.2 (0.5, 2.9)

Risk of re-fracture is now similar within group

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/